

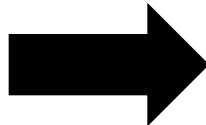
UNDERSTANDING GIT

Joe Gibson (SSCO/HFCS)

joseph.gibson@nasa.gov

Git...

- Is frustrating
- Is complicated
- Gets in your way
- Produces grey hair and stress



- Is intuitive
- Is easy
- Guides development
- Produces happy developers

Complaints/Concerns/Critiques

- “It’s going to take me months to learn Git!”
 - Yes.
- “Git sucks, let’s move back to SVN/CVS/carrier-pigeon”
 - Let’s not be rash!
- “Git seems like a cult... Do we have to drink some sort of poisoned fruit punch?
 - Nope, no punch.
- “What is the difference between a fetch and a merge and a pull and what the hell does checkout actually do???”
 - Glad you asked...

Understanding Git

- Requires:
 - Knowledge of the basic Git commands
 - Knowledge of the basic workflow
 - One-on-one training (please continue to ask us!)
 - *Knowledge of how Git really works under the hood*
 - i.e. What is Git and what is a Git repository?

What is Git?

- Git is a Content Addressable File System
- The content (i.e. files) is stored as (key, value) pairs, where the key (a SHA-1 hash) is derived from the content itself

Write



Read

Key	af5e7e	7a39b4	401fe2	d0f1e6	91e466
Value	A.txt	B.txt	A.txt	C.txt	C.txt

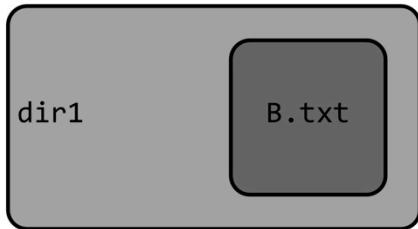
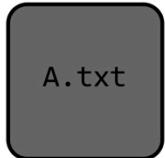
What is a Git repository?

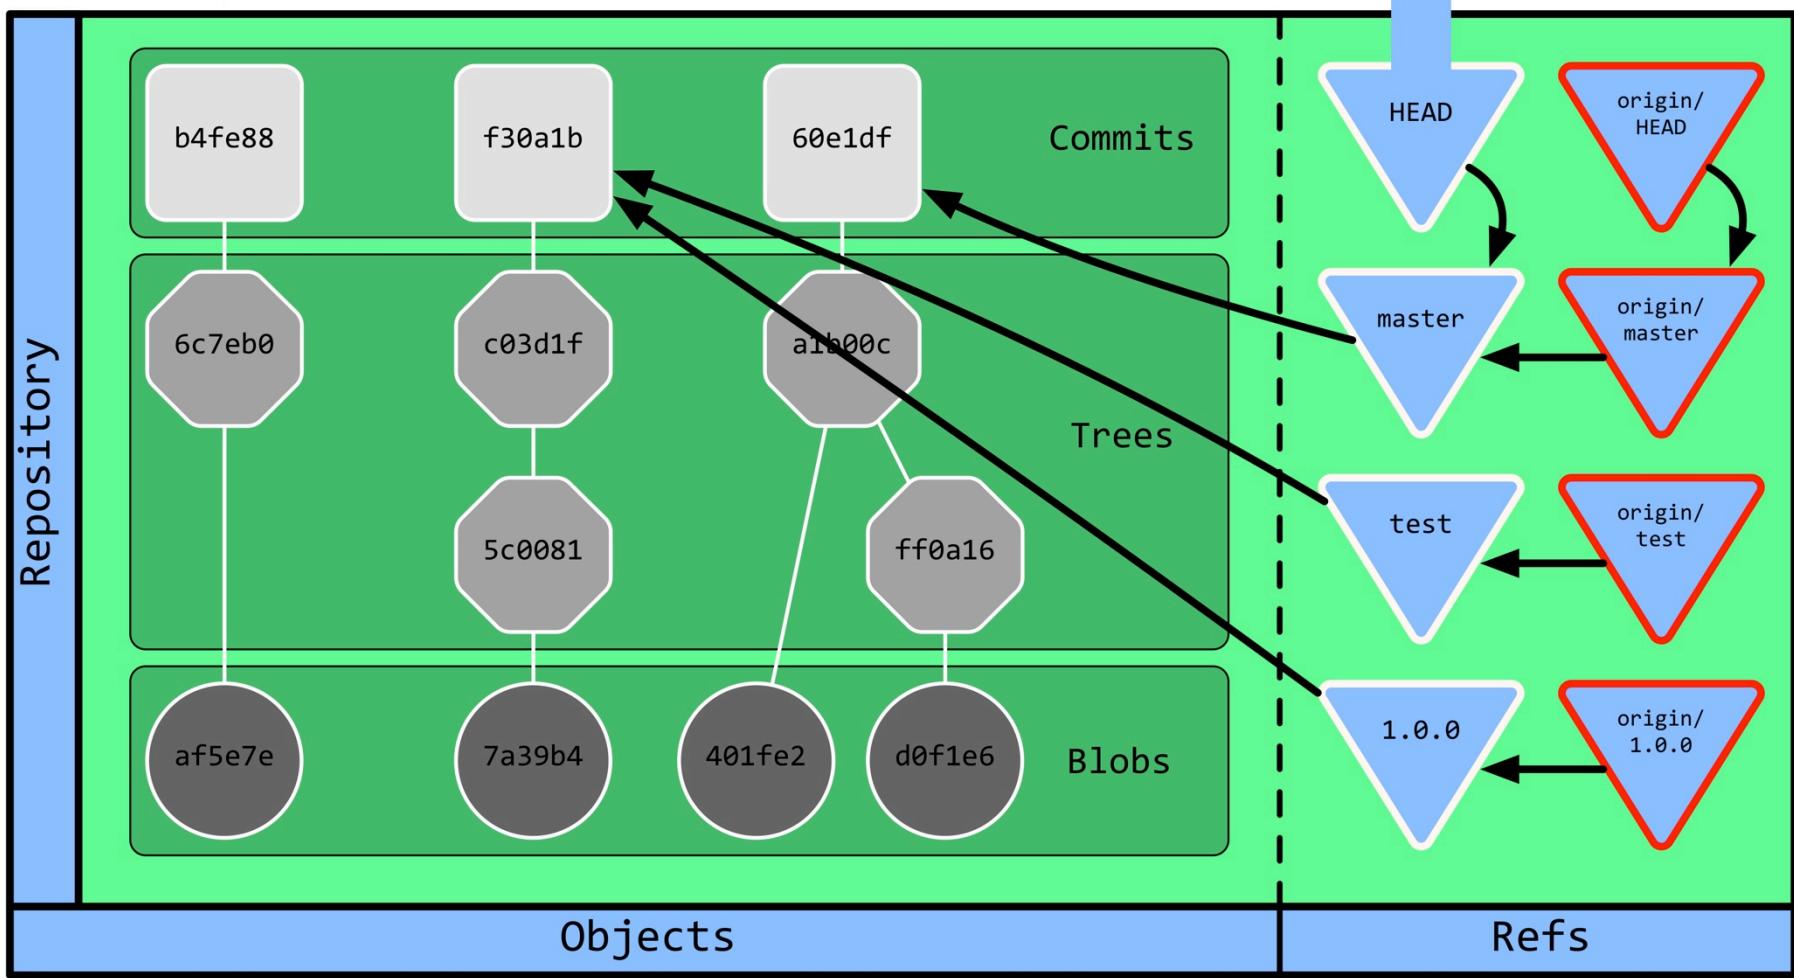
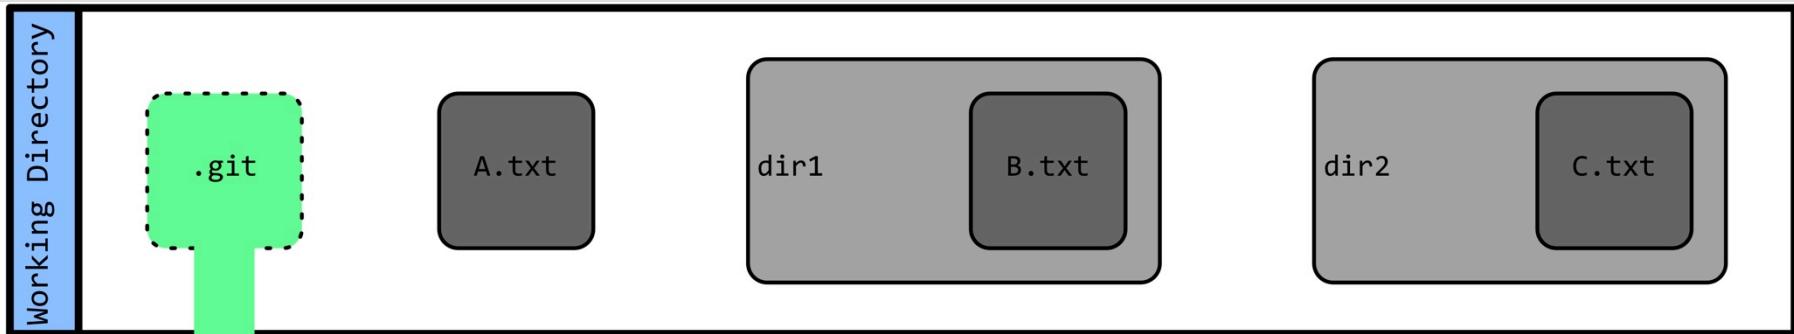
- Each directory we consider a repository has a `.git` folder
- The contents of the `.git` directory is really the repository, and the files at the top level constitute the working directory/working tree

What is a Git repository?

- The `.git` folder contains:
 - **References** (Refs)
 - **Objects**
 - **Commits**
 - **Trees**
 - **Blobs**

Working Directory





```
$ tree -a
.
├── .git
│   ├── COMMIT_EDITMSG
│   ├── HEAD
│   ├── ORIG_HEAD
│   ├── config
│   ├── description
│   ├── hooks
│   ├── index
│   ├── info
│   │   └── exclude
│   ├── logs
│   │   ├── HEAD
│   │   └── refs
│   │       └── heads
│   │           ├── master
│   │           └── test
│   ├── objects
│   │   ├── 34
│   │   │   └── 5e6aef713208c8d50cdea23b85e6ad831f0449
│   │   ├── 3e
│   │   │   └── 6d9614ec3cf76e440aca6dafa8e6e0c14f34a
│   │   ├── 43
│   │   │   └── 222cd2e96acac8c83f5162226721f3351c2208
│   │   ├── 46
│   │   │   └── 2221e11fd20cc6dc0203c38fe85b88047721d0
│   │   ├── info
│   │   └── pack
│   └── refs
│       ├── heads
│       │   ├── master
│       │   └── test
│       └── tags
│           └── 1.0.0
└── A.txt
└── dir1
    └── B.txt
└── dir2
    └── C.txt
```

Refs

- **remotes**
- **branches** (i.e. ‘**heads**’ of branches)
- **tags**
- Either way... They’re just pointers!

```
$ tree .git/refs/
refs/
├── heads
│   └── geons
│       └── master
├── remotes
│   └── origin
│       ├── geons
│       ├── HEAD
│       └── master
└── tags
```

Branches

- **Branches** point to the latest commit in a series of commits, i.e. ‘head’ of commit chain

```
$ cat .git/refs/heads/geons
c74945f3f0878aadb4a6f7126c9cd414fe9569b4

$ git cat-file -p c74945
tree c96564d63ad92d90d5225151e7b64a5994ba4fc9
parent cd4f876ba710b2fc48b414336cdd0d30781a3bee
author Joe Gibson <joseph.gibson@nasa.gov> 1469022963 -0400
committer Joe Gibson <joseph.gibson@nasa.gov> 1469022963 -0400
```

Remove GEONS CCSDS binary file and update CMakeList

HEAD

- **HEAD** is a special **ref**
 - Points to another ref (usually)
 - Always ends up pointing to the commit from which the **working directory** is populated
- HEADs point to refs
- Refs point to commits

HEADs and Refs

```
$ cat .git/HEAD
```

```
ref: refs/heads/master
```

```
$ cat .git/refs/heads/master
```

```
cd4f876ba710b2fc48b414336cdd0d30781a3bee
```

```
$ cat .git/refs/remotes/origin/HEAD
```

```
ref: refs/remotes/origin/master
```

```
$ cat .git/refs/remotes/origin/master
```

```
cd4f876ba710b2fc48b414336cdd0d30781a3bee
```

Objects

- **Commits**
- **Trees**
- **Blobs**

Commits

- Tree SHA-1
- Parent Commit SHA-1
- Author/Committer Information
- Blank Line
- Commit Message

```
$ git cat-file -p cd4f876
tree 723a574a17275f194faef6de2c8a2884fd320a37
parent f6e5d7fd3dcf3b1c483feb56ed00113ccf809827
author Joe Gibson <joseph.gibson@nasa.gov> 1469022841 -0400
committer Joe Gibson <joseph.gibson@nasa.gov> 1469022841 -0400
```

Roll back to using OS_Stat instead of CFS_IsValidFilename

Trees

- Basically a directory
- **Sub-tree SHA-1s**
- **Blob SHA-1s**

```
$ git cat-file -p 723a57
100644 blob 73322dc022ec7ecb80b254824756b0b8a94e4dca CMakeLists.txt
040000 tree f15e3959143fa3c7b416e0797ad3eb4328900d7e fsw
```

Blobs

- Content!

```
$ git cat-file -p 73322dc

cmake_minimum_required(VERSION 2.6.4)
project(FC CXX)

# Force all c files to be compiled with g++
file(GLOB_RECURSE C_FILES ${CMAKE_CURRENT_SOURCE_DIR} *.c)
set_source_files_properties(${C_FILES} PROPERTIES LANGUAGE CXX)

include_directories(fsw/src)
...
```

Blobs

- As a side note, the SHA-1 hash for the blob is created by concatenating and hashing the following string:
 - The word “blob”
 - A space
 - The size of the file in bytes
 - A null byte (\0)
 - The full contents of the file
- Note that the filename is **NOT** used in calculating the SHA-1...
 - This makes it easier to do things like renaming, moving, and reducing the total number of blobs

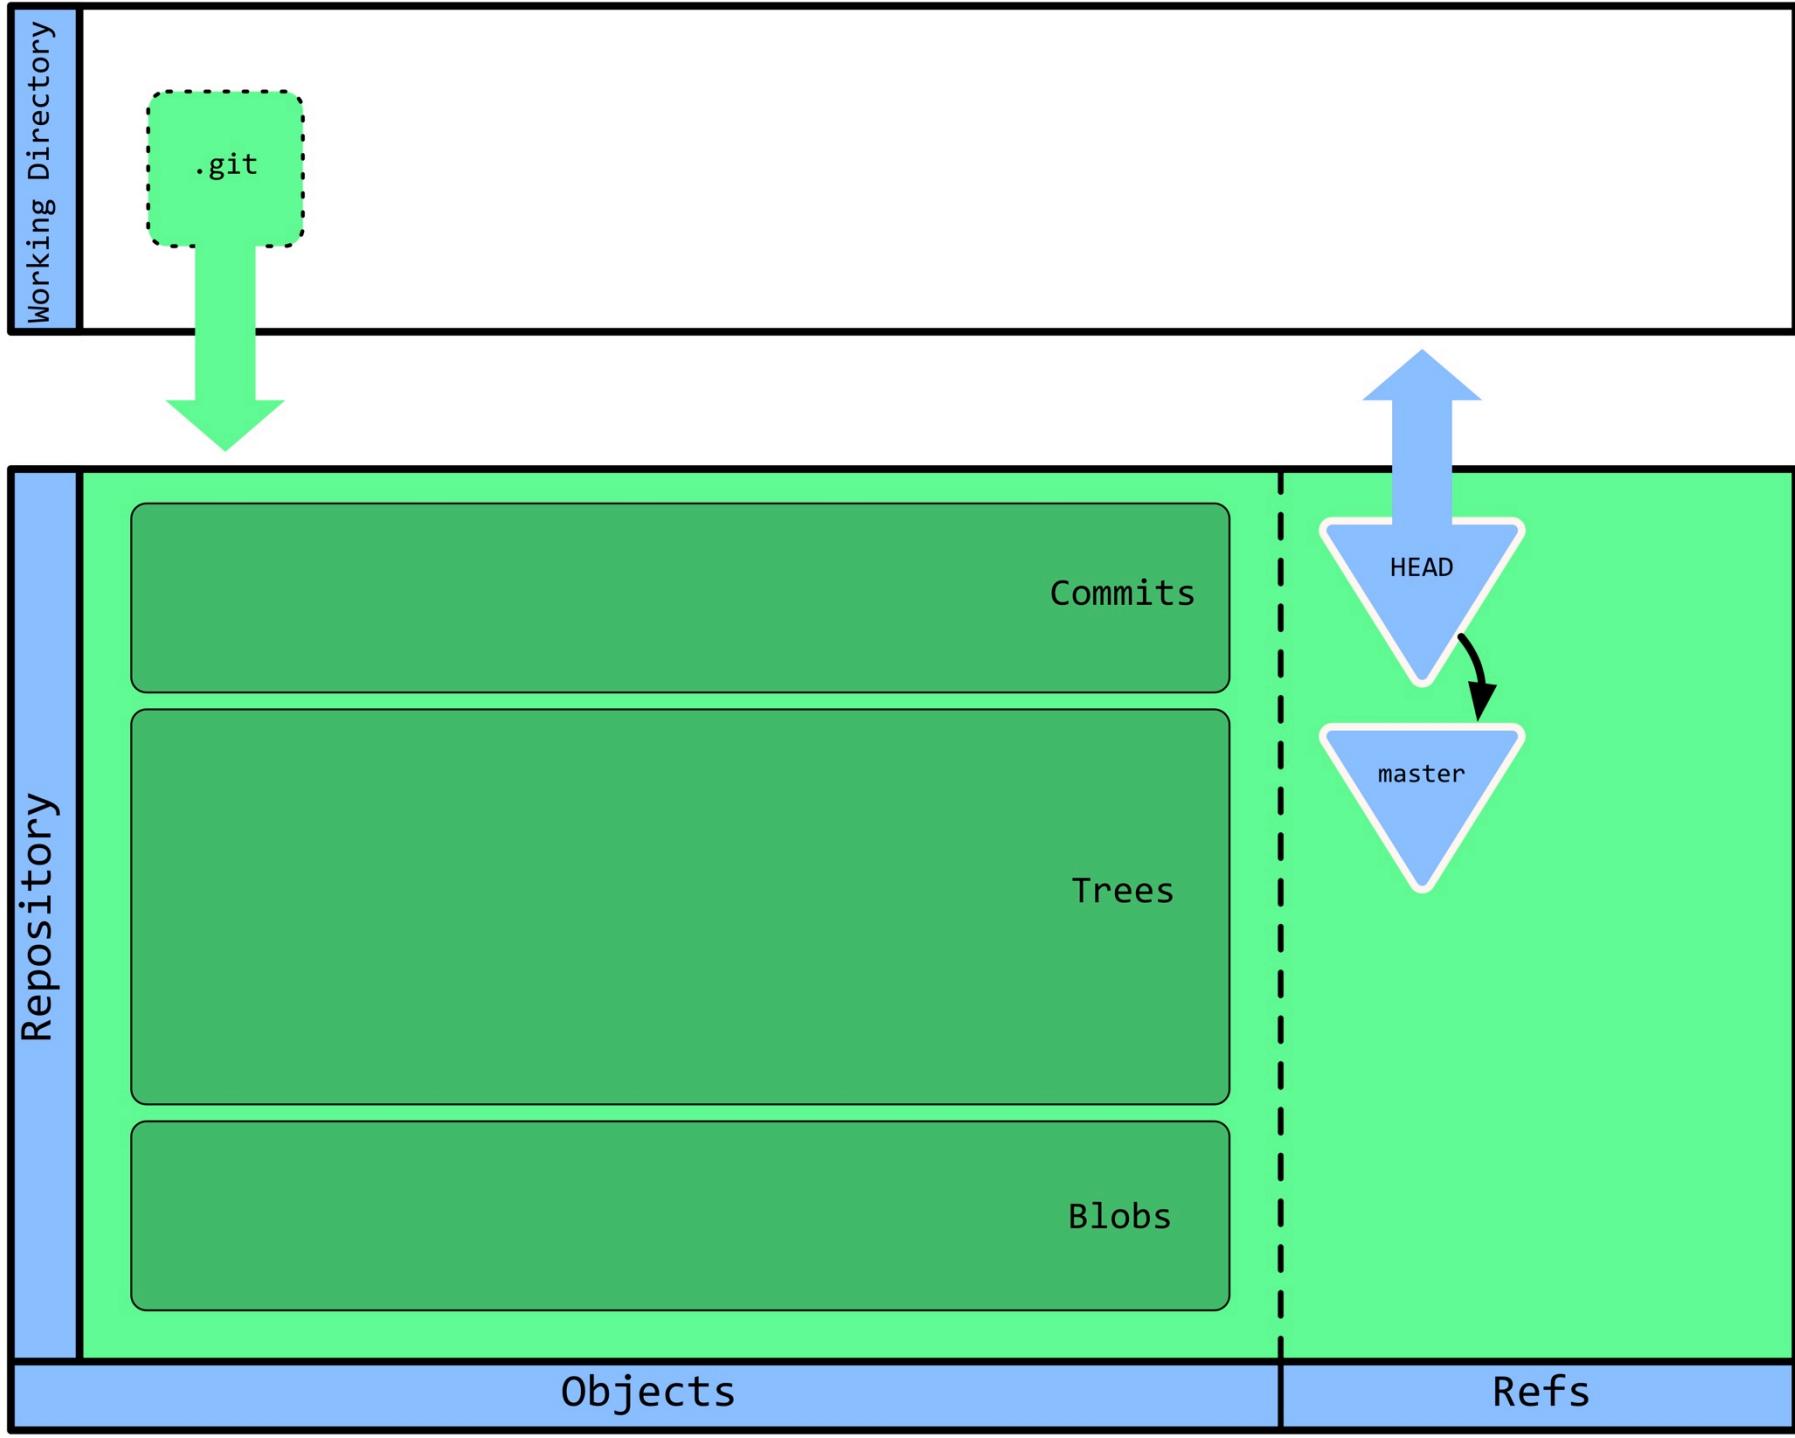
```
$ cat C.txt
Test
$ git hash-object C.txt
345e6aef713208c8d50cdea23b85e6ad831f0449
$ echo -e "blob $(wc -c C.txt | awk '{print $1}')\0$(cat C.txt)" | shasum
345e6aef713208c8d50cdea23b85e6ad831f0449
$ echo -e "blob 5\0Test" | shasum
345e6aef713208c8d50cdea23b85e6ad831f0449
```

Okay, so what?

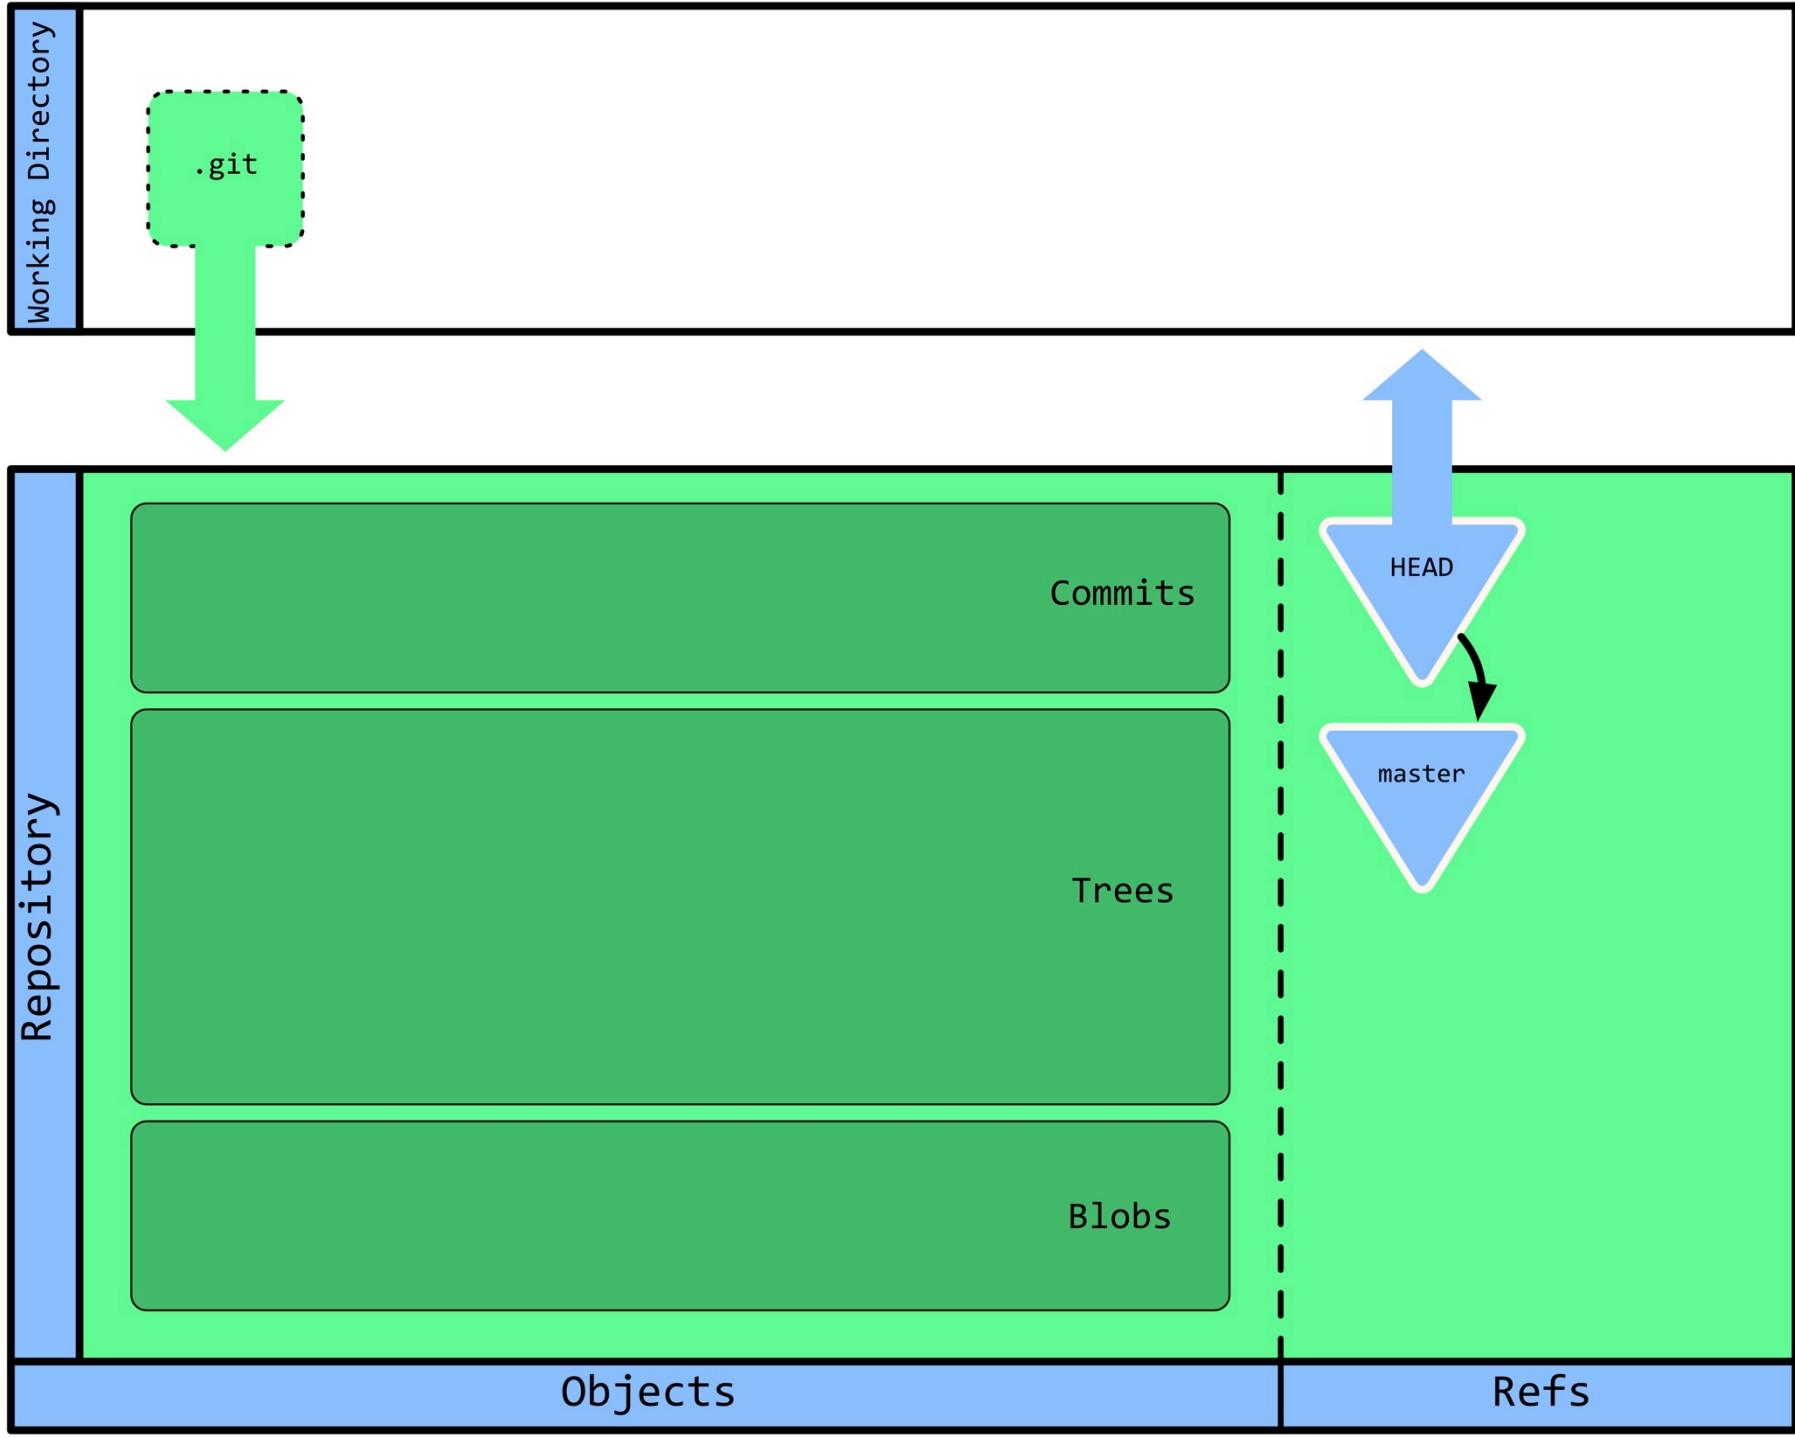
A Detailed Walkthrough...

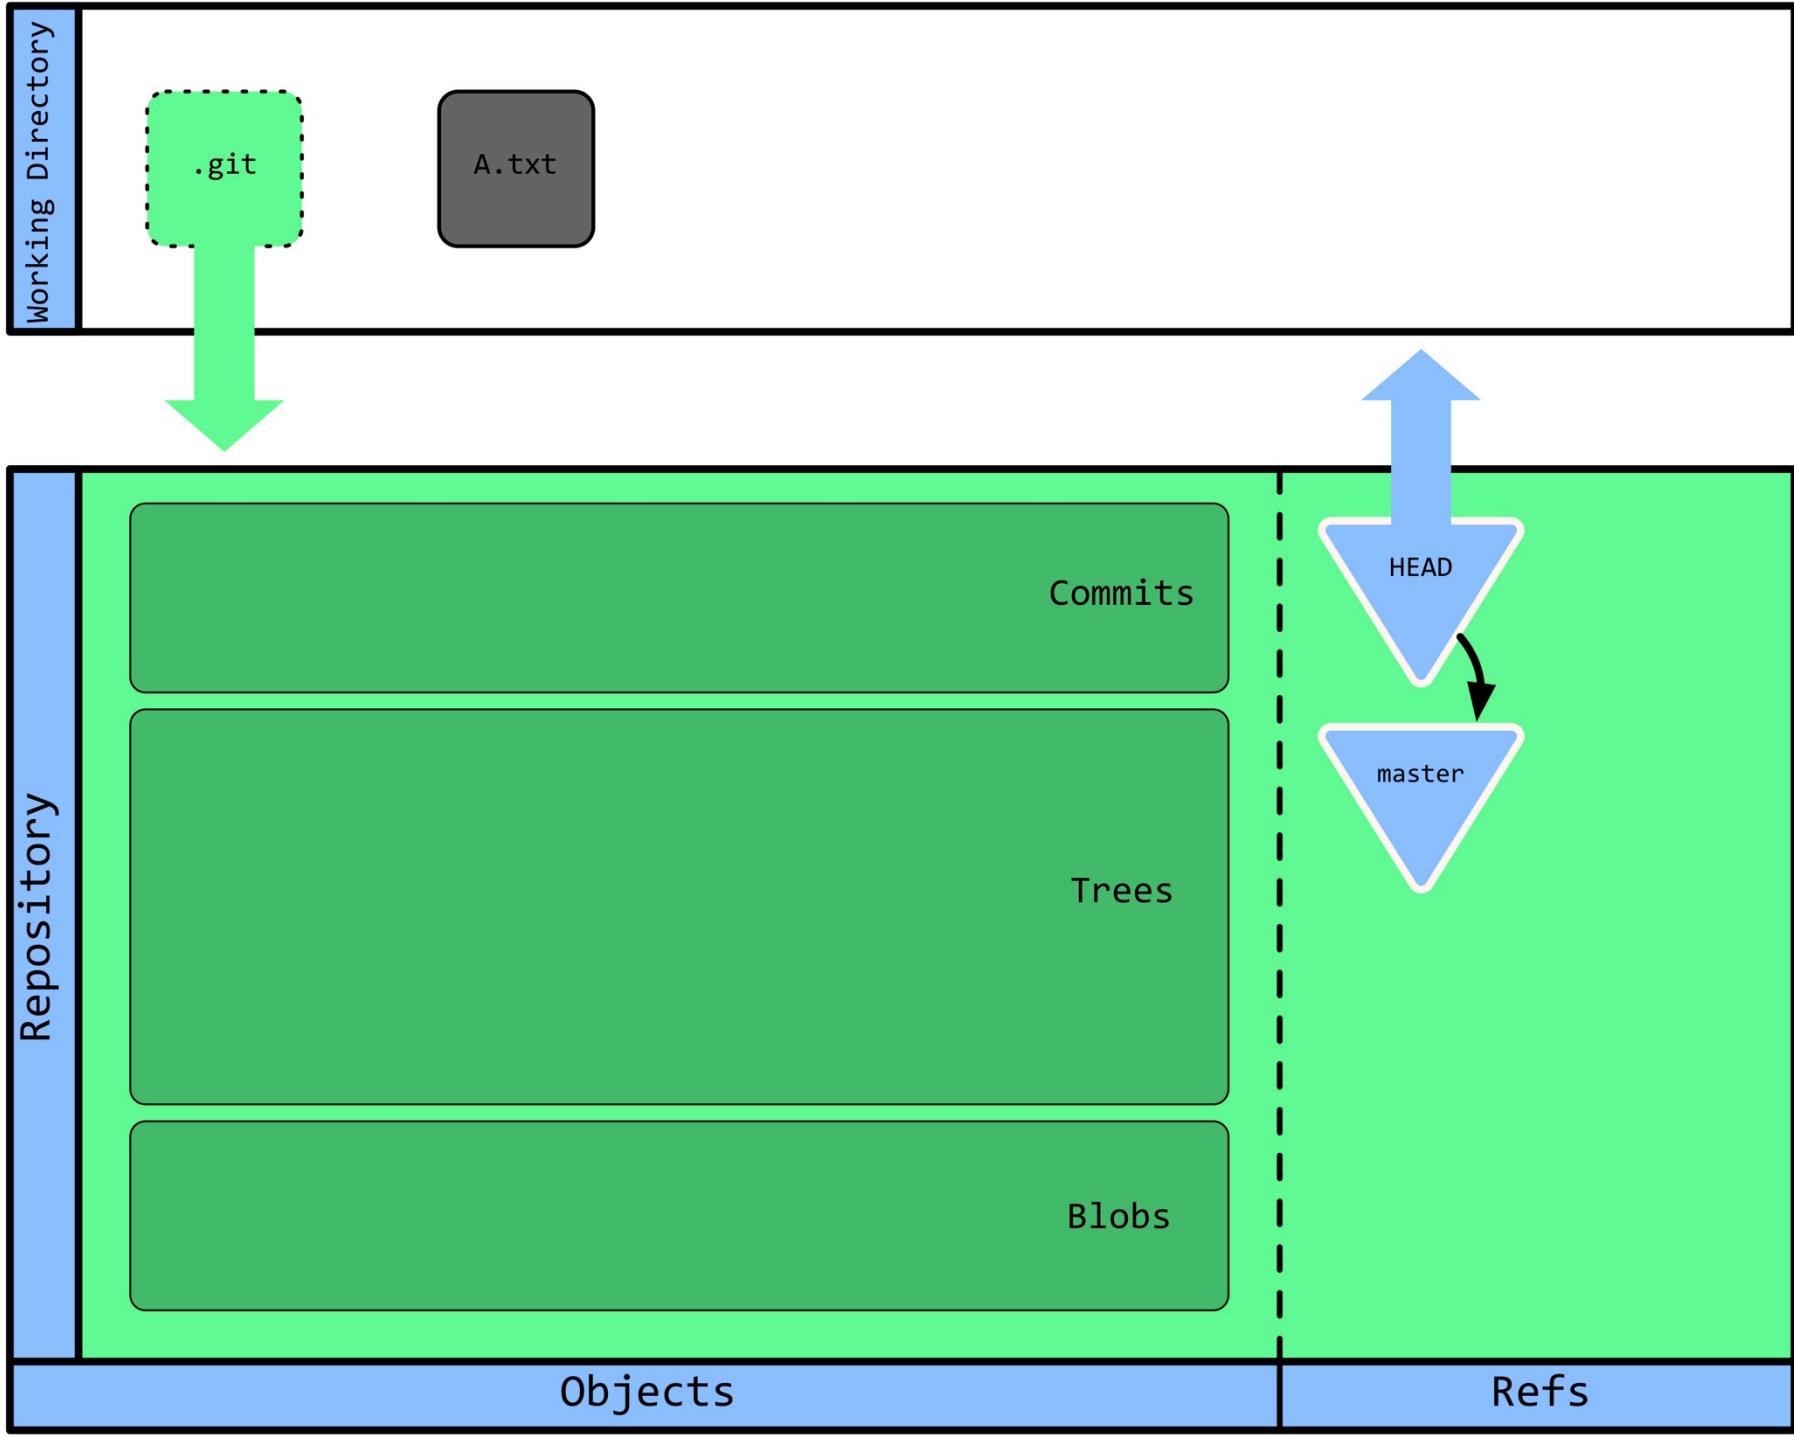
- Let's take a look at the following Git operations in gory detail:
 - `git add` a new file
 - `git commit` the new file
 - `git remote add` a remote called `origin`
 - `git push` to the remote
 - `git branch` to a new branch called `test`
 - `git checkout` the new branch
 - `git add` a new directory and new file
 - `git commit` the new directory and file
 - `git merge` into `master`
 - `git tag` the commit
 - `git push` the changes to the remote
 - `git fetch` for new changes
 - `git merge` with the remote branch to update ours
 - `git checkout` a previous commit

We'll start with an empty repo

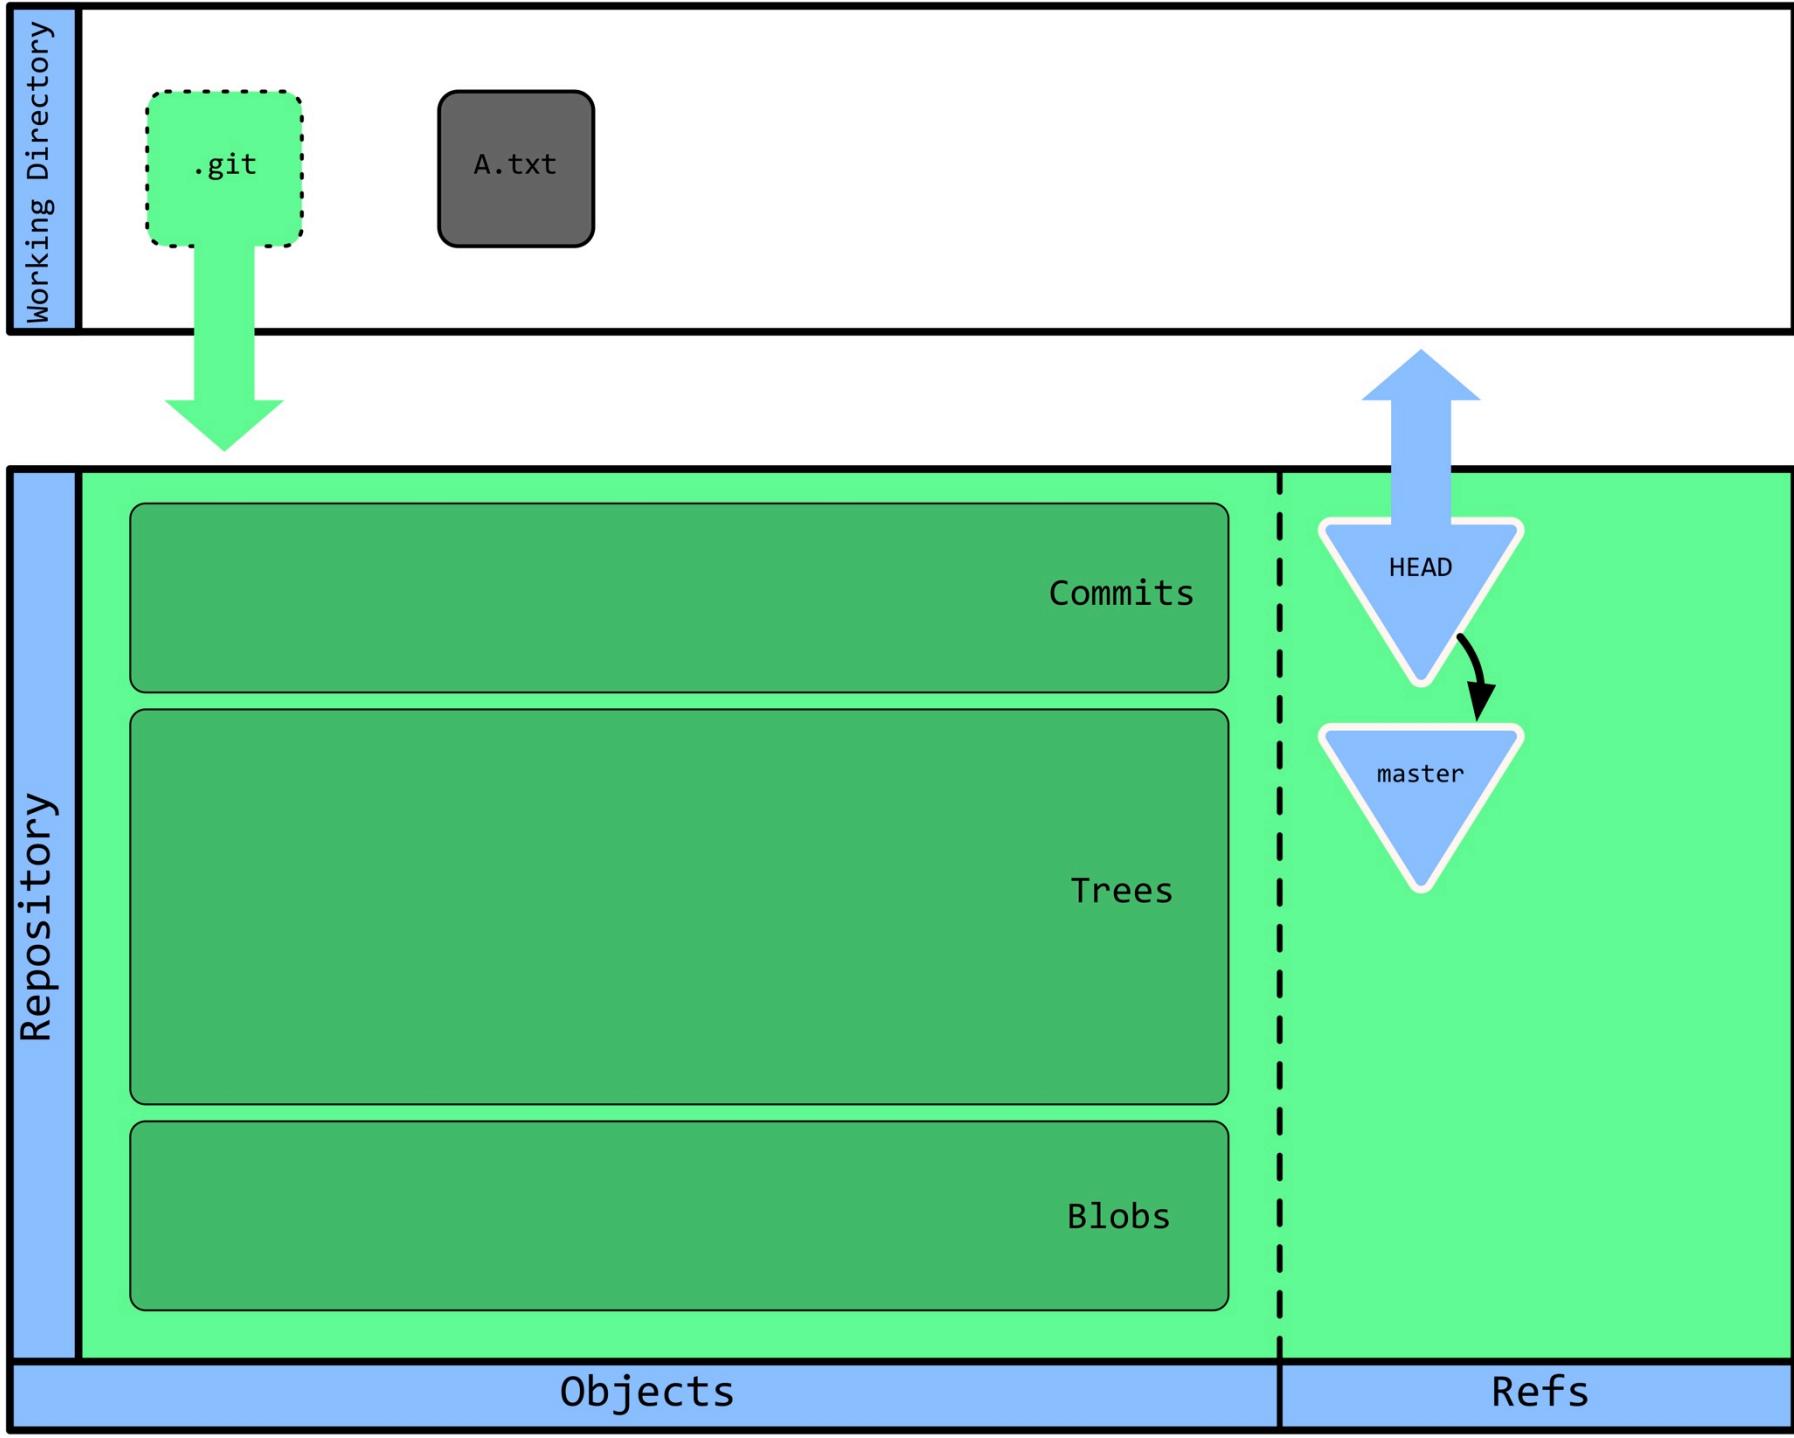


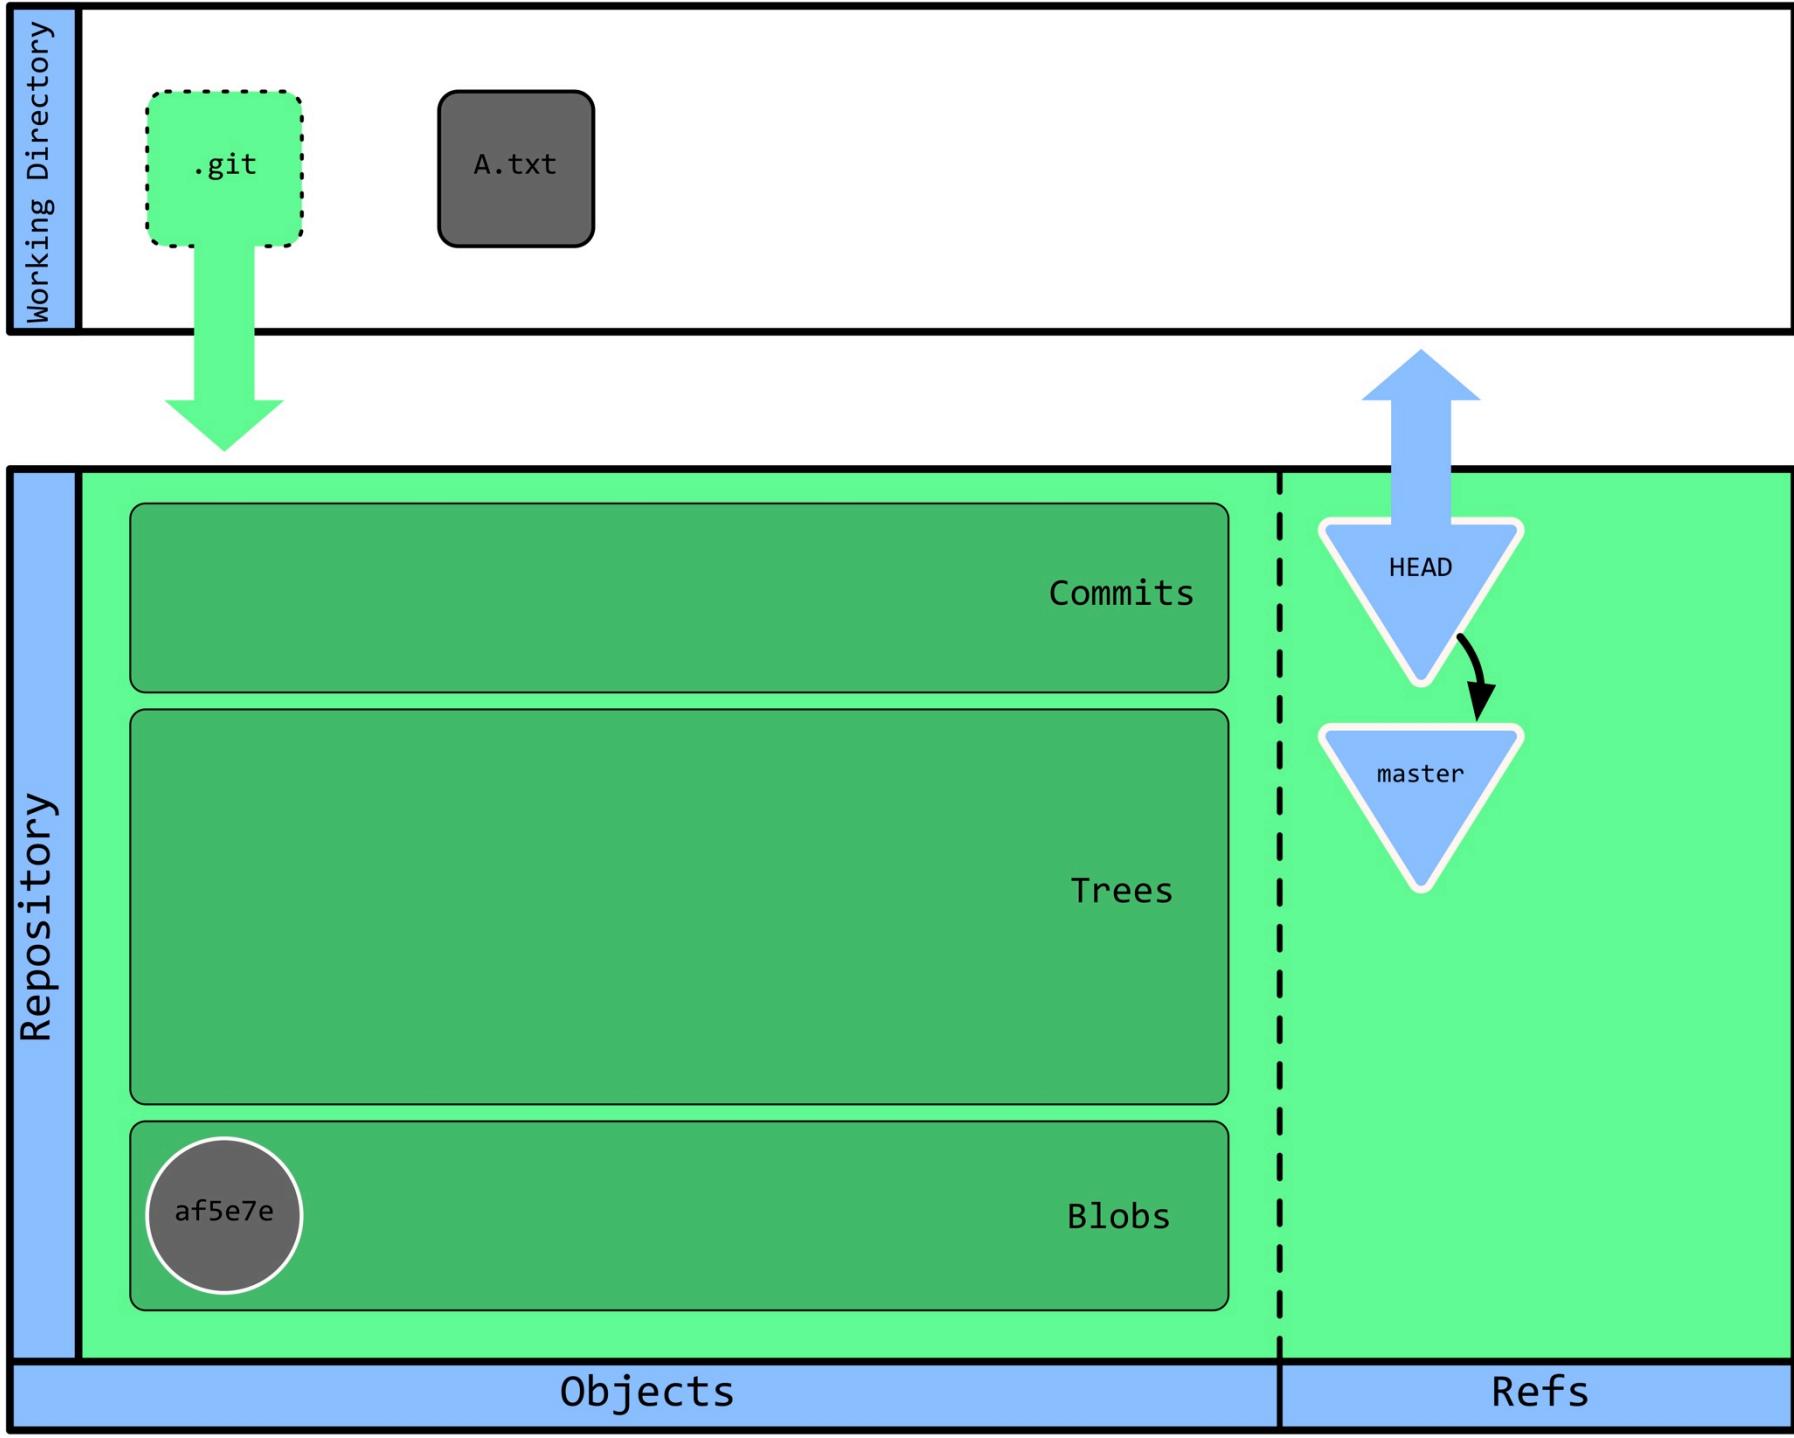
We create a file: A.txt



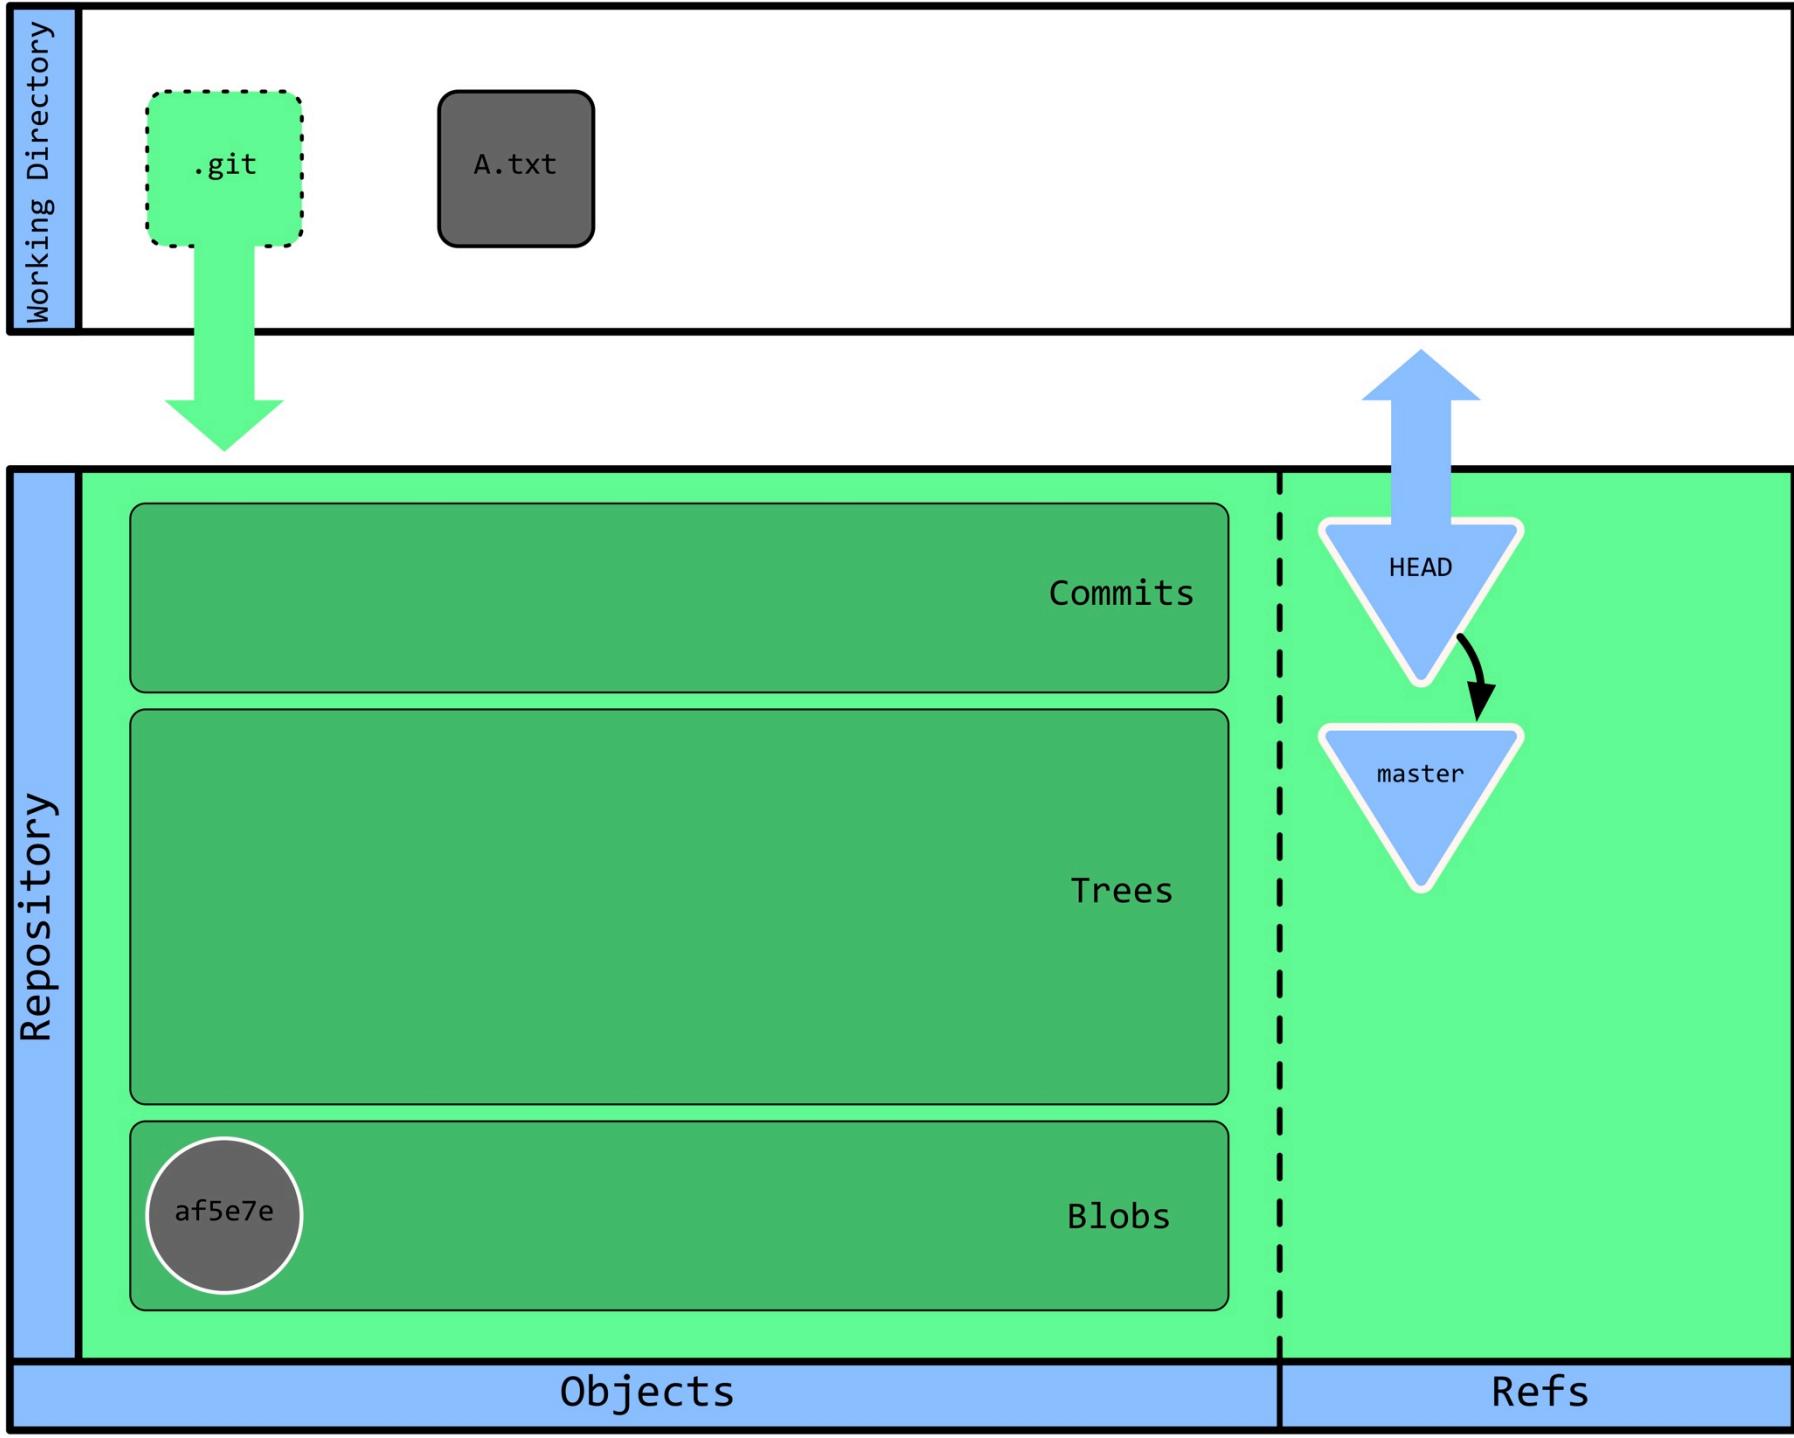


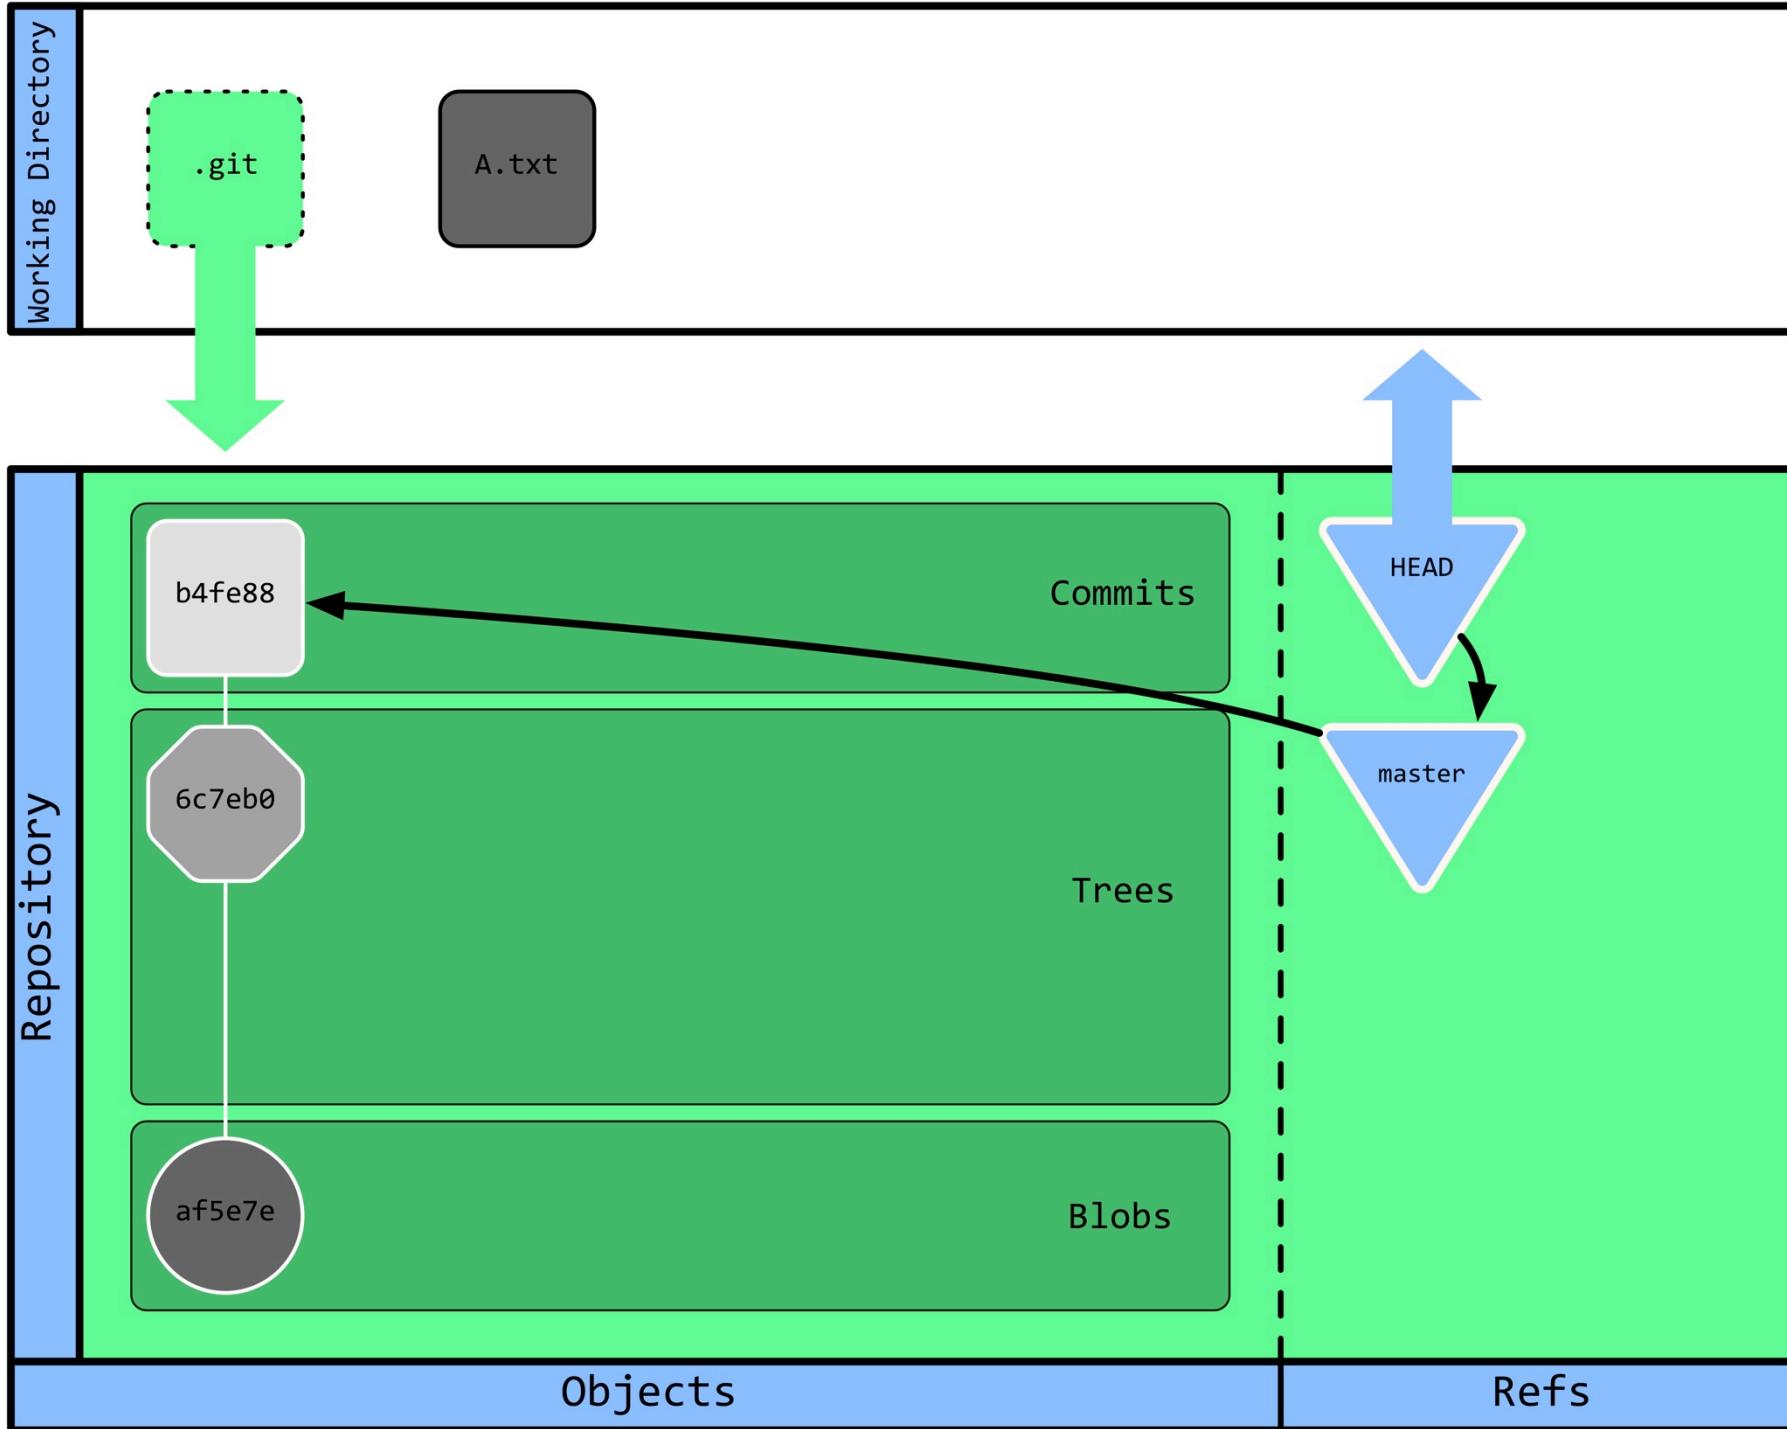
git add A.txt





```
git commit -m “...”
```

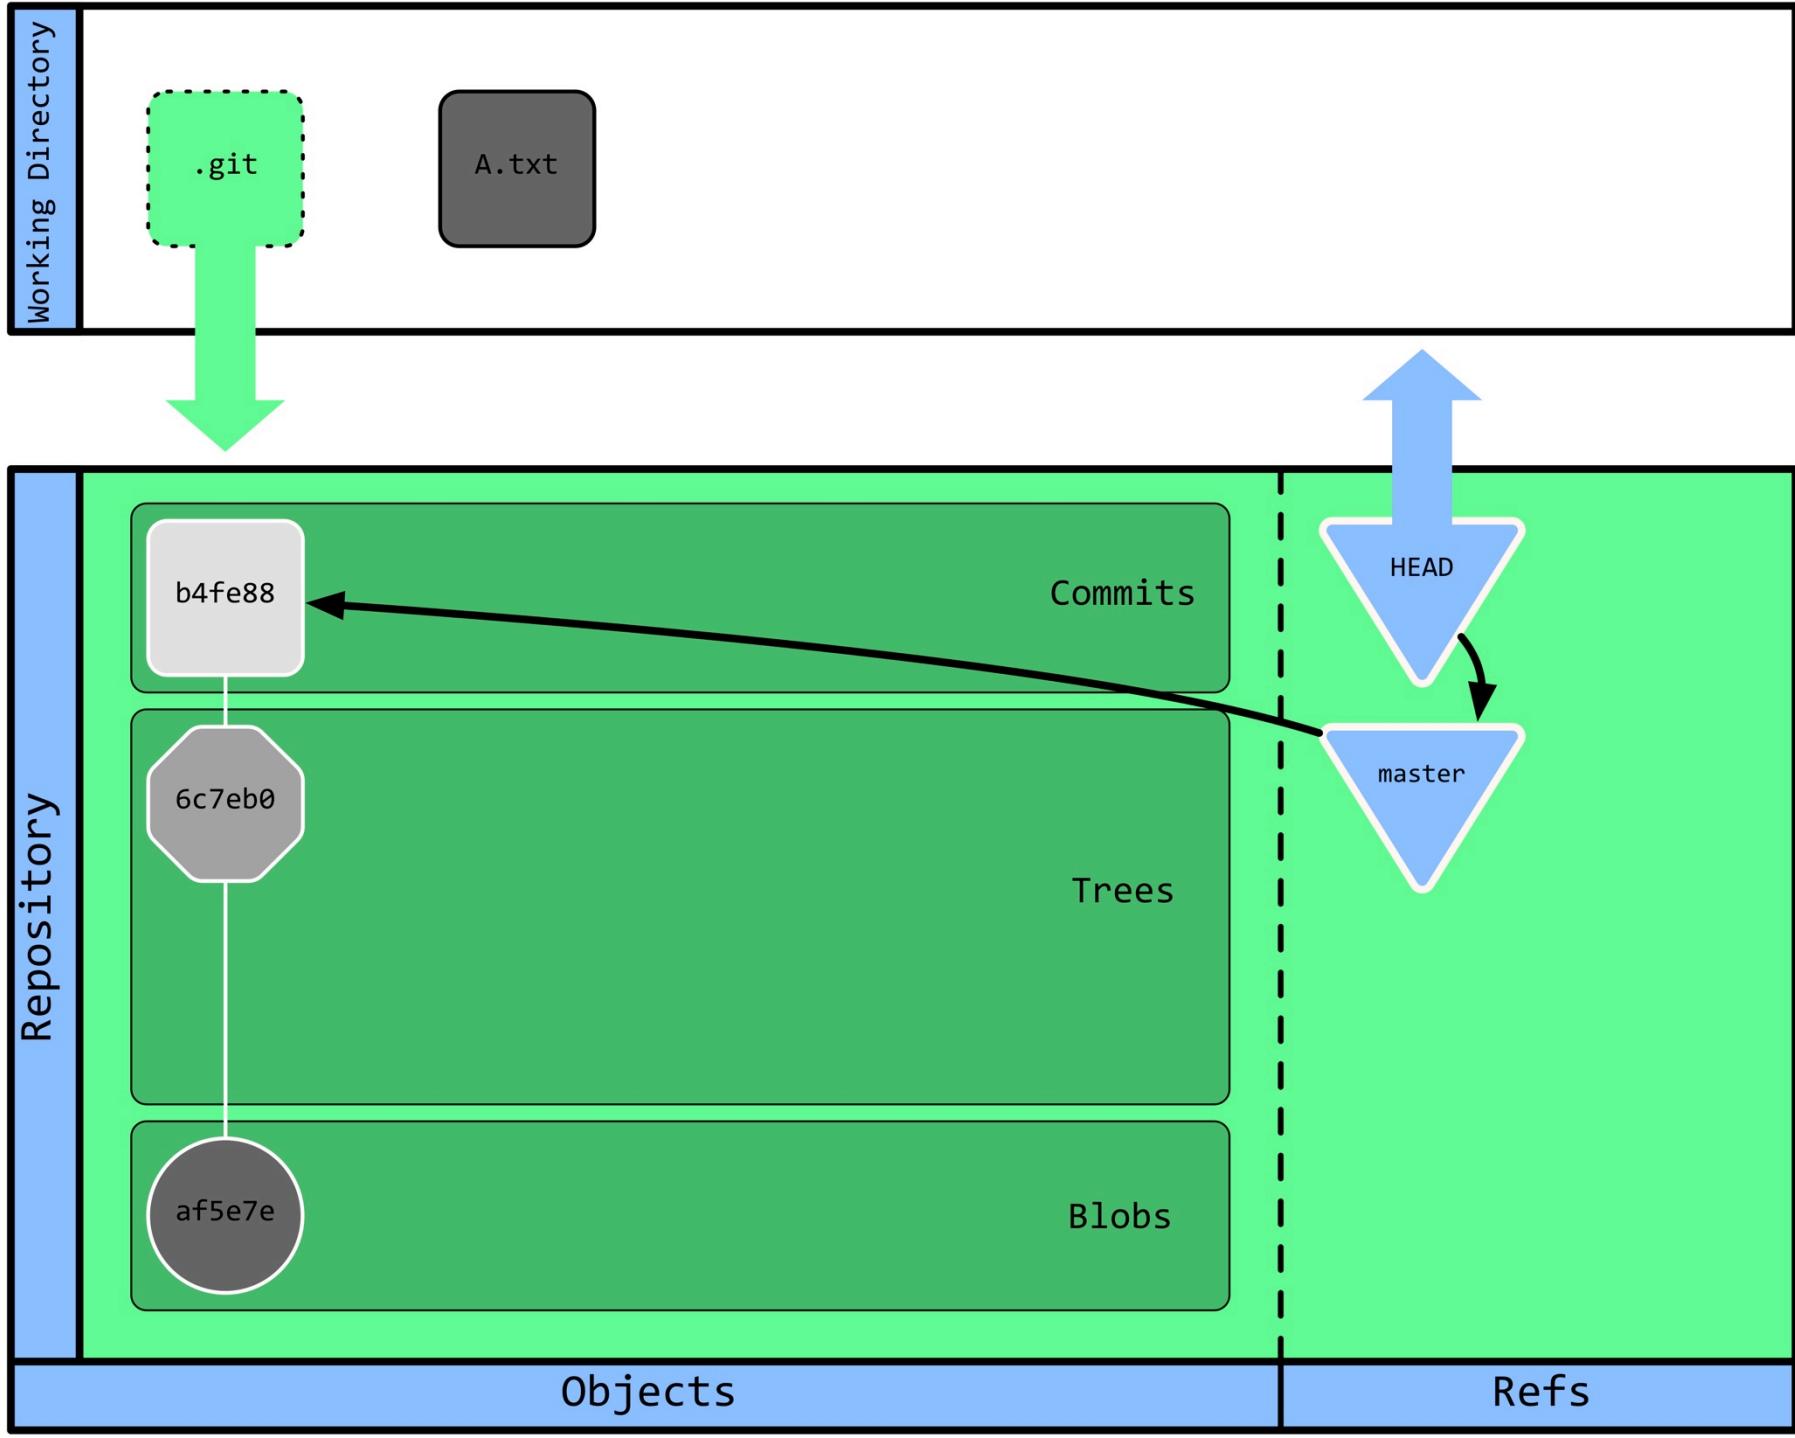


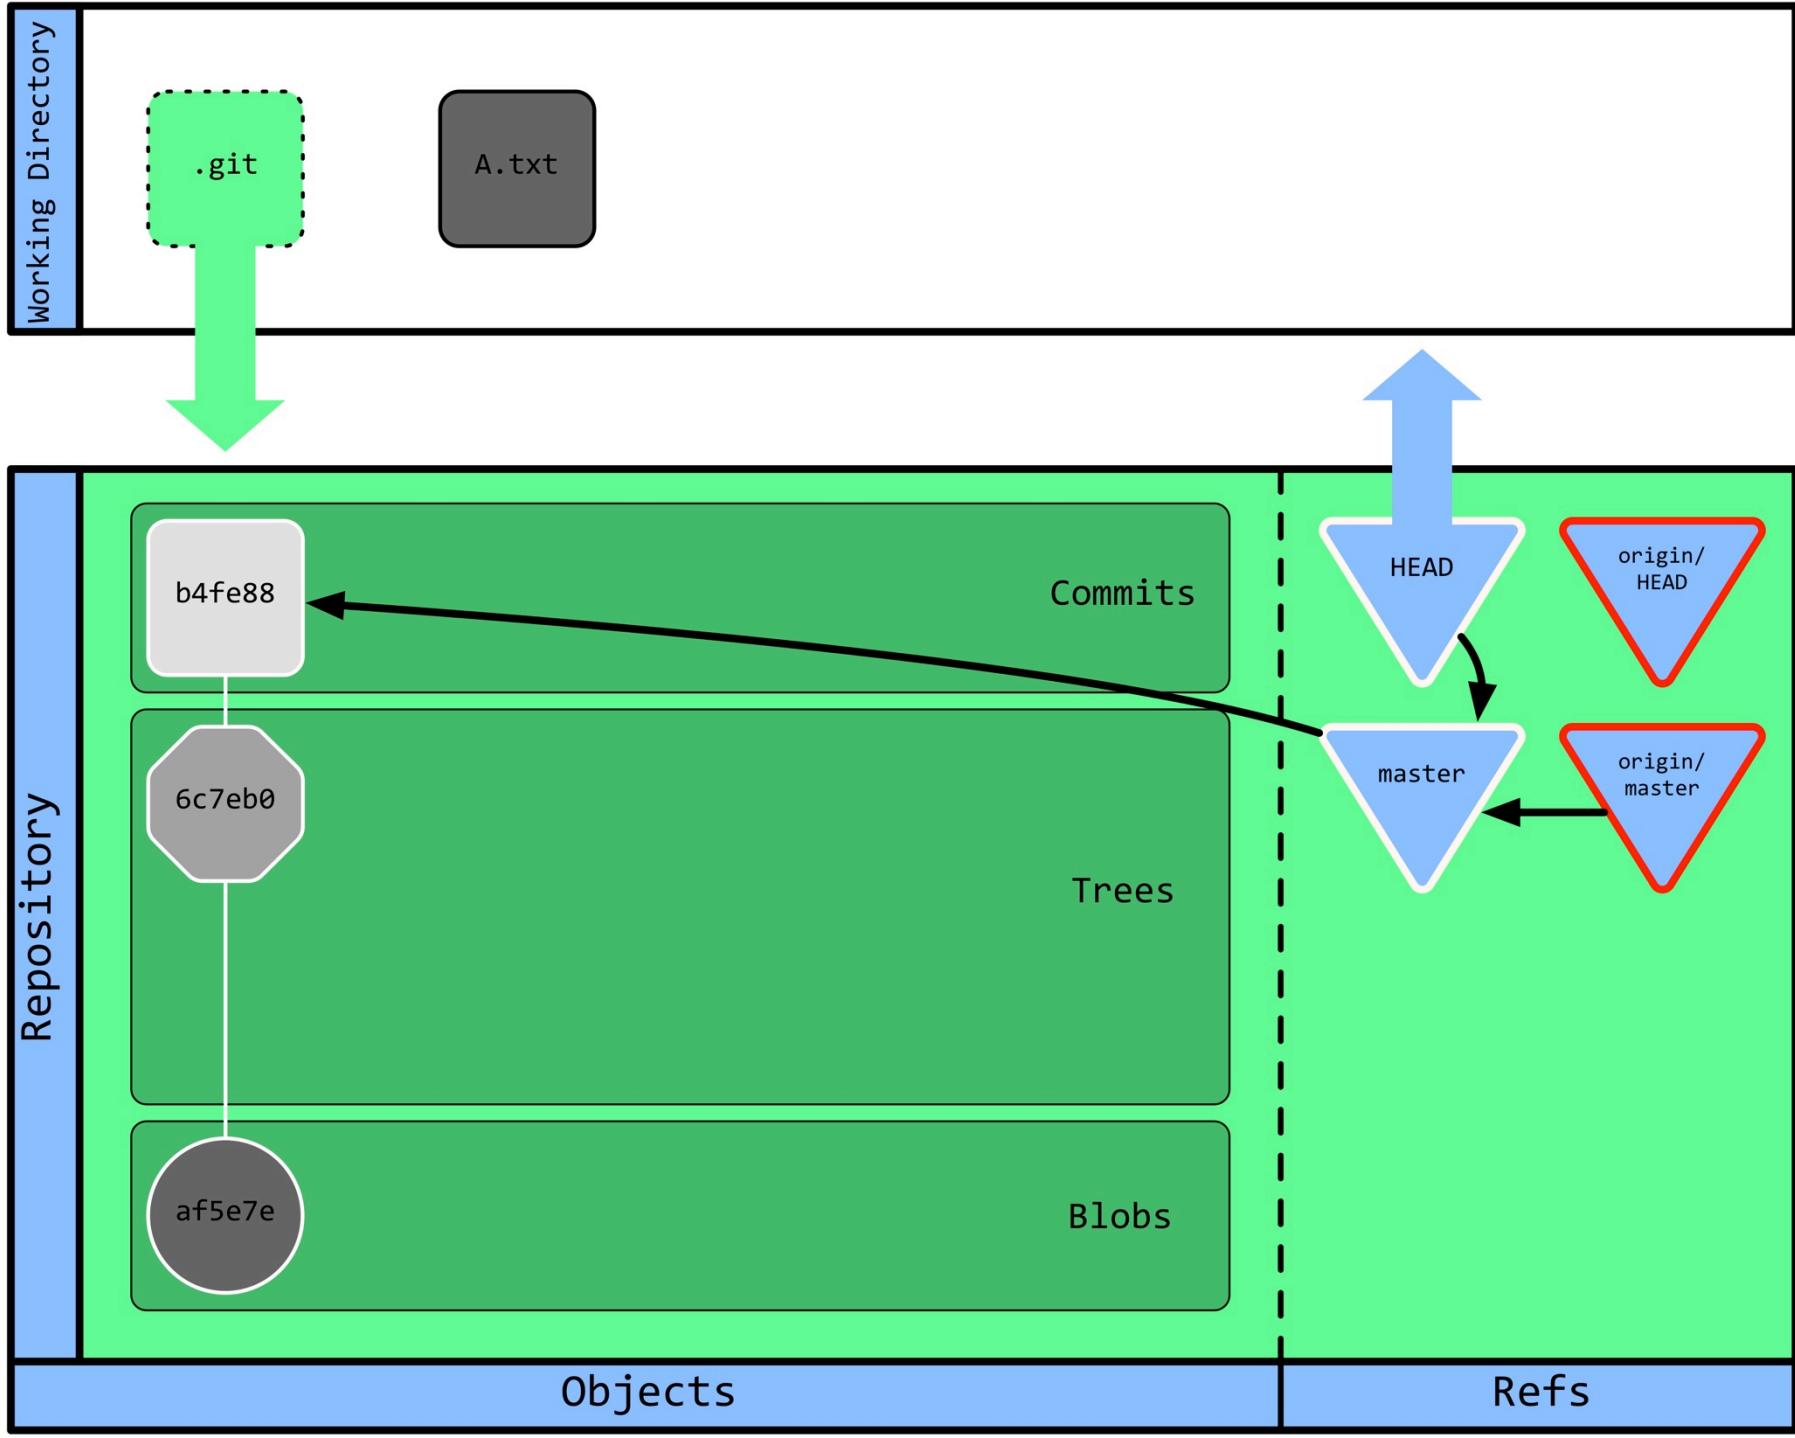


Let's add a remote: `origin`

git remote add origin

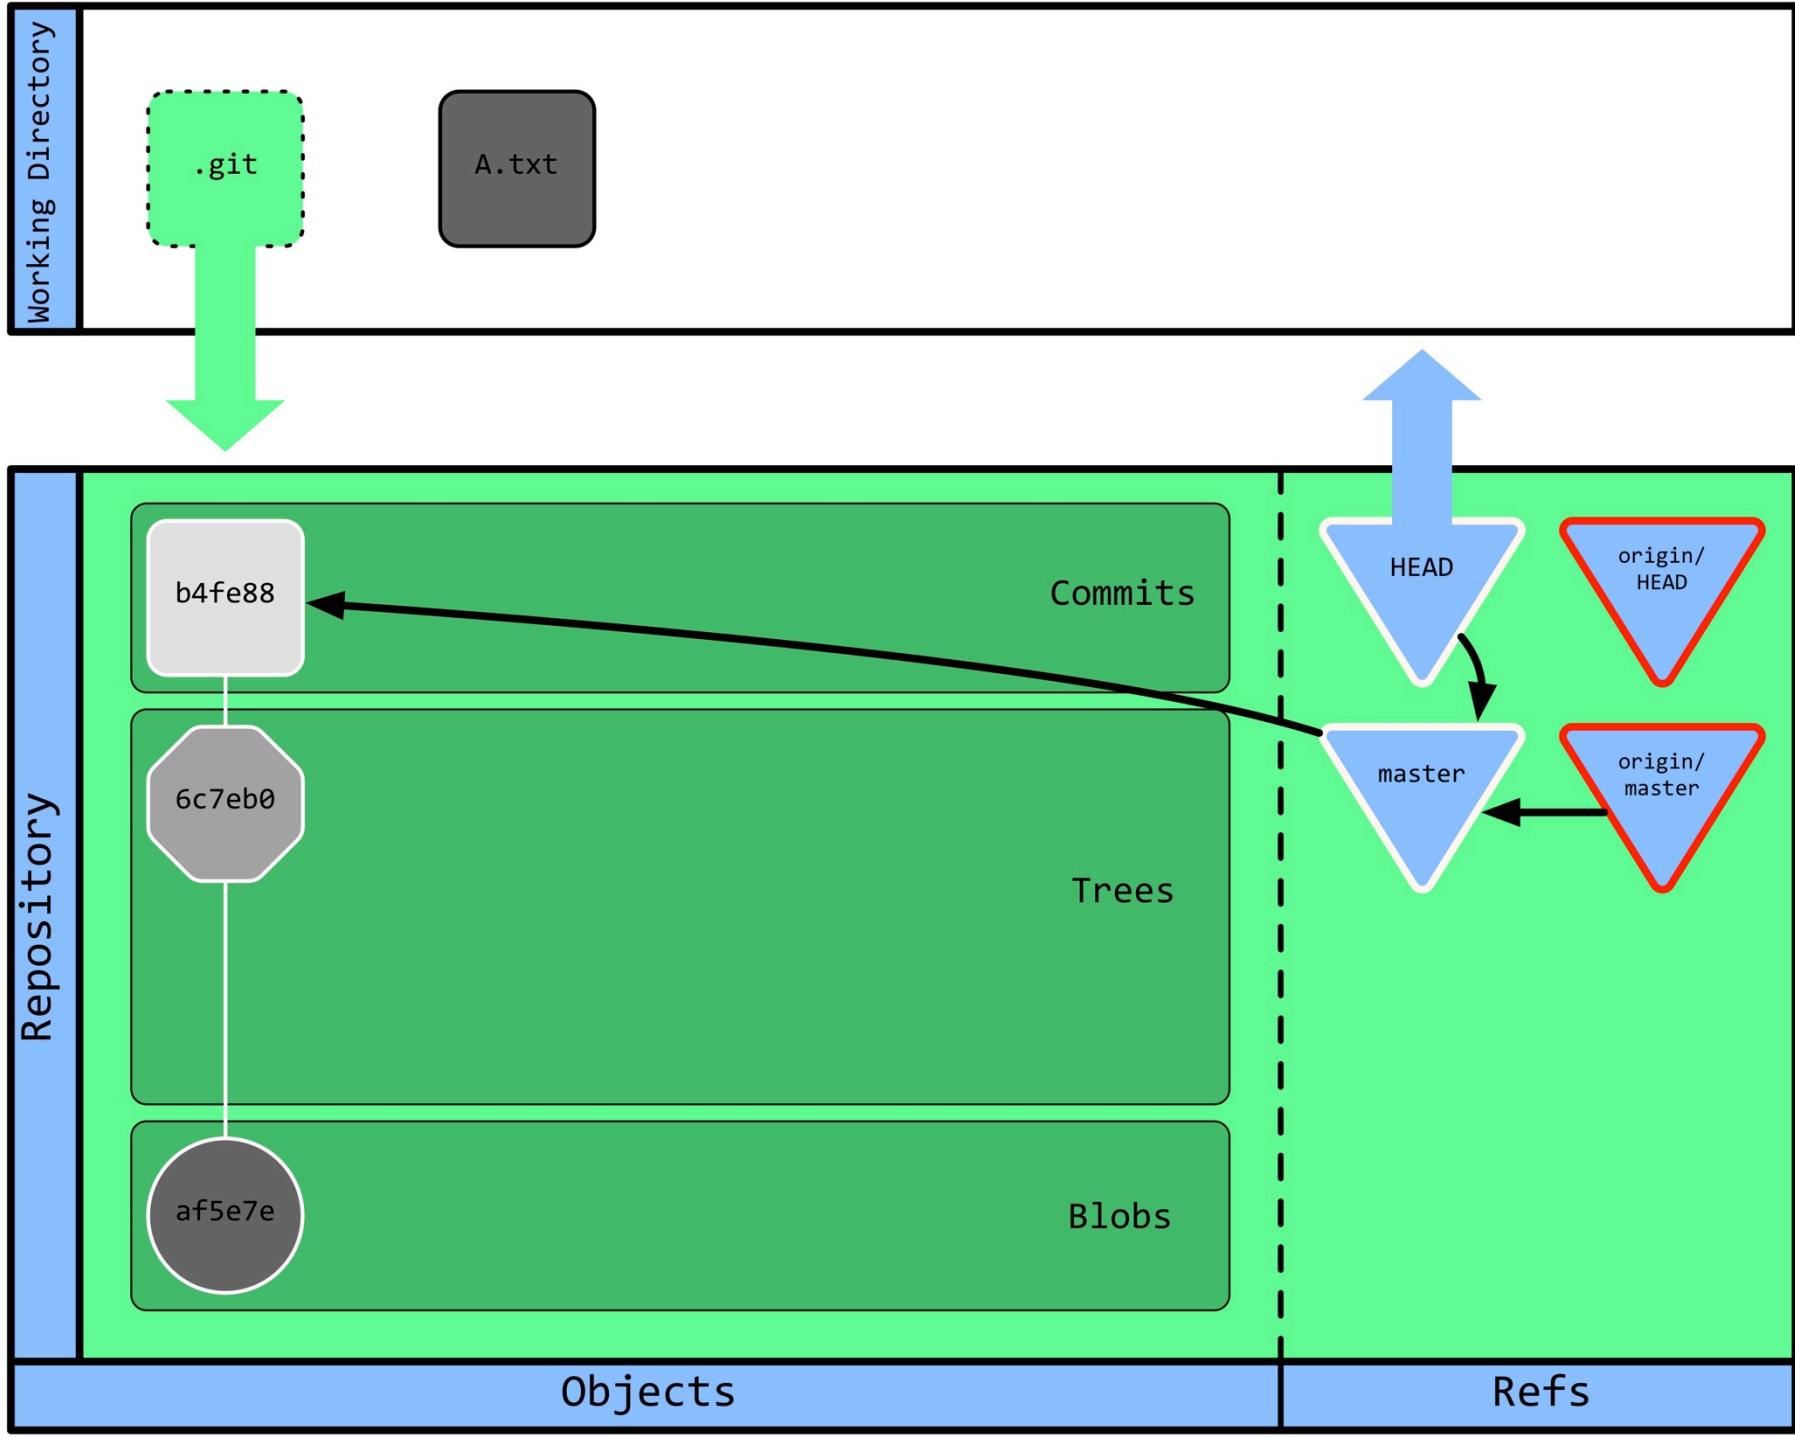
git push origin master

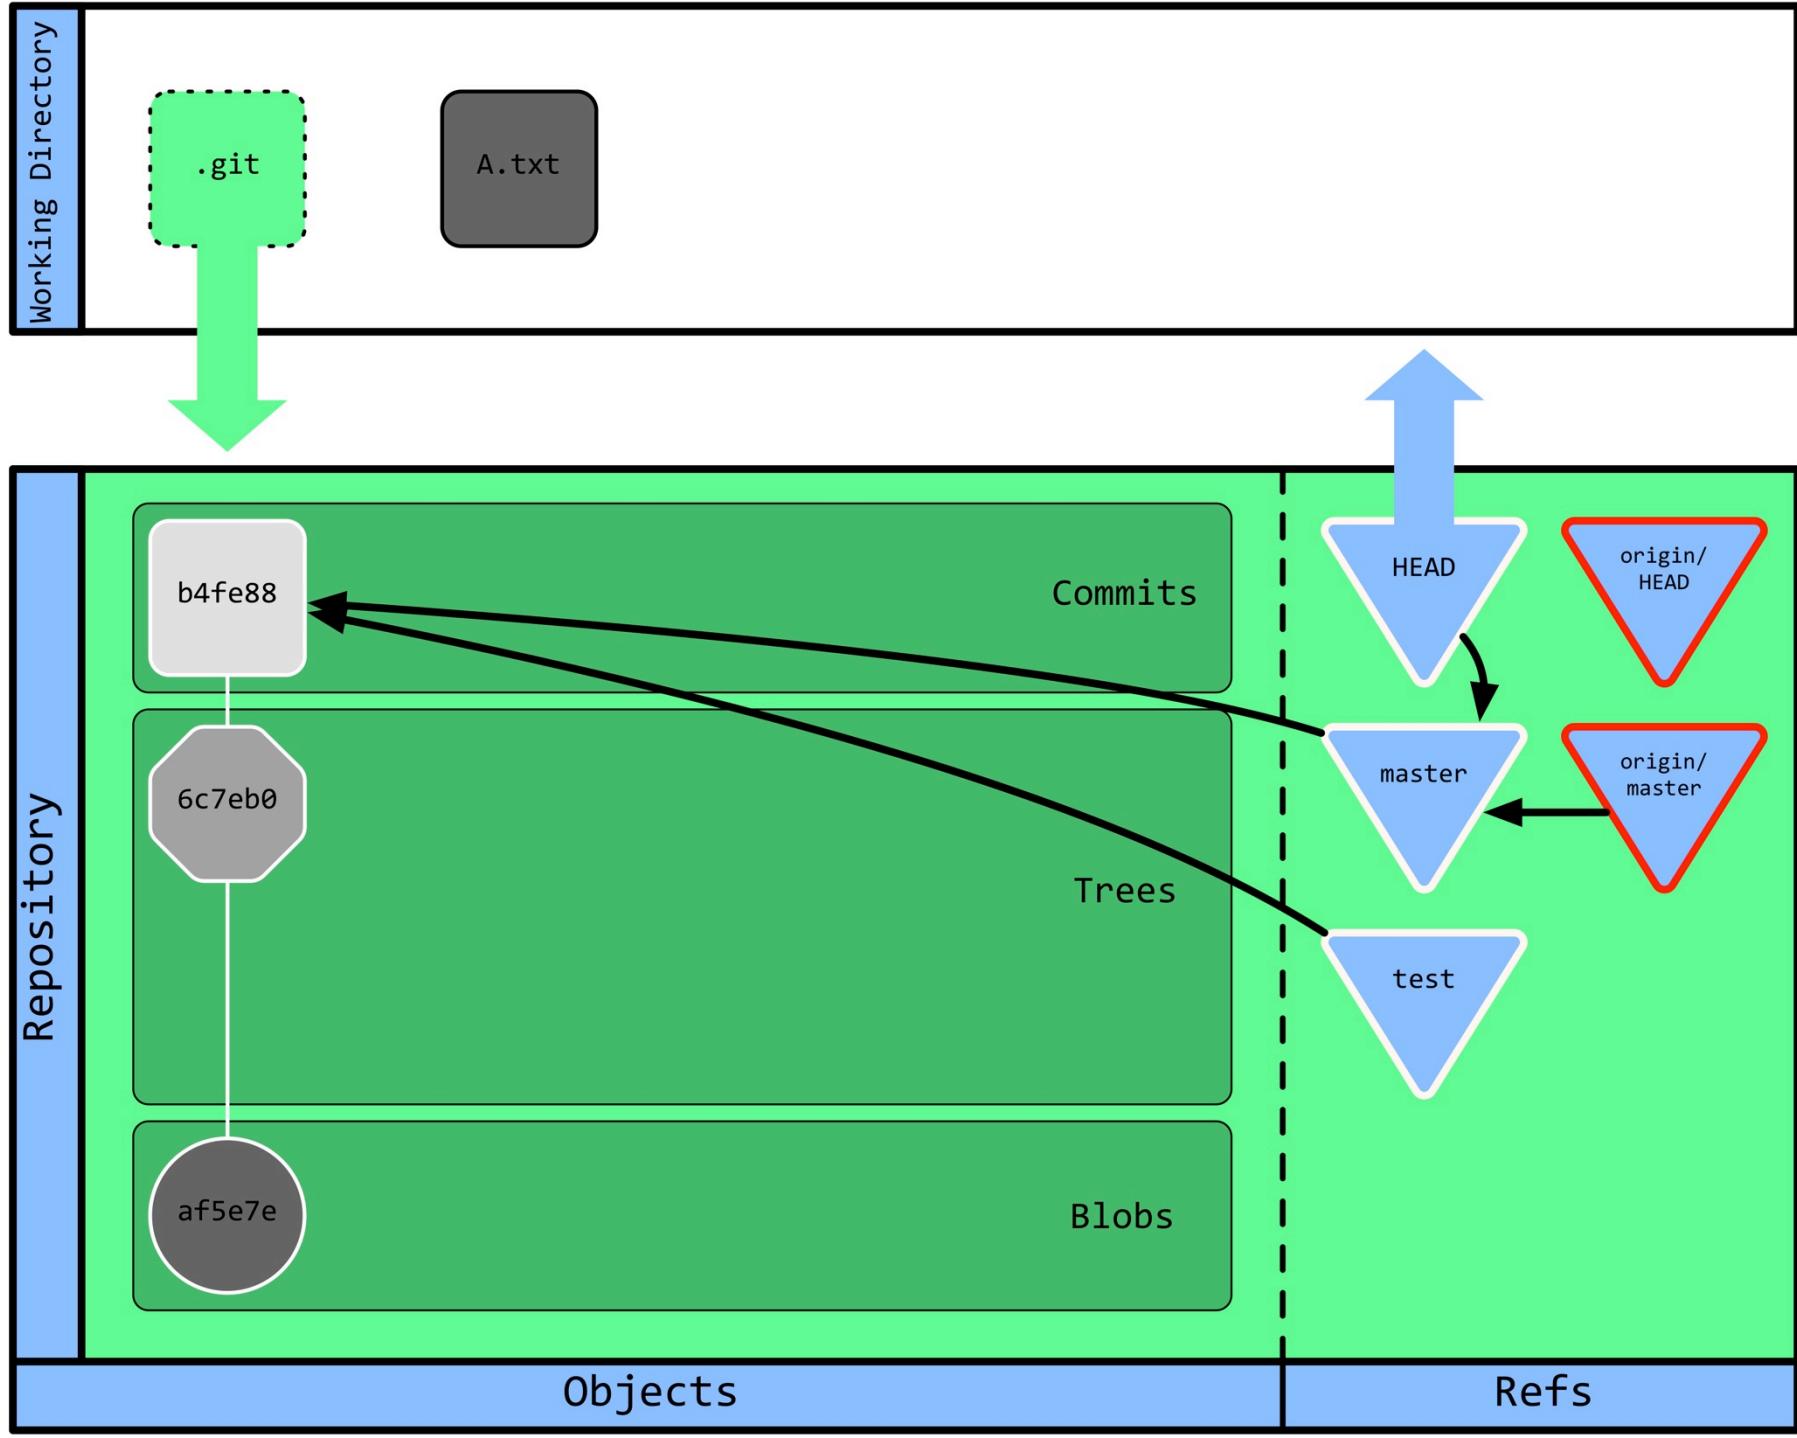




Let's add a branch: test

git branch test

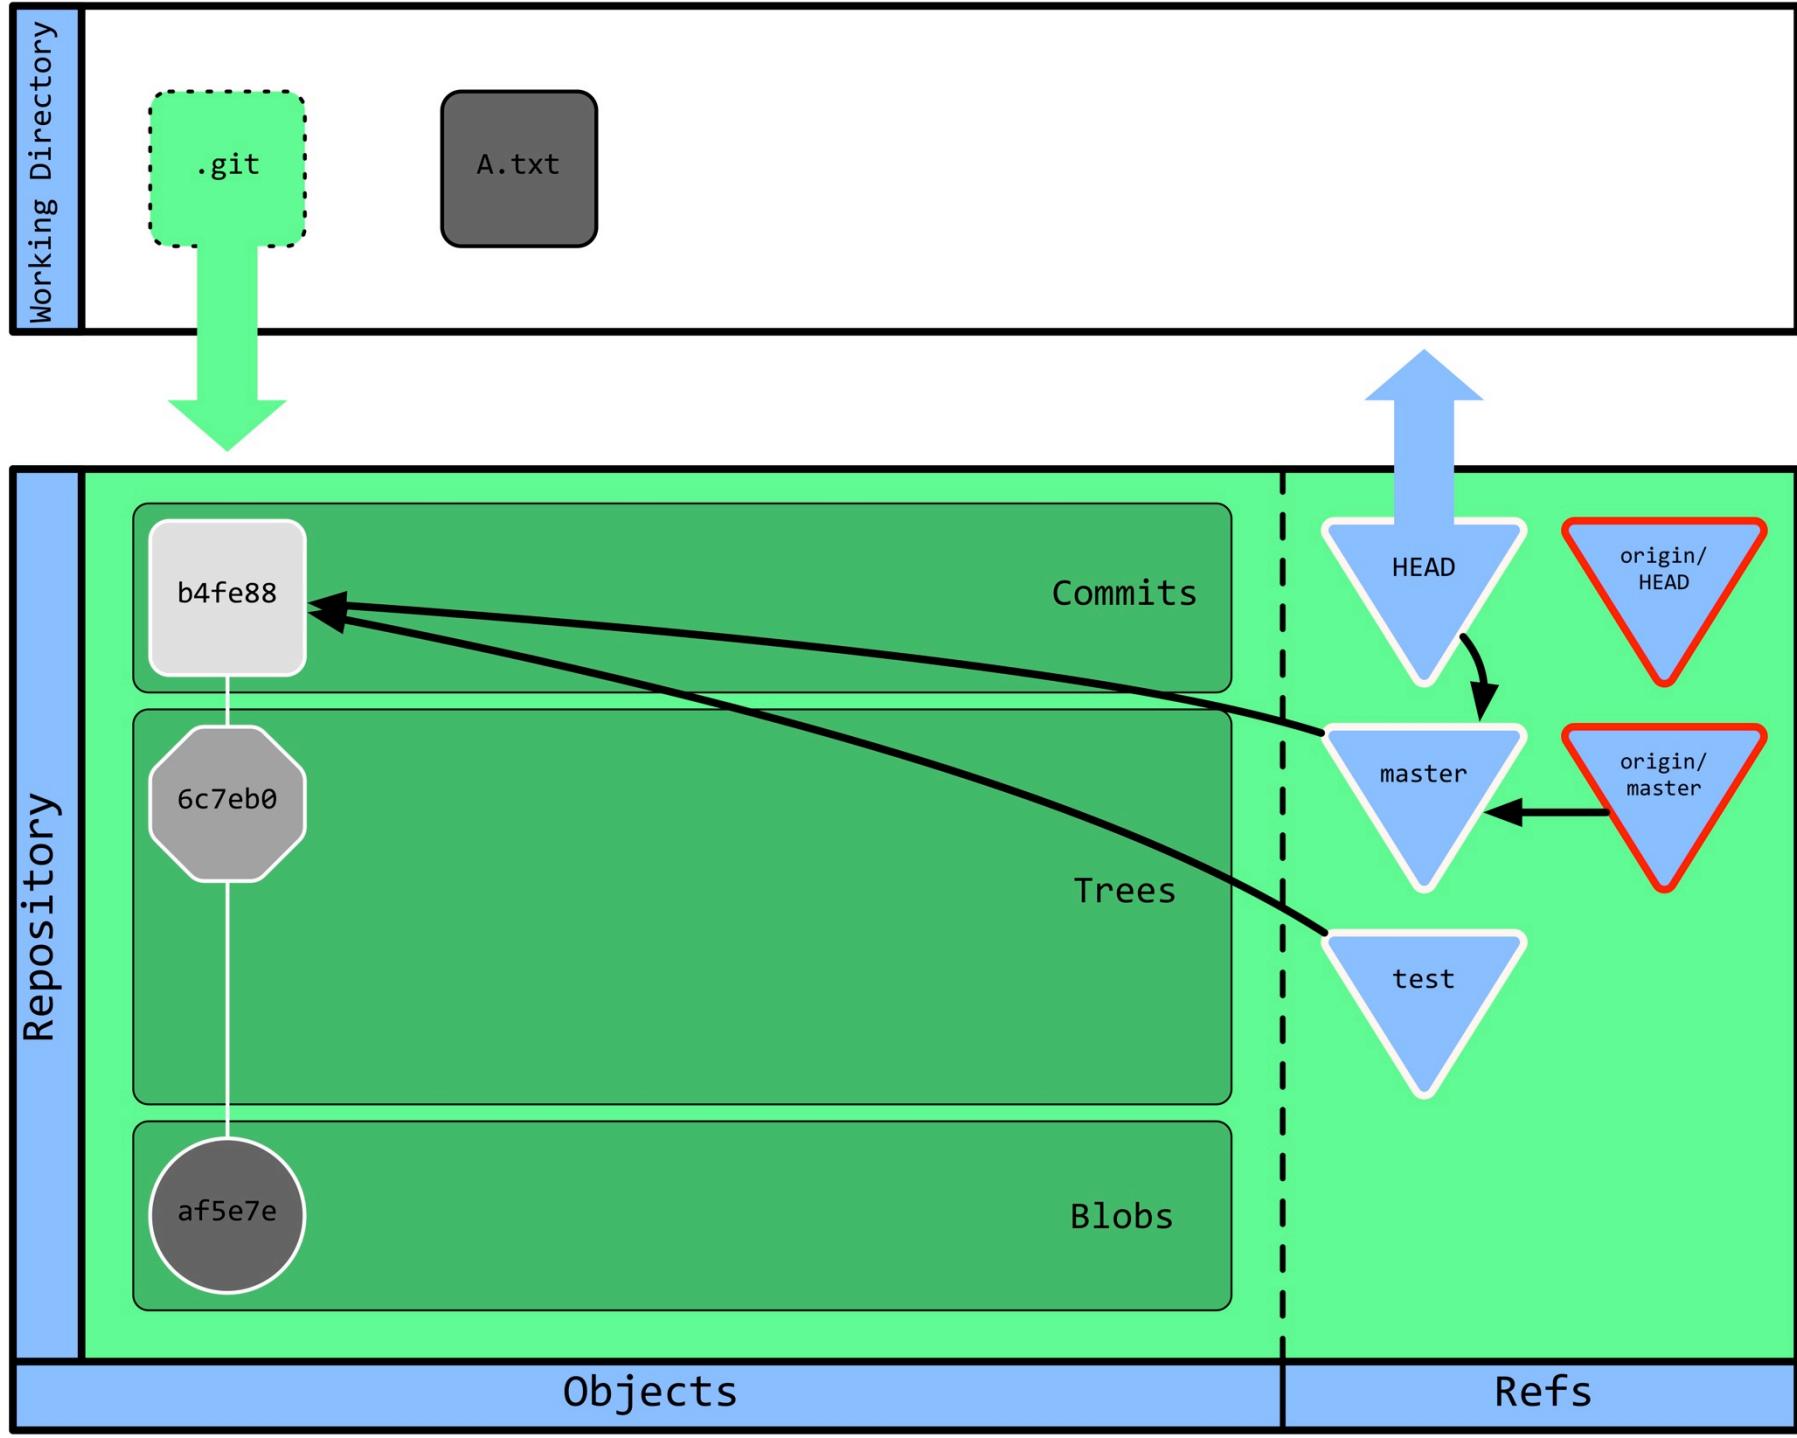


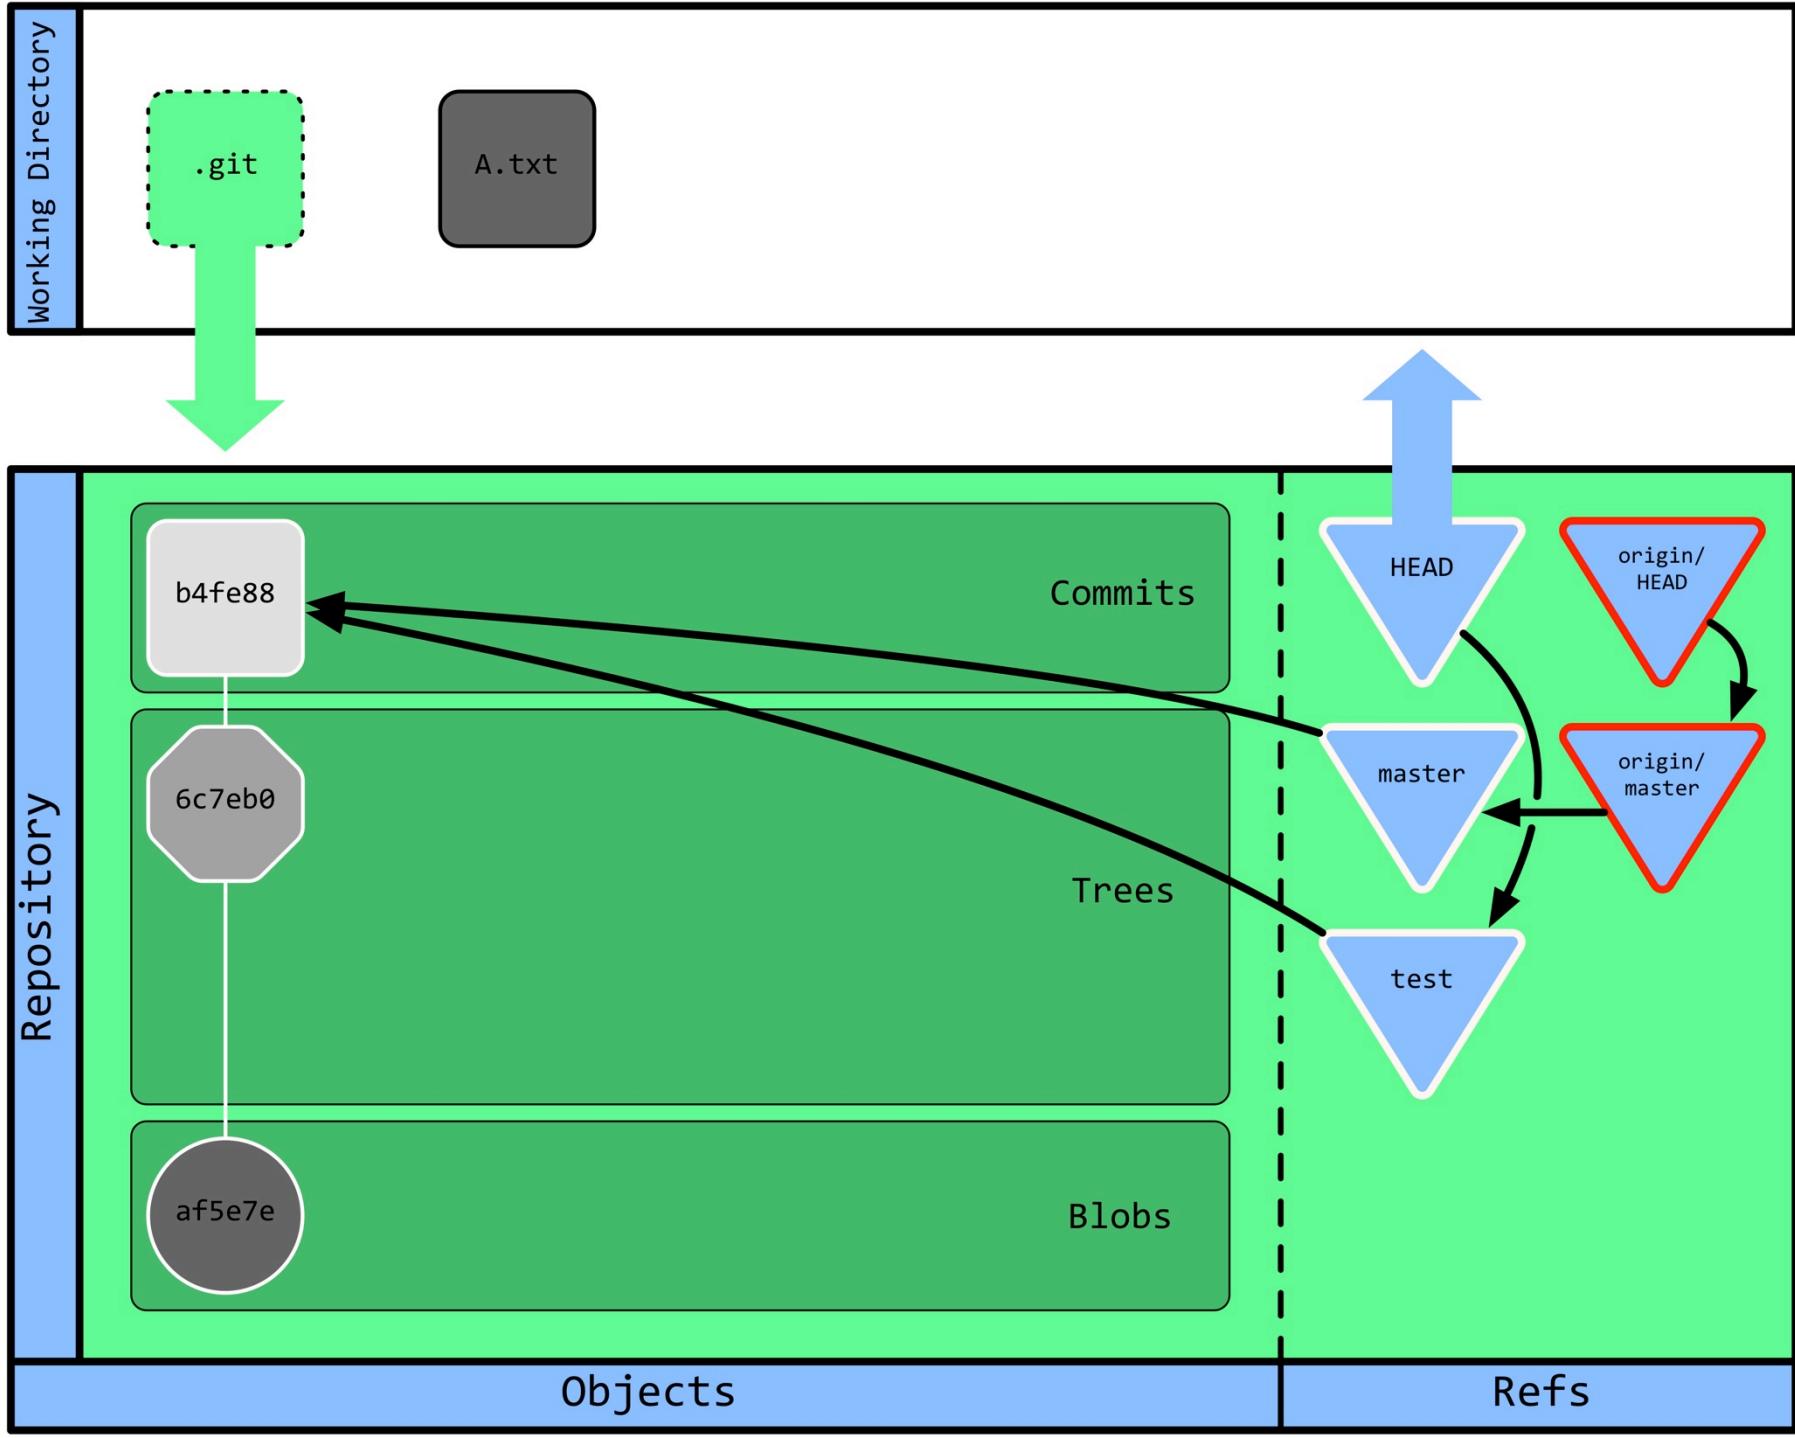


That's seriously what a branch is in Git

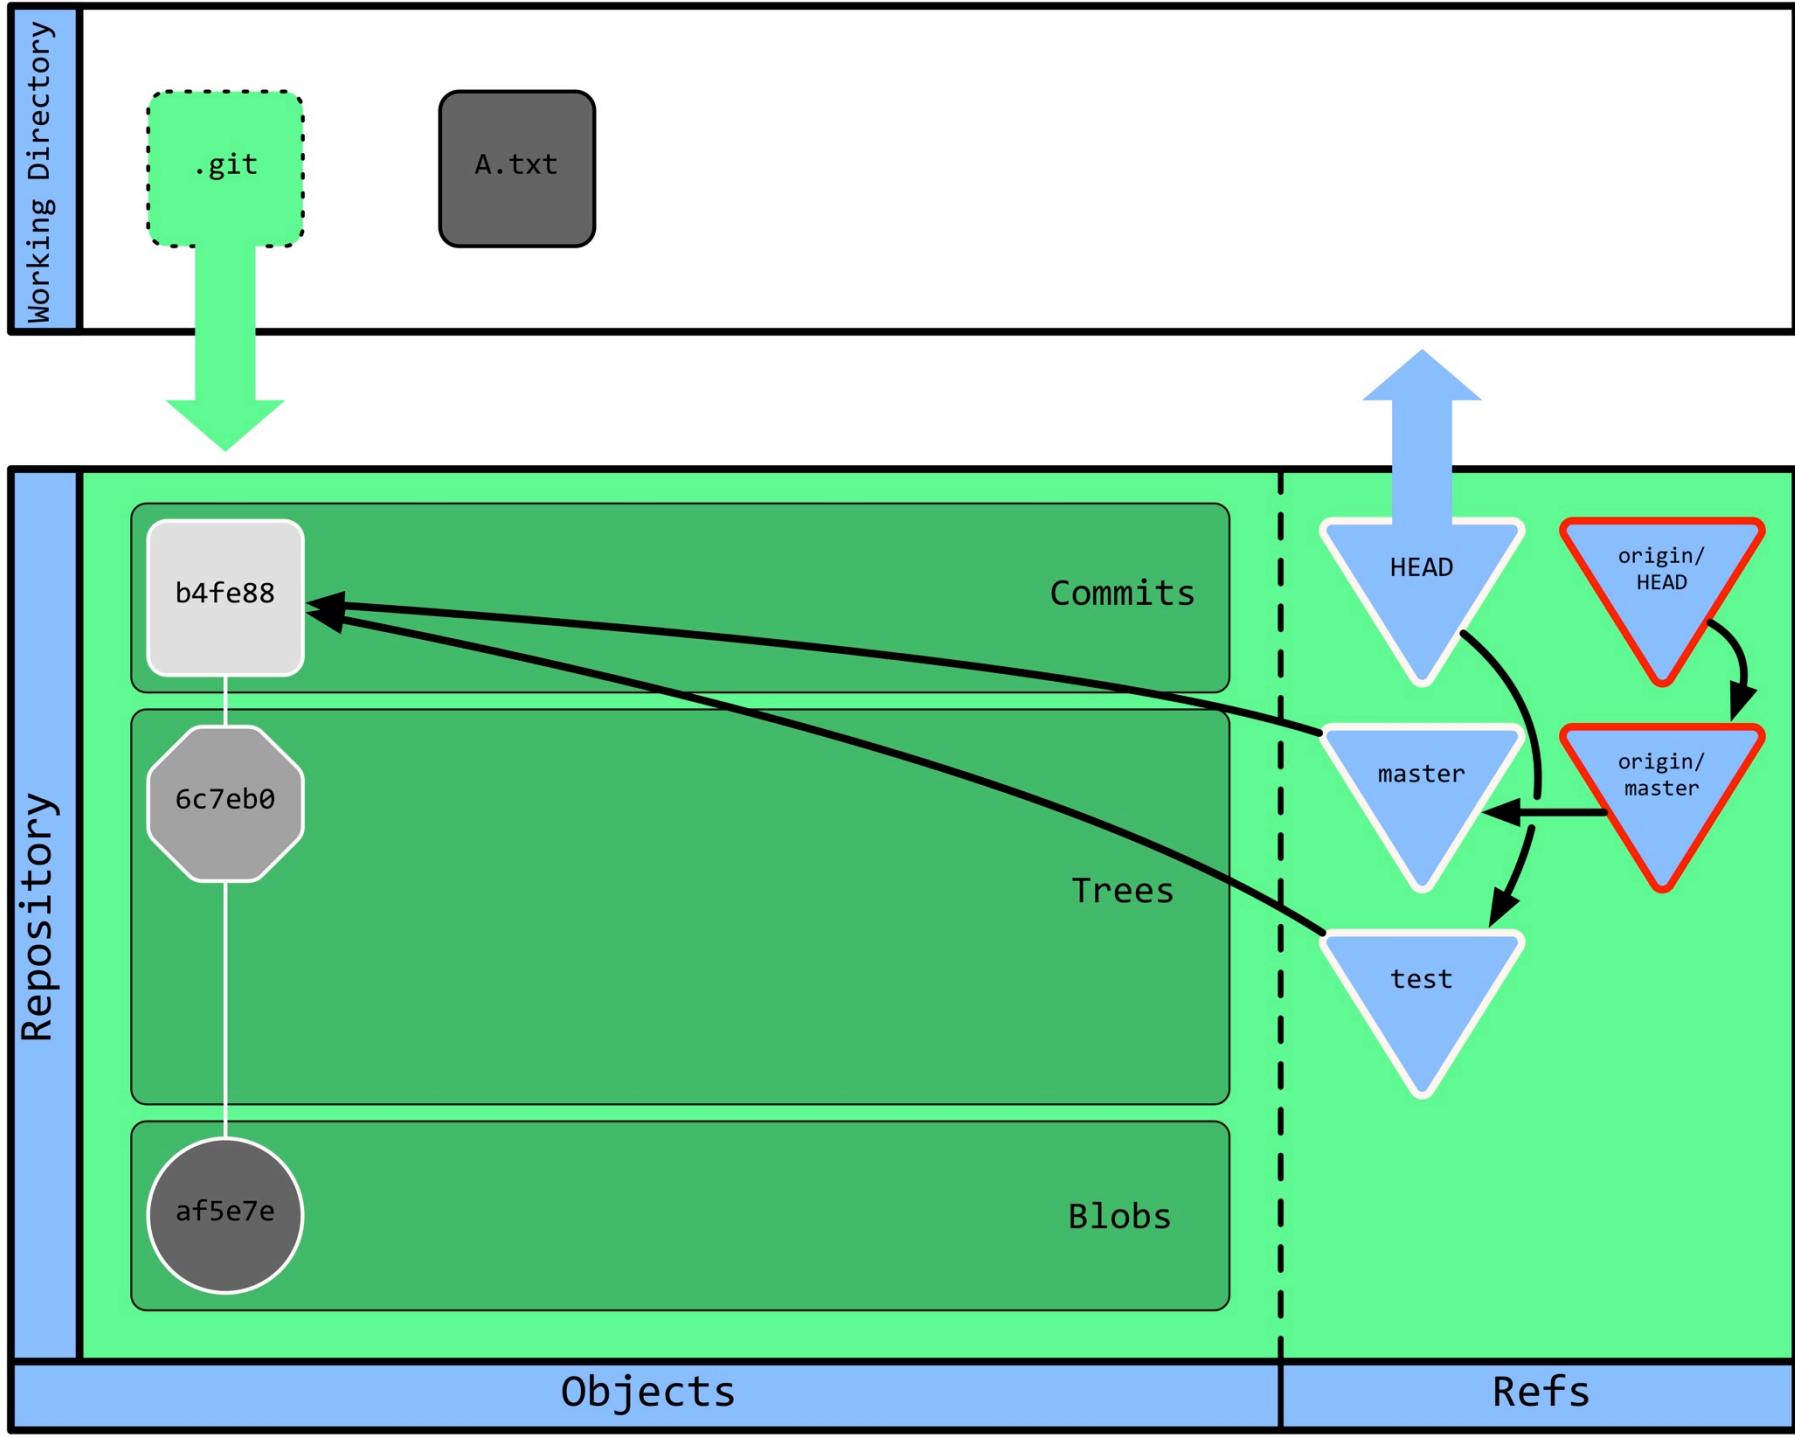
We can checkout the test branch

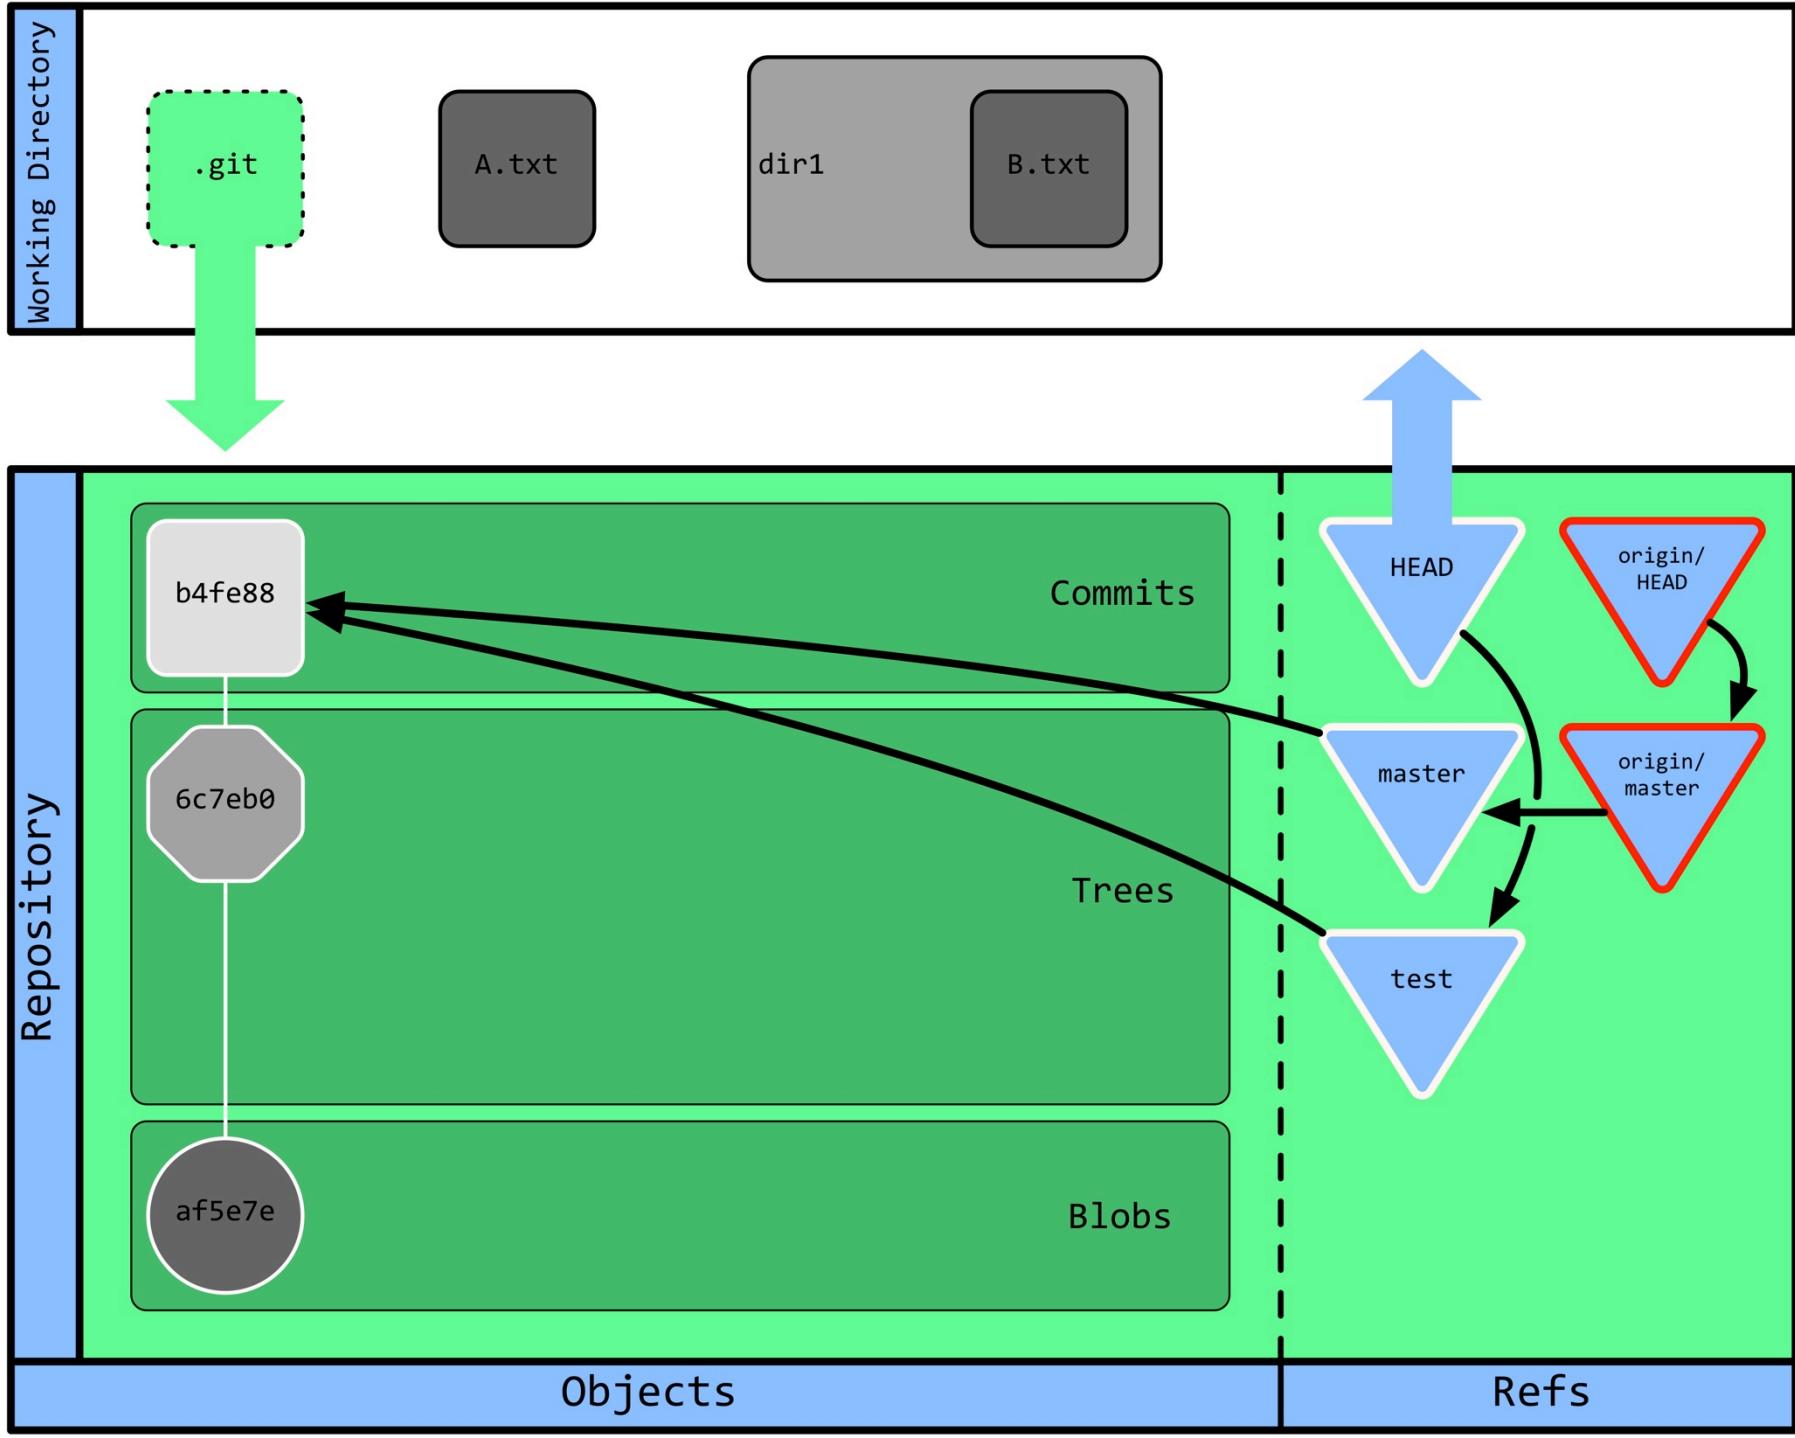
git checkout test



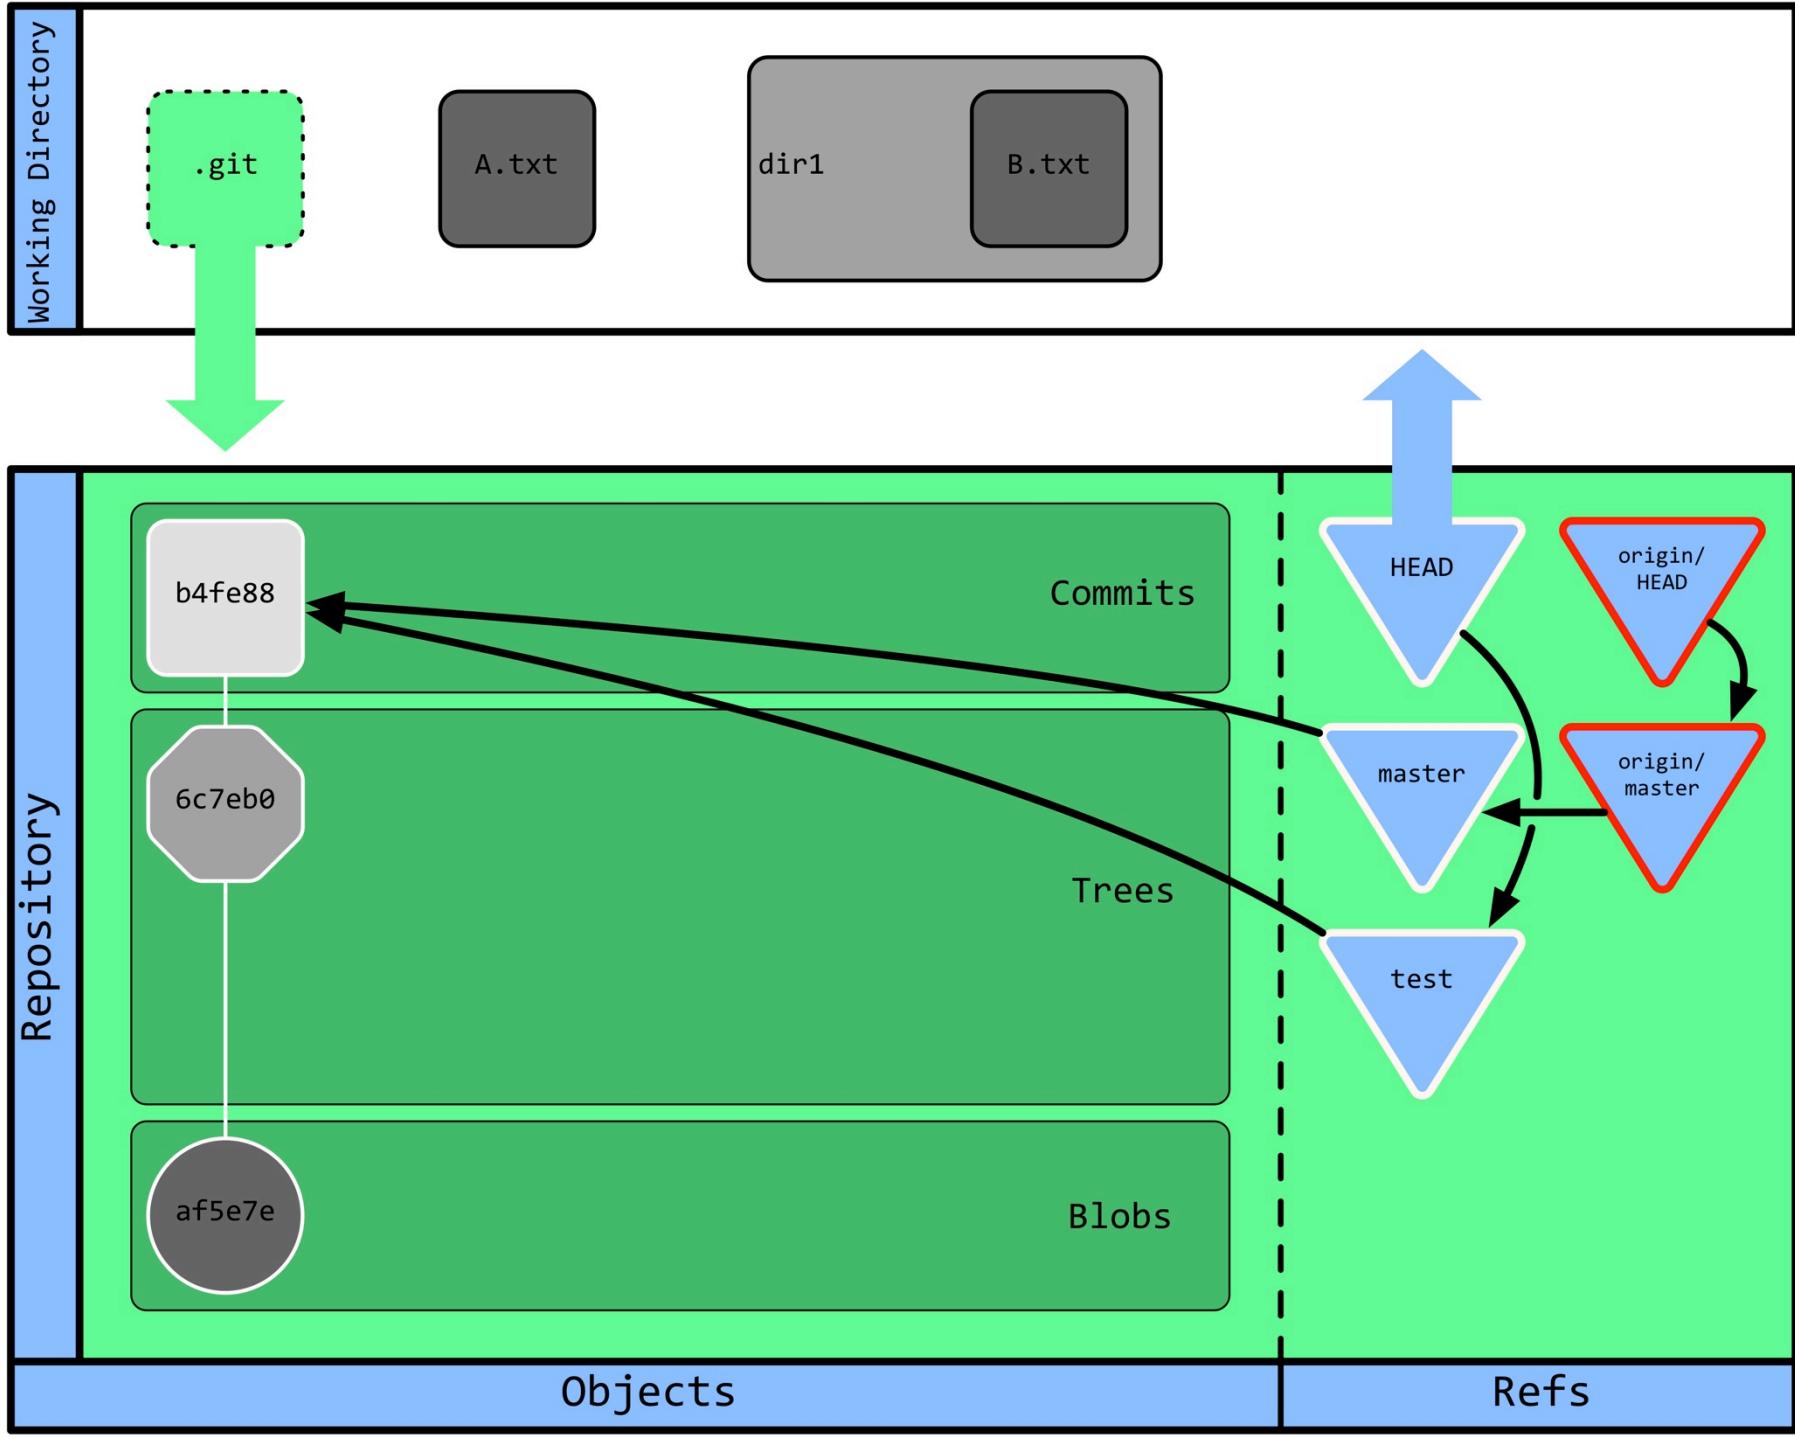


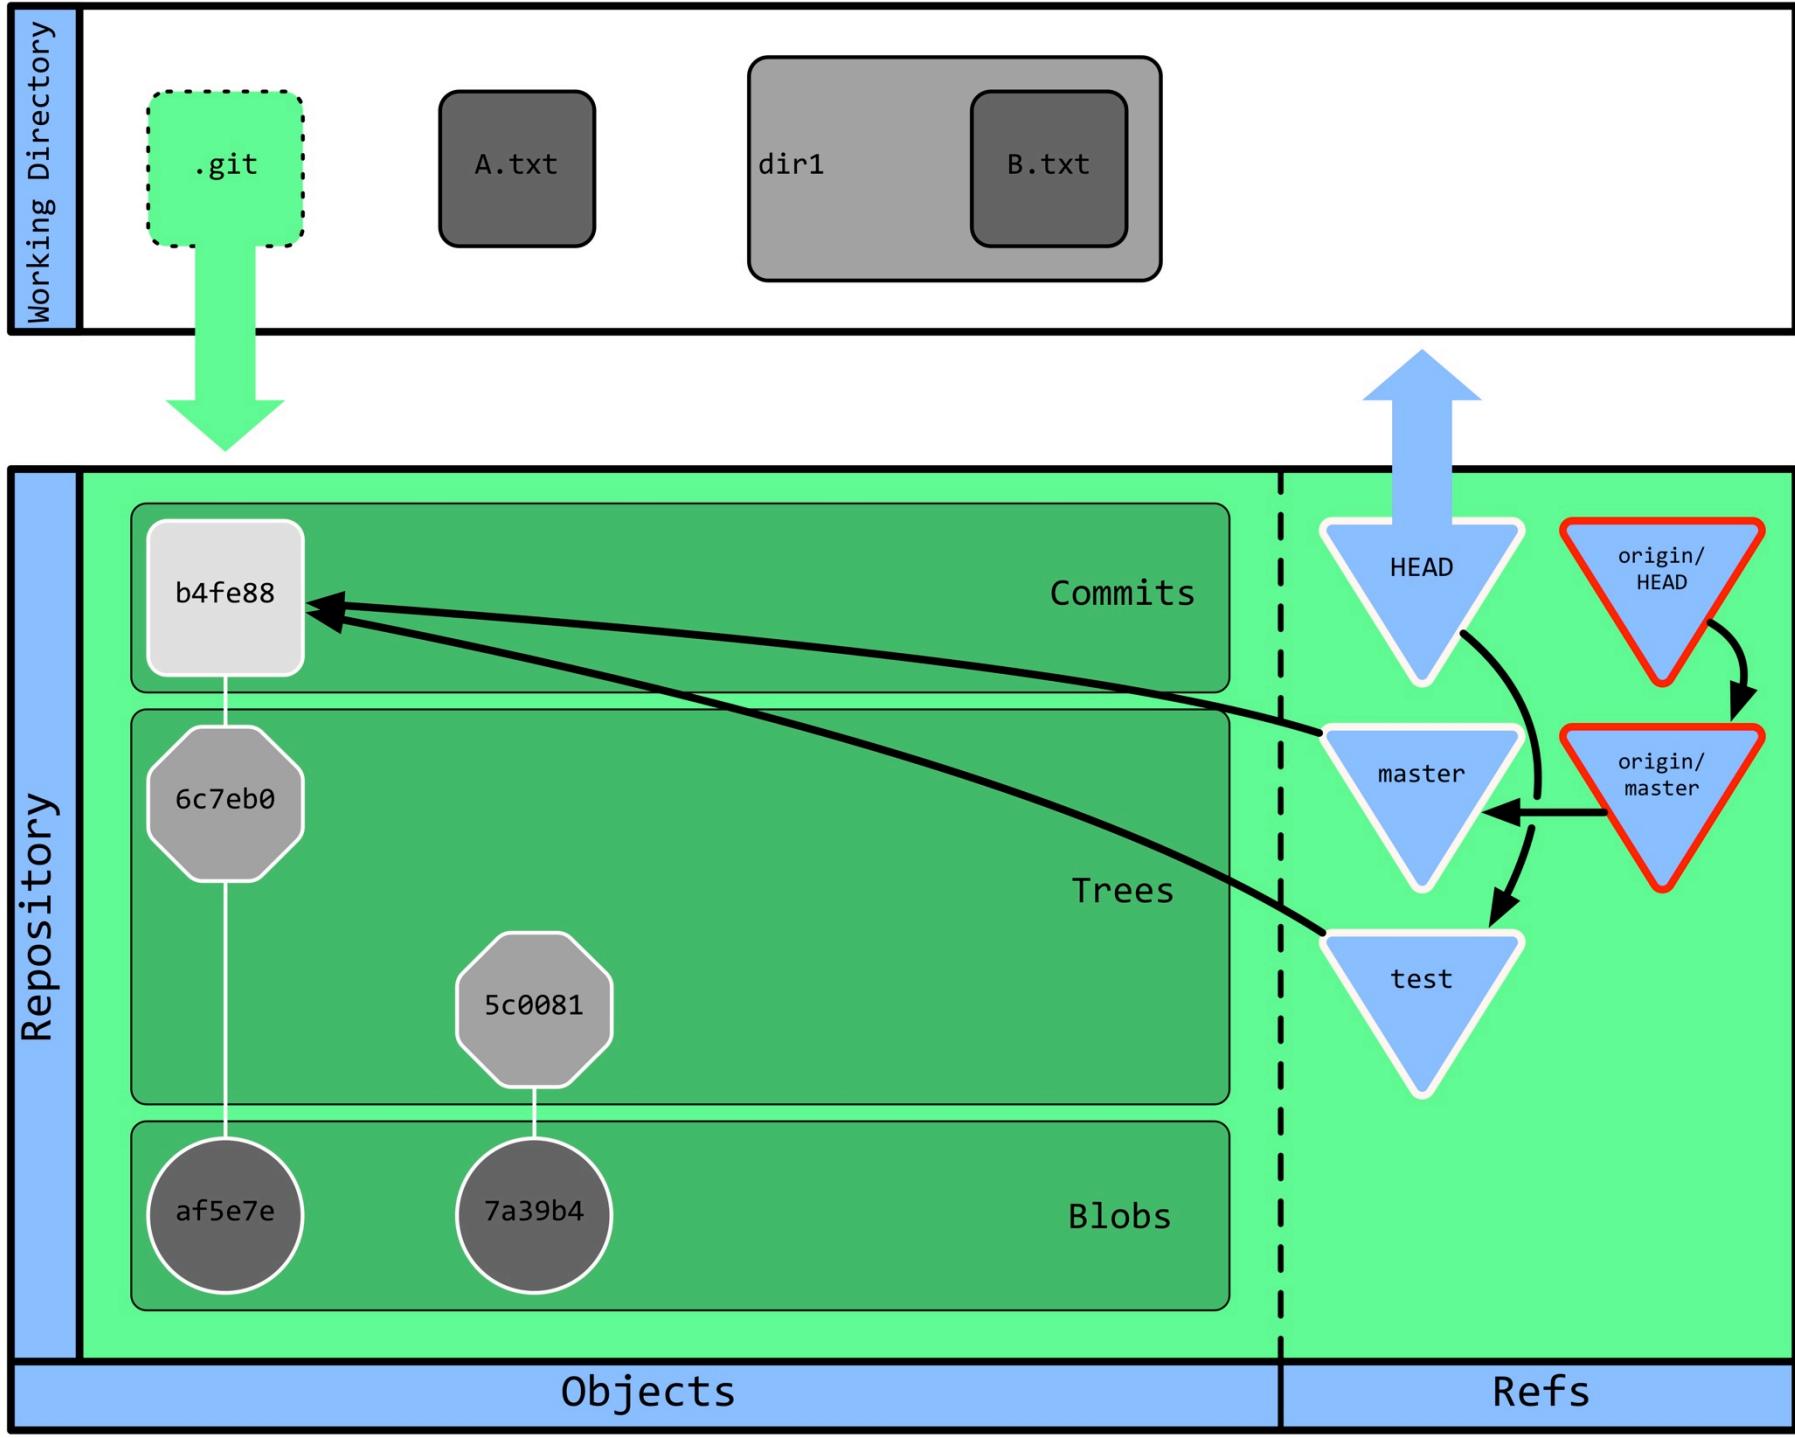
**Let's create a directory: dir1
and a file: B.txt**



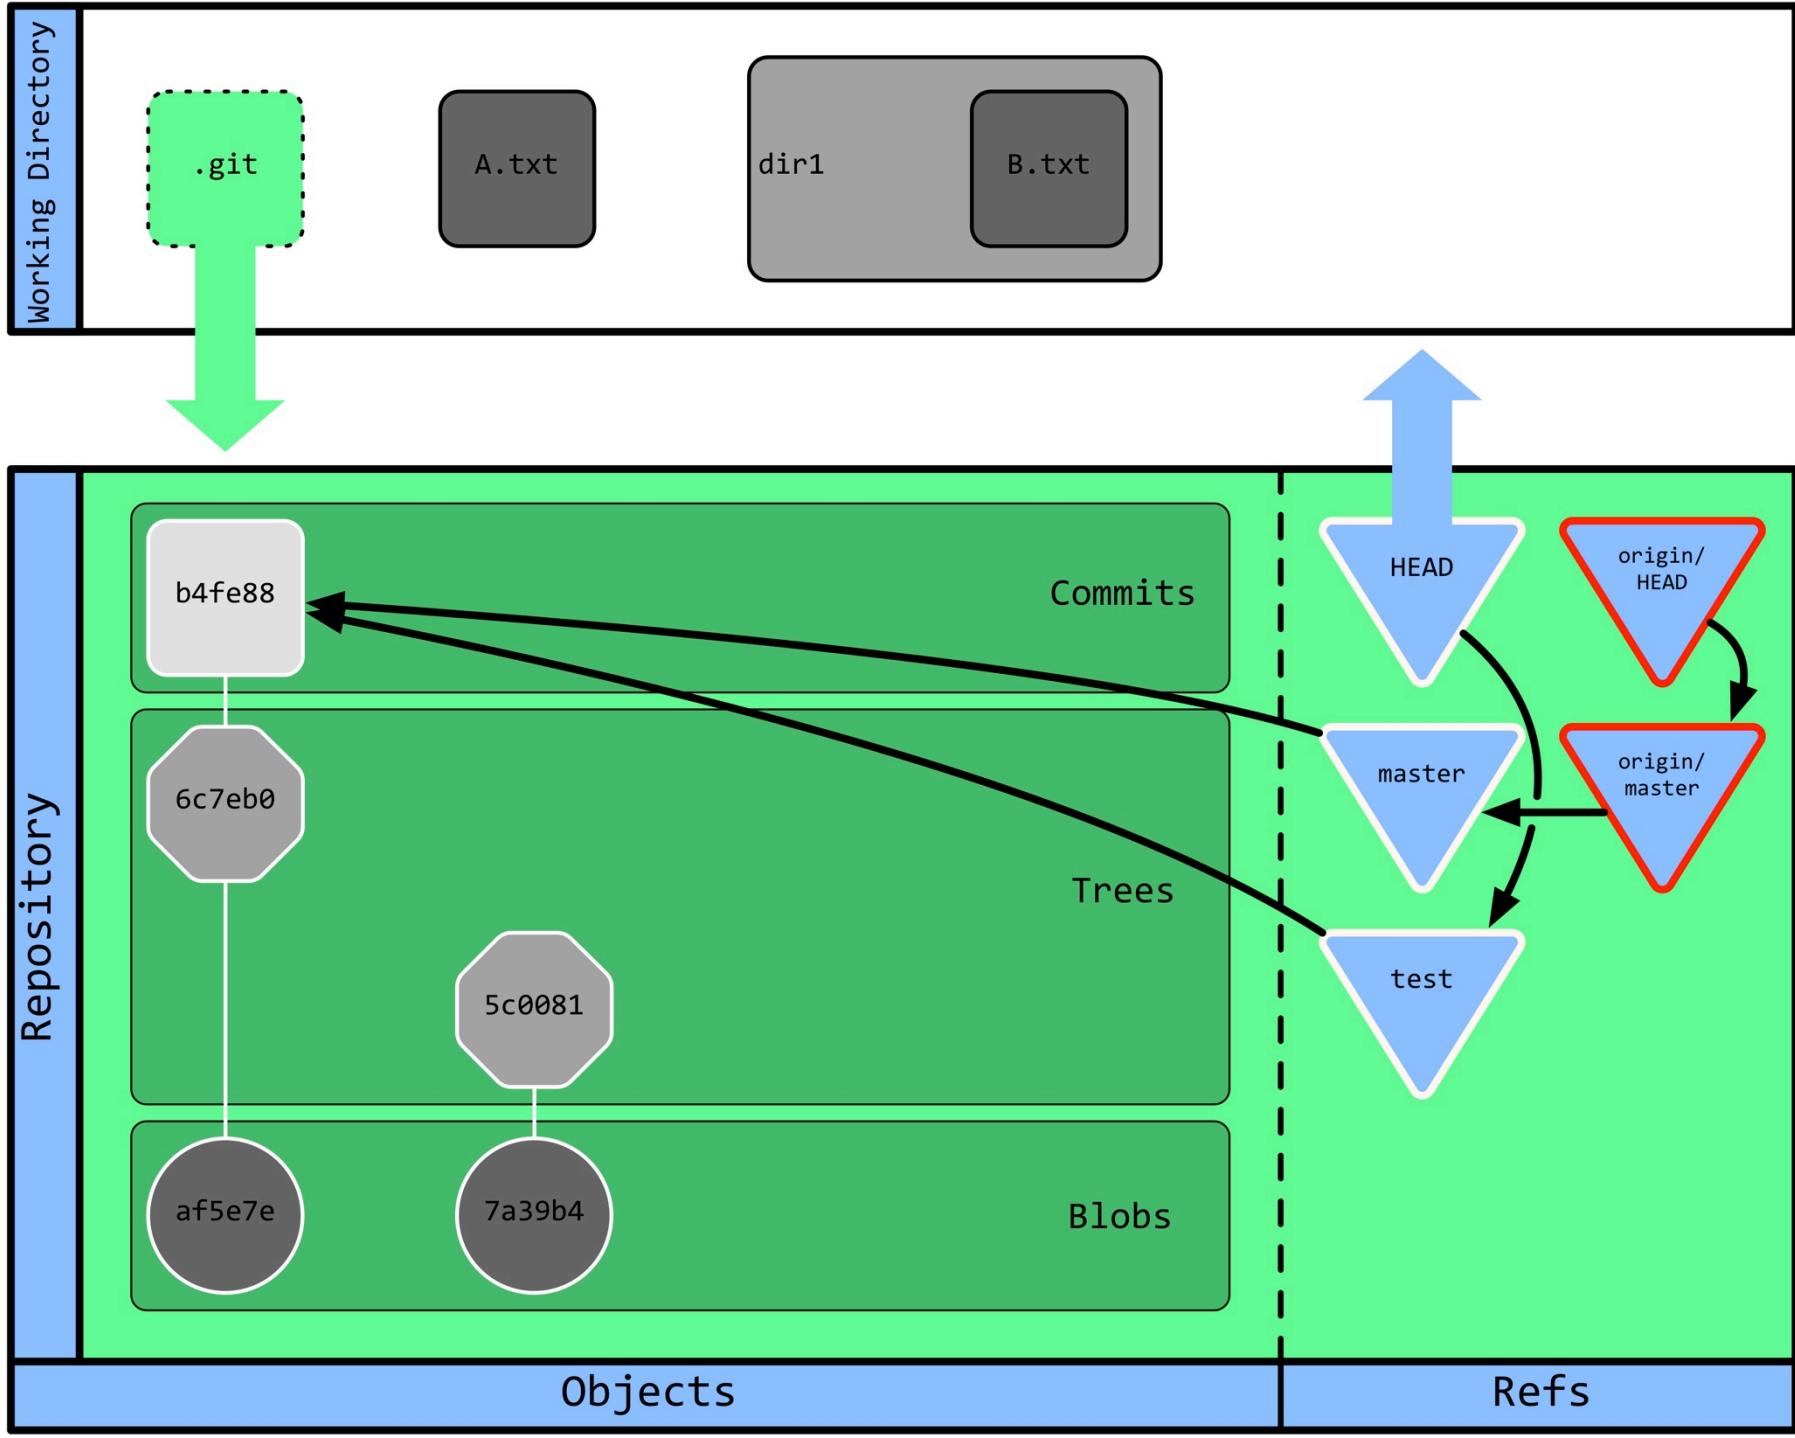


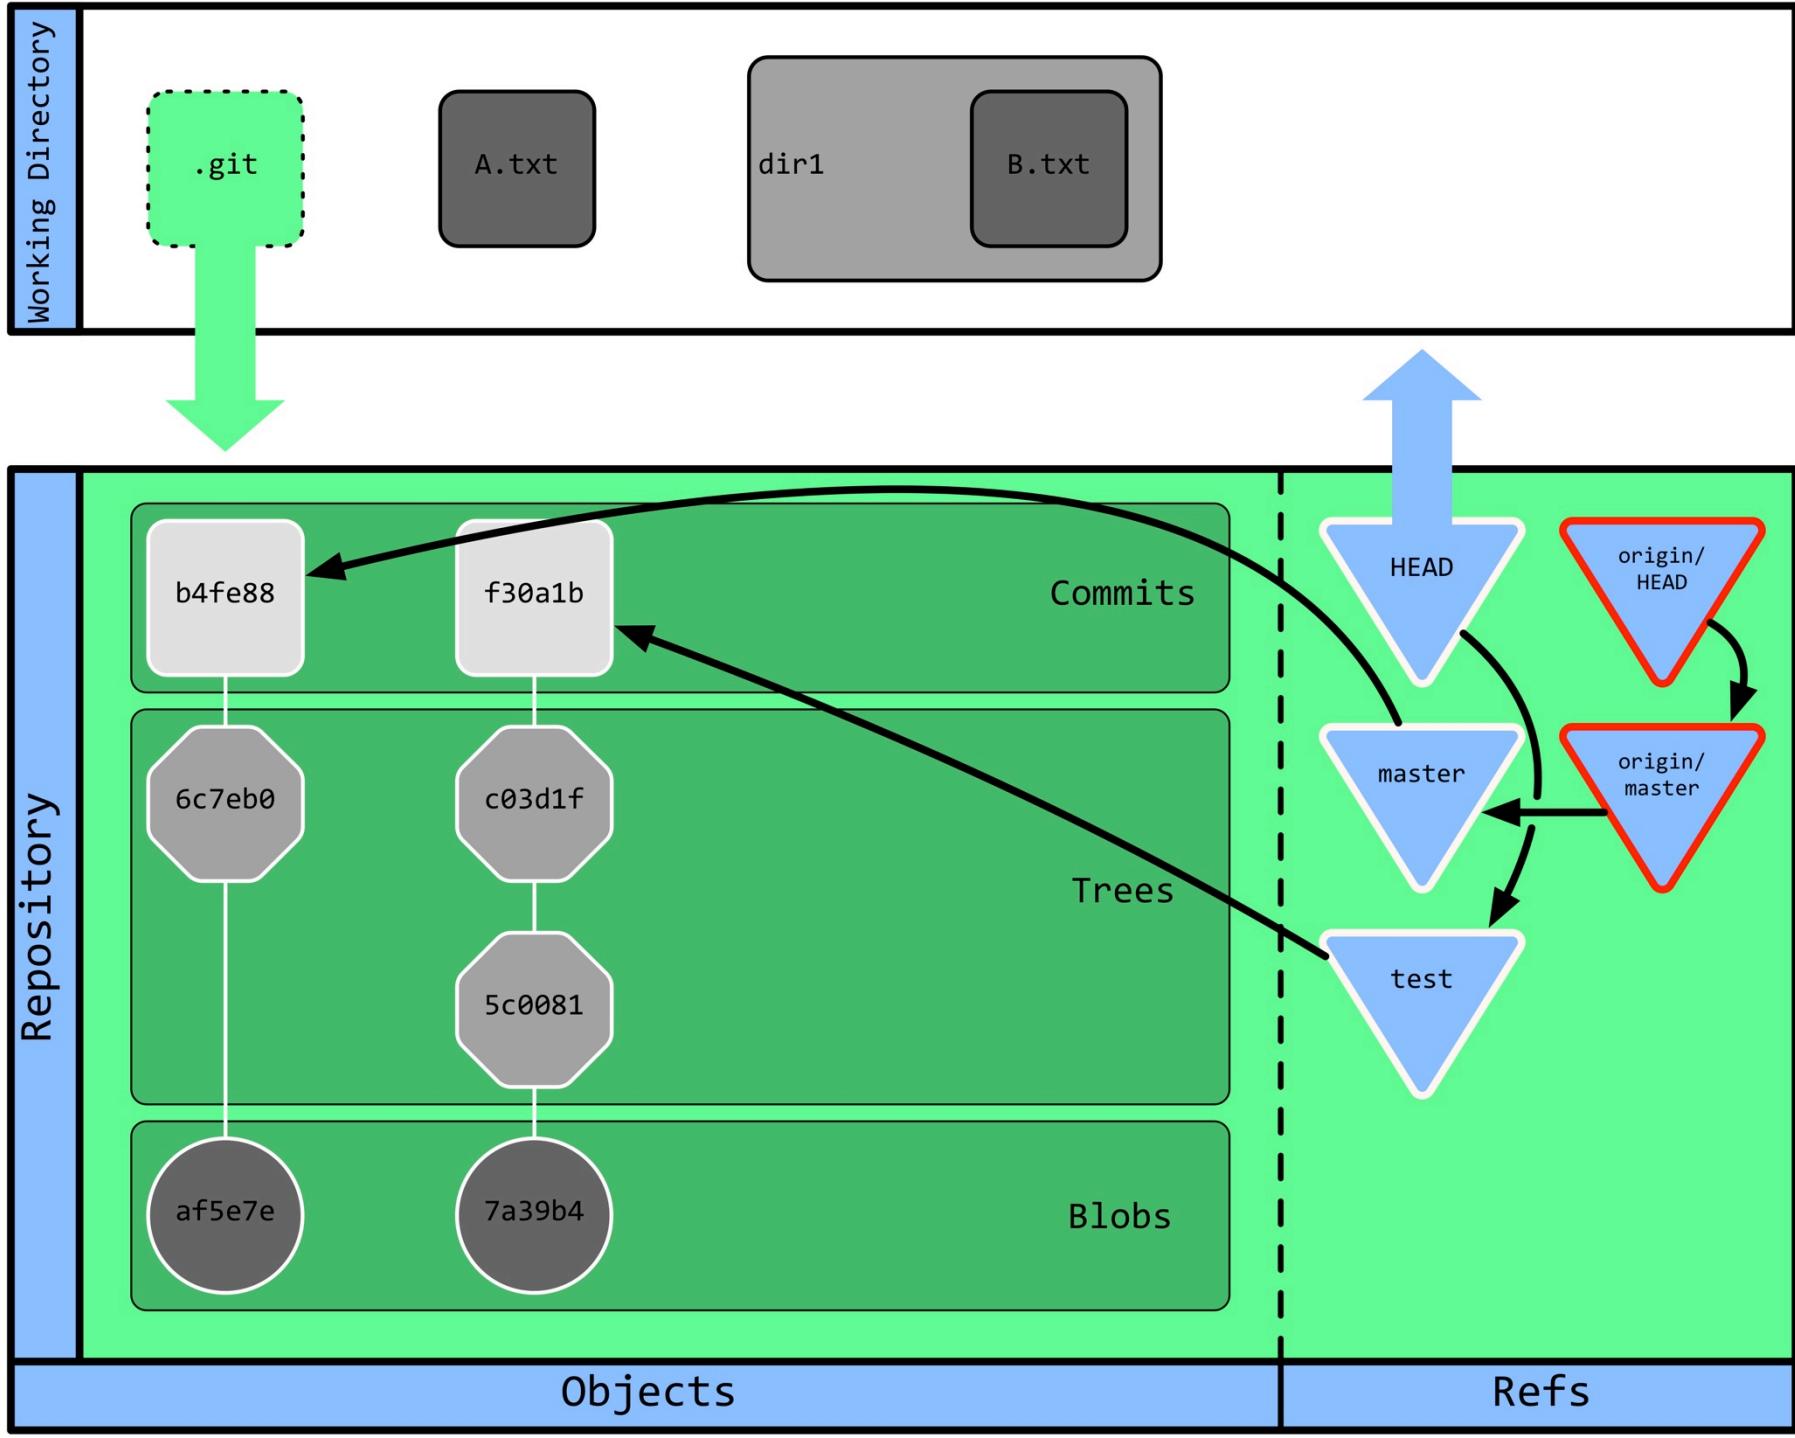
git add dir1/B.txt



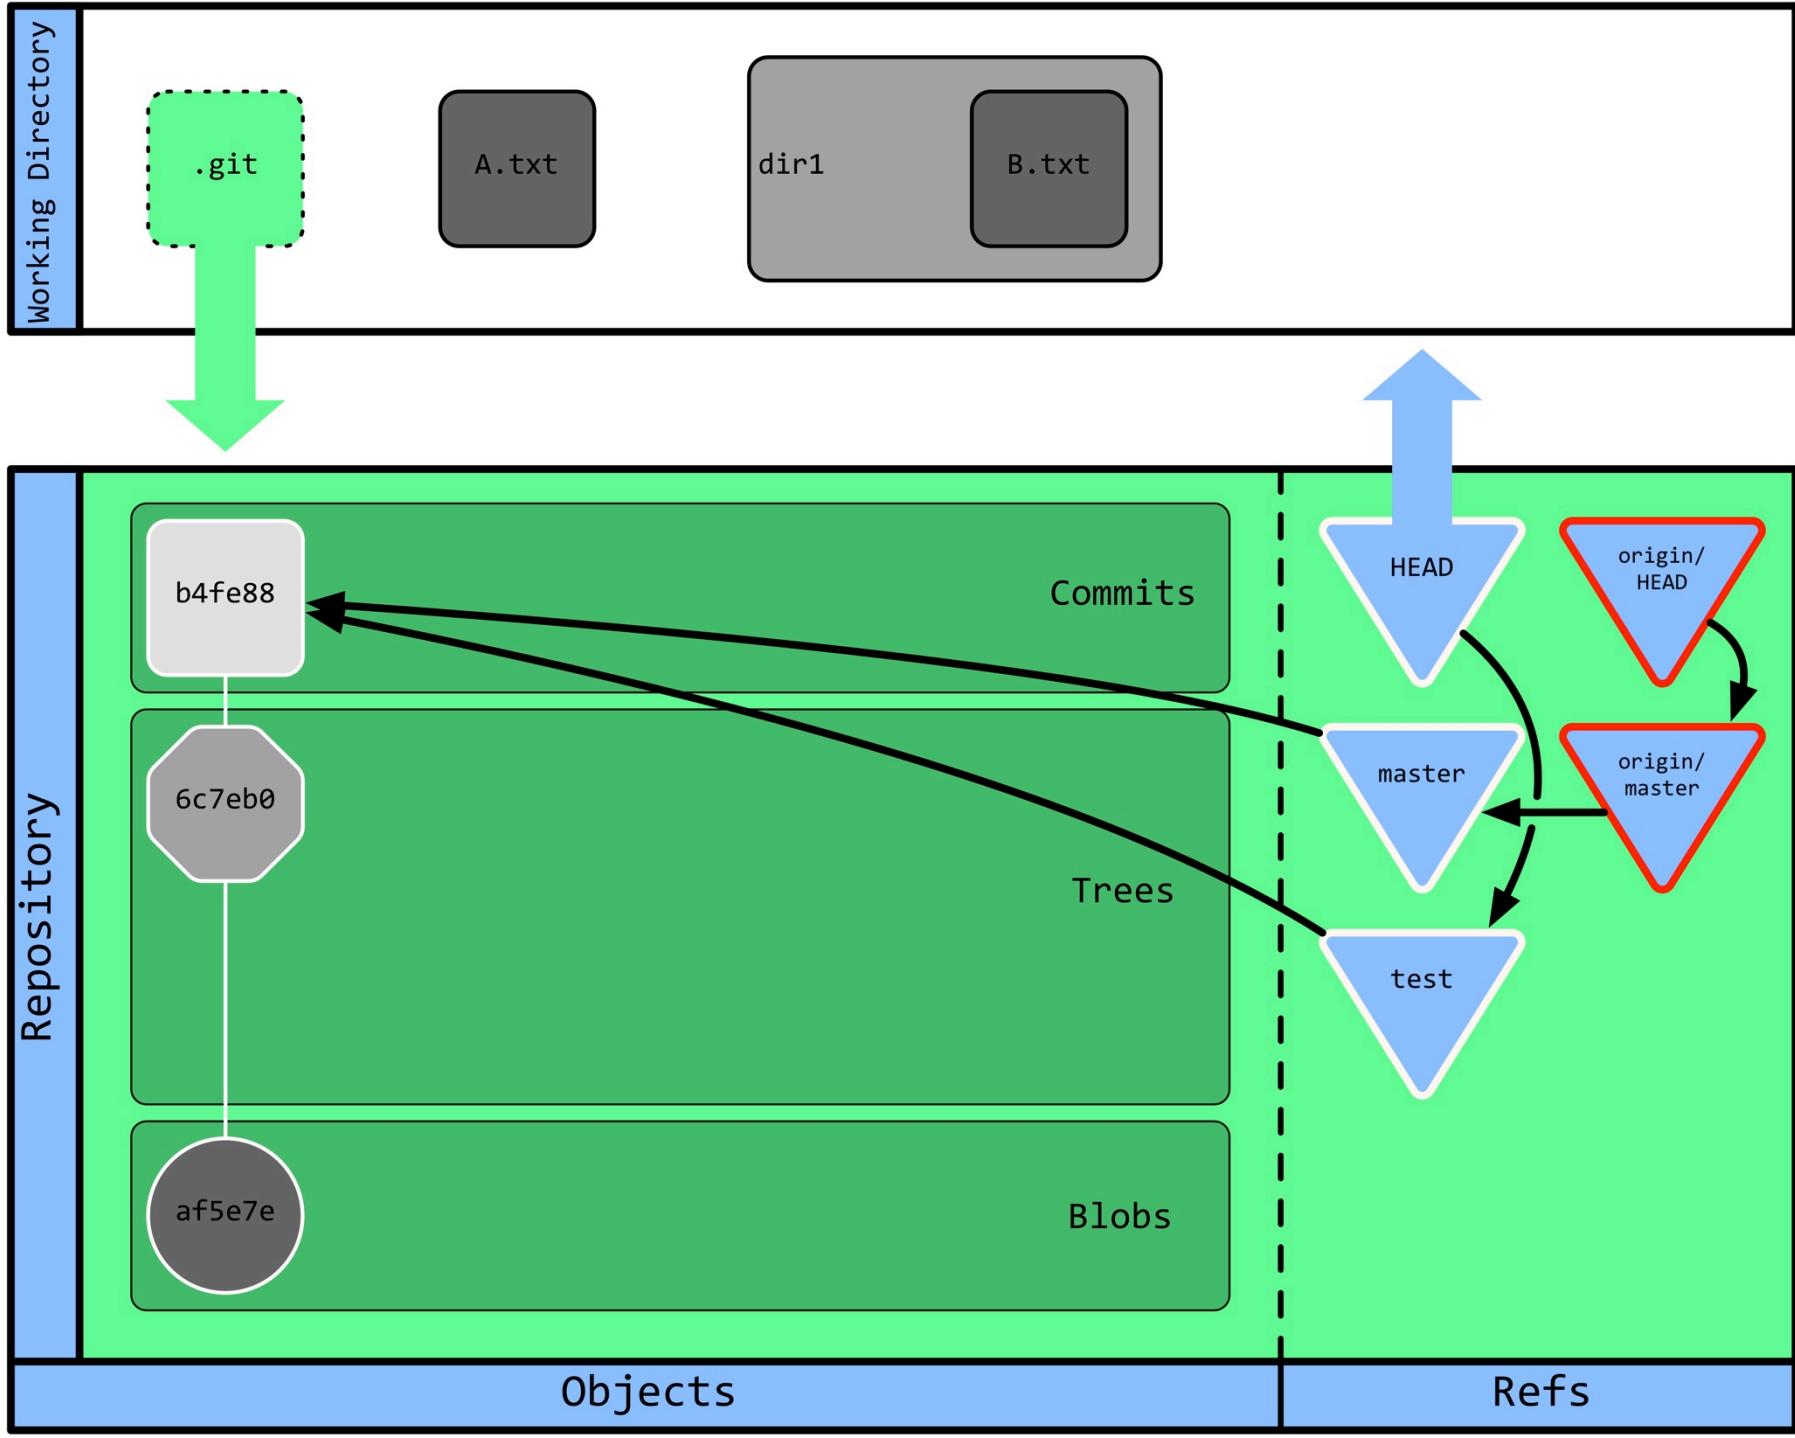


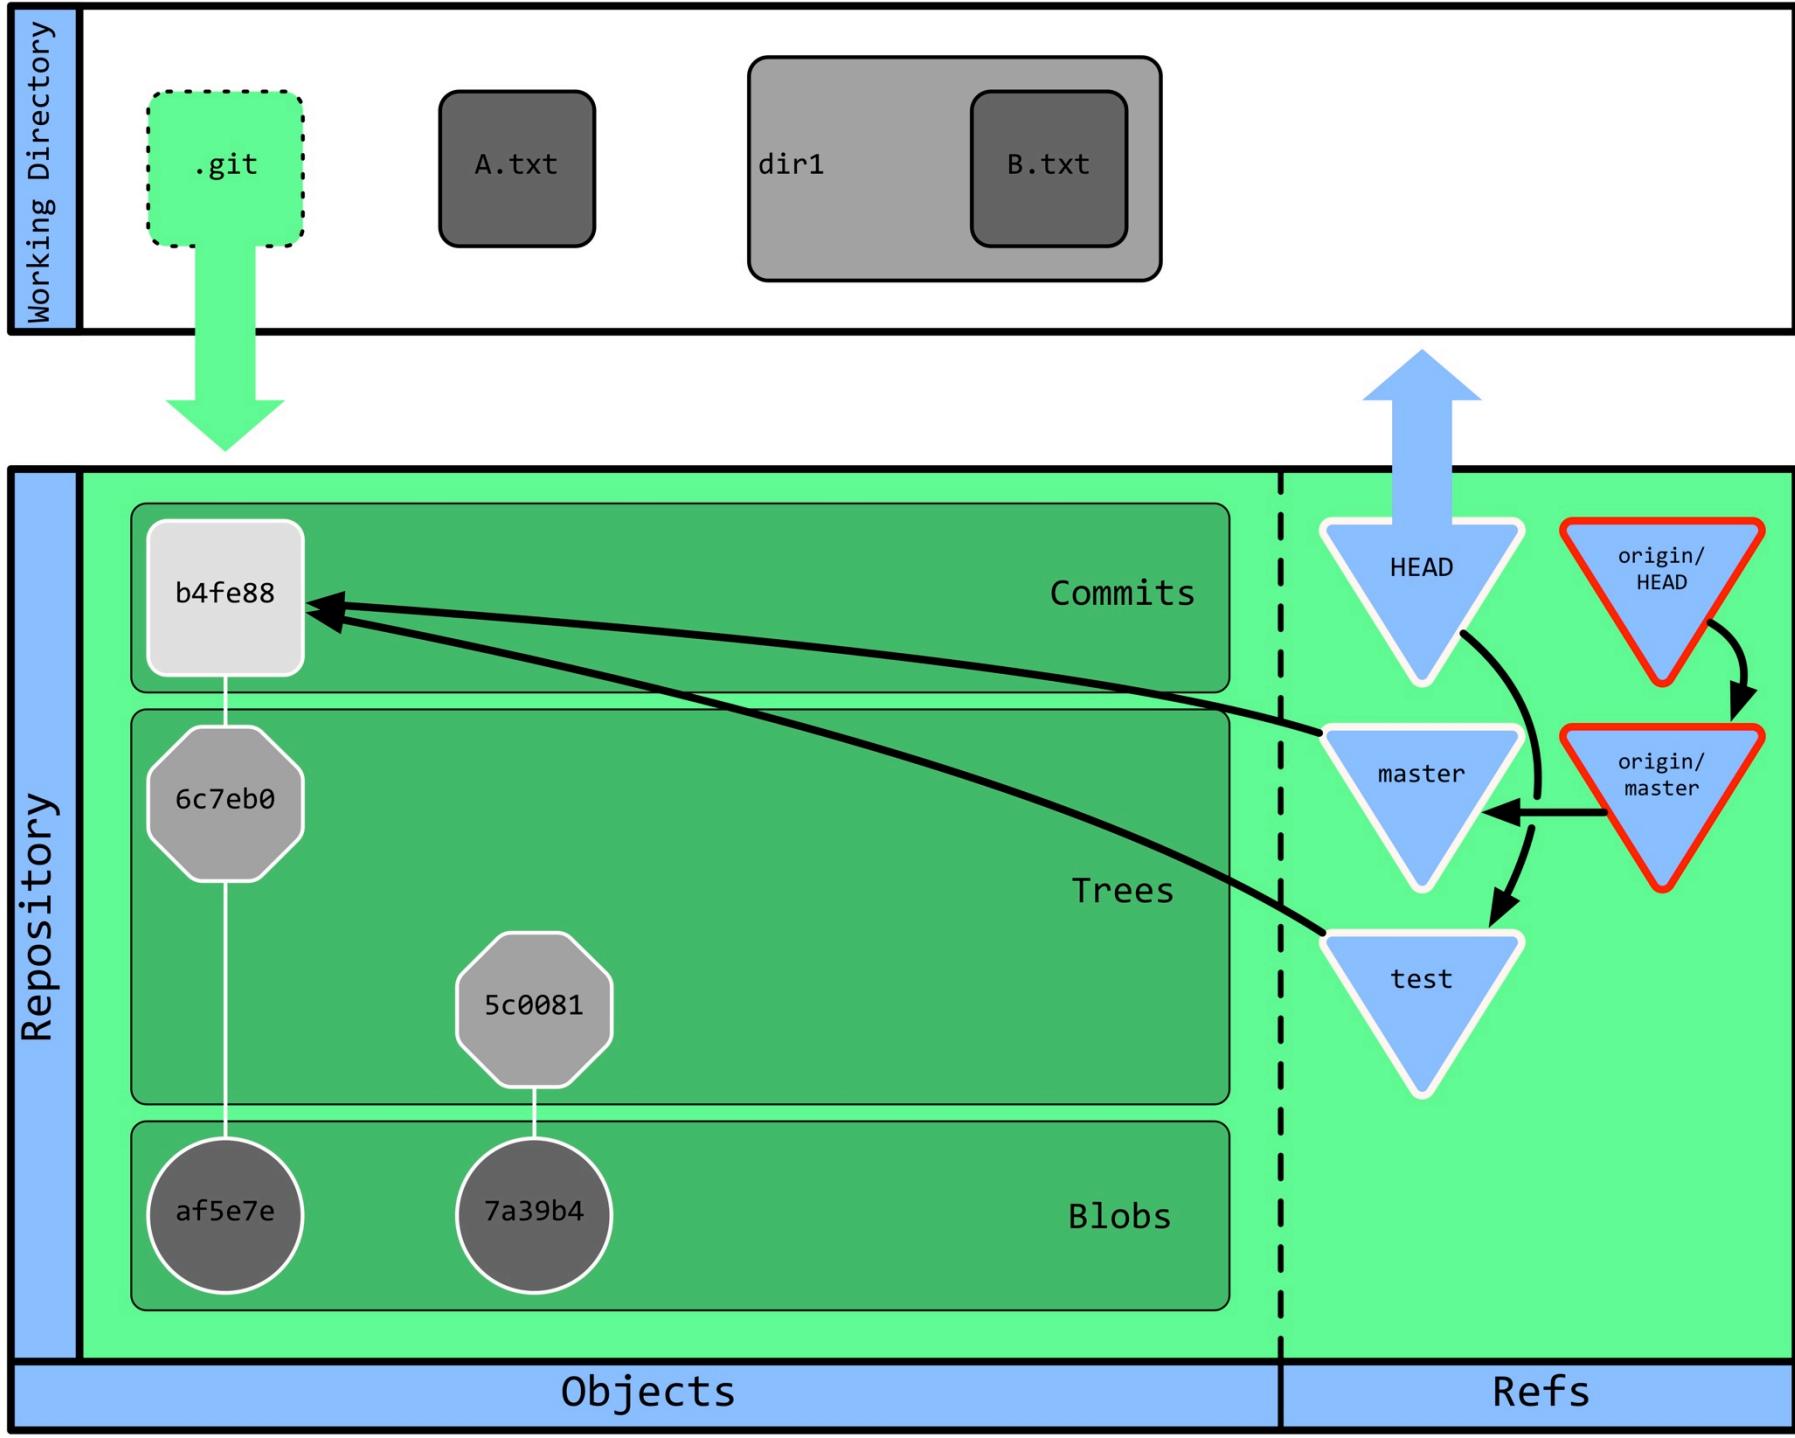
```
git commit -m “...”
```





**Wait...What's happening between the
add and commit?**





The Staging Area/Index

- As soon as a blob or tree is created by calling `add`, it is added to the **index**
- It's a literal database...

```
$ git add dir1/B.txt
$ git ls-files --stage
100644 e69de29bb2d1d6434b8b29ae775ad8c2e48c5391 0      A.txt
100644 4c5477a837a981cb182f4bad7c4425ec379a9cac 0      dir1/B.txt
$ git commit -m "Add file B"
```

The Staging Area/Index

- Calling `git commit` takes the blobs and trees in the index, adds a top-level tree object and then wraps them all up in a commit object

Branches, again

- Also notice that the **test** branch pointer automatically moved after the commit...

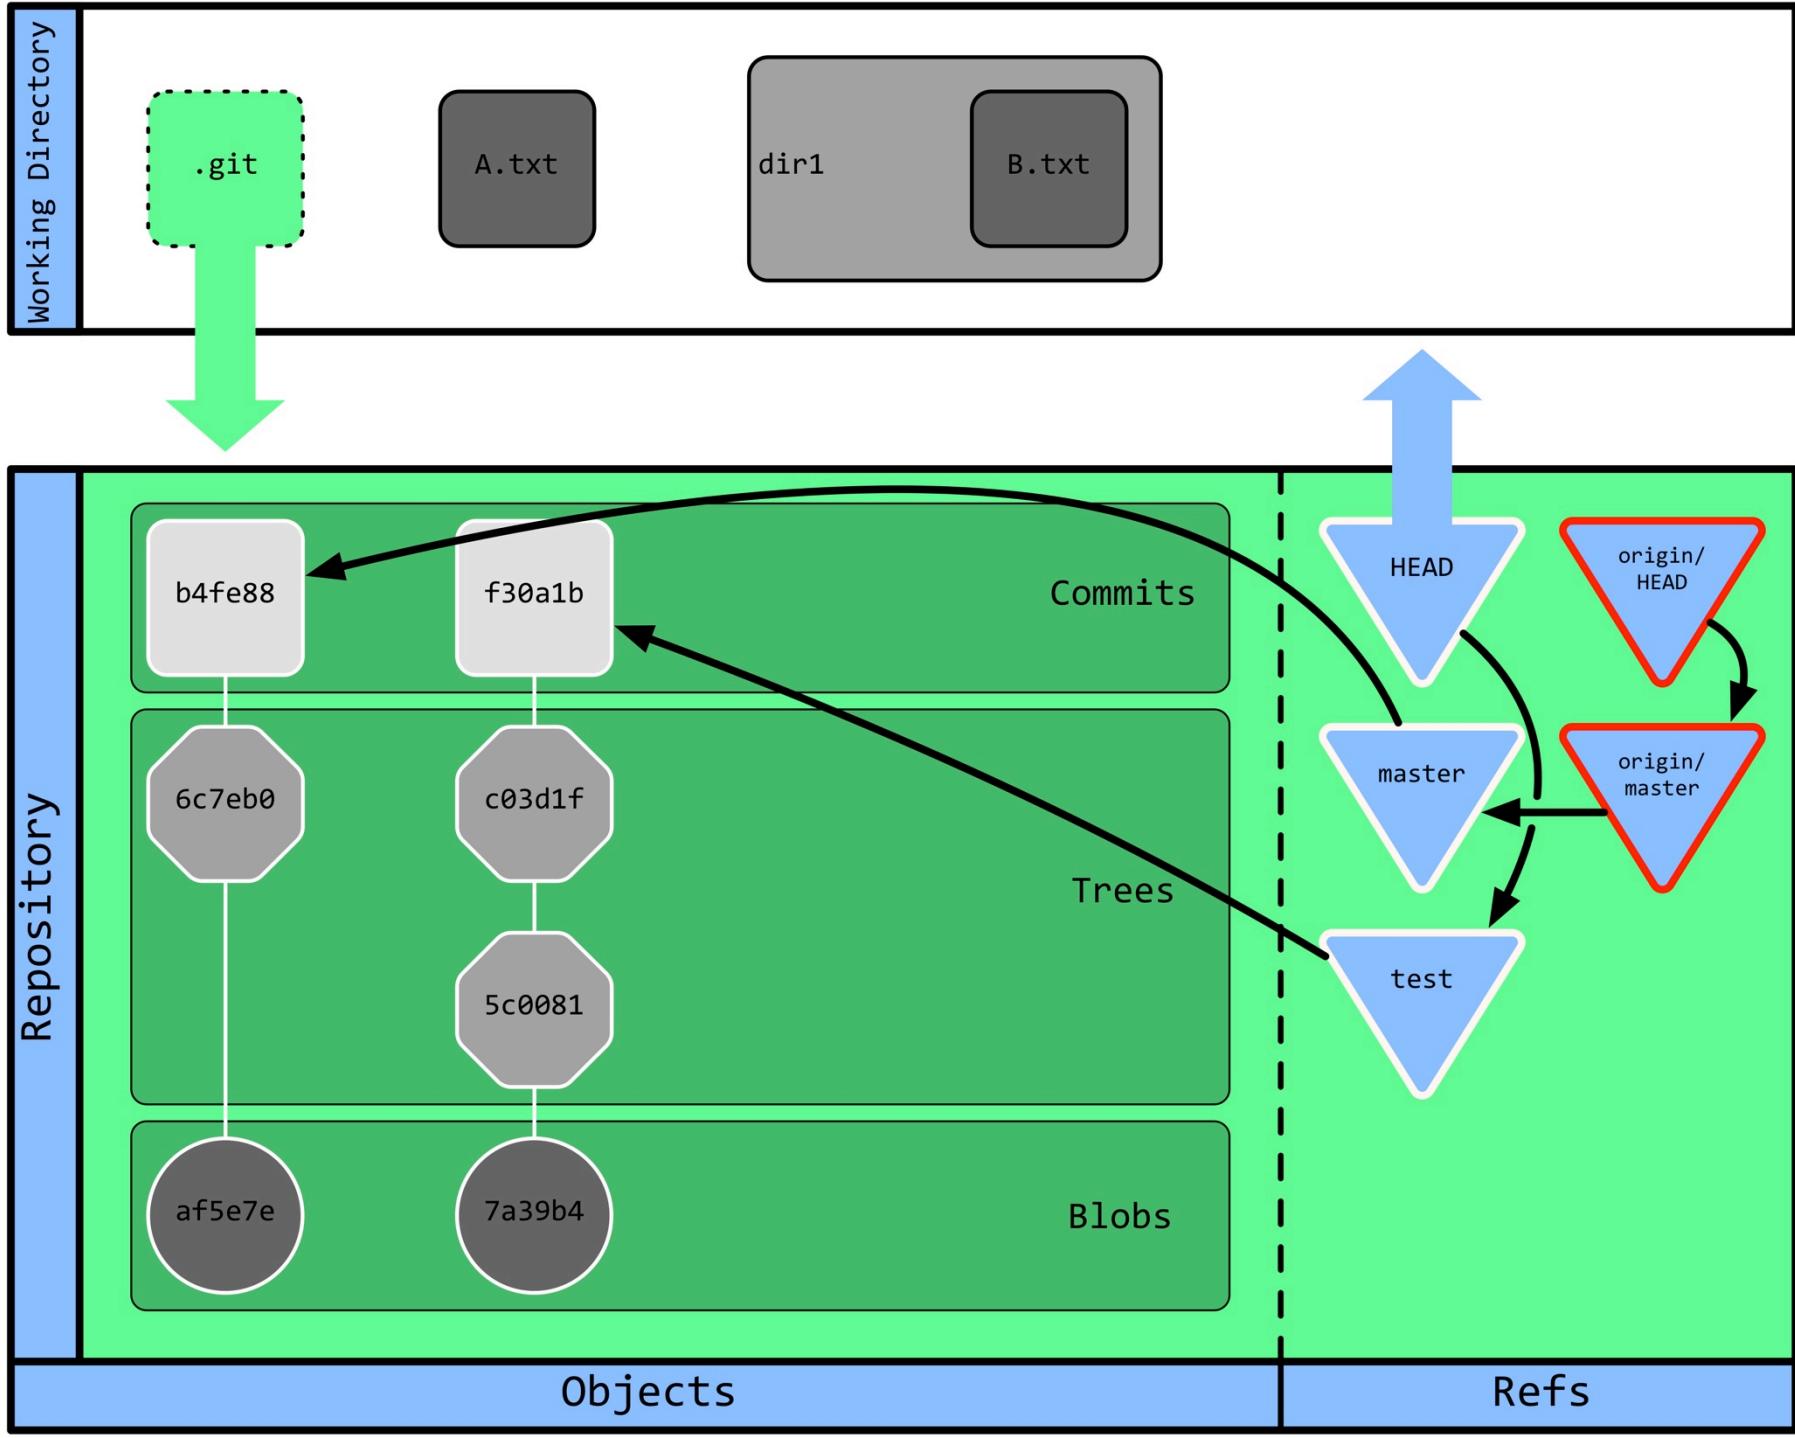
```
$ git ls
* 3e6d961 - (HEAD -> test) Add file B (6 minutes ago) <Joe Gibson>
* 901afad - (master) Add file A (8 minutes ago) <Joe Gibson>

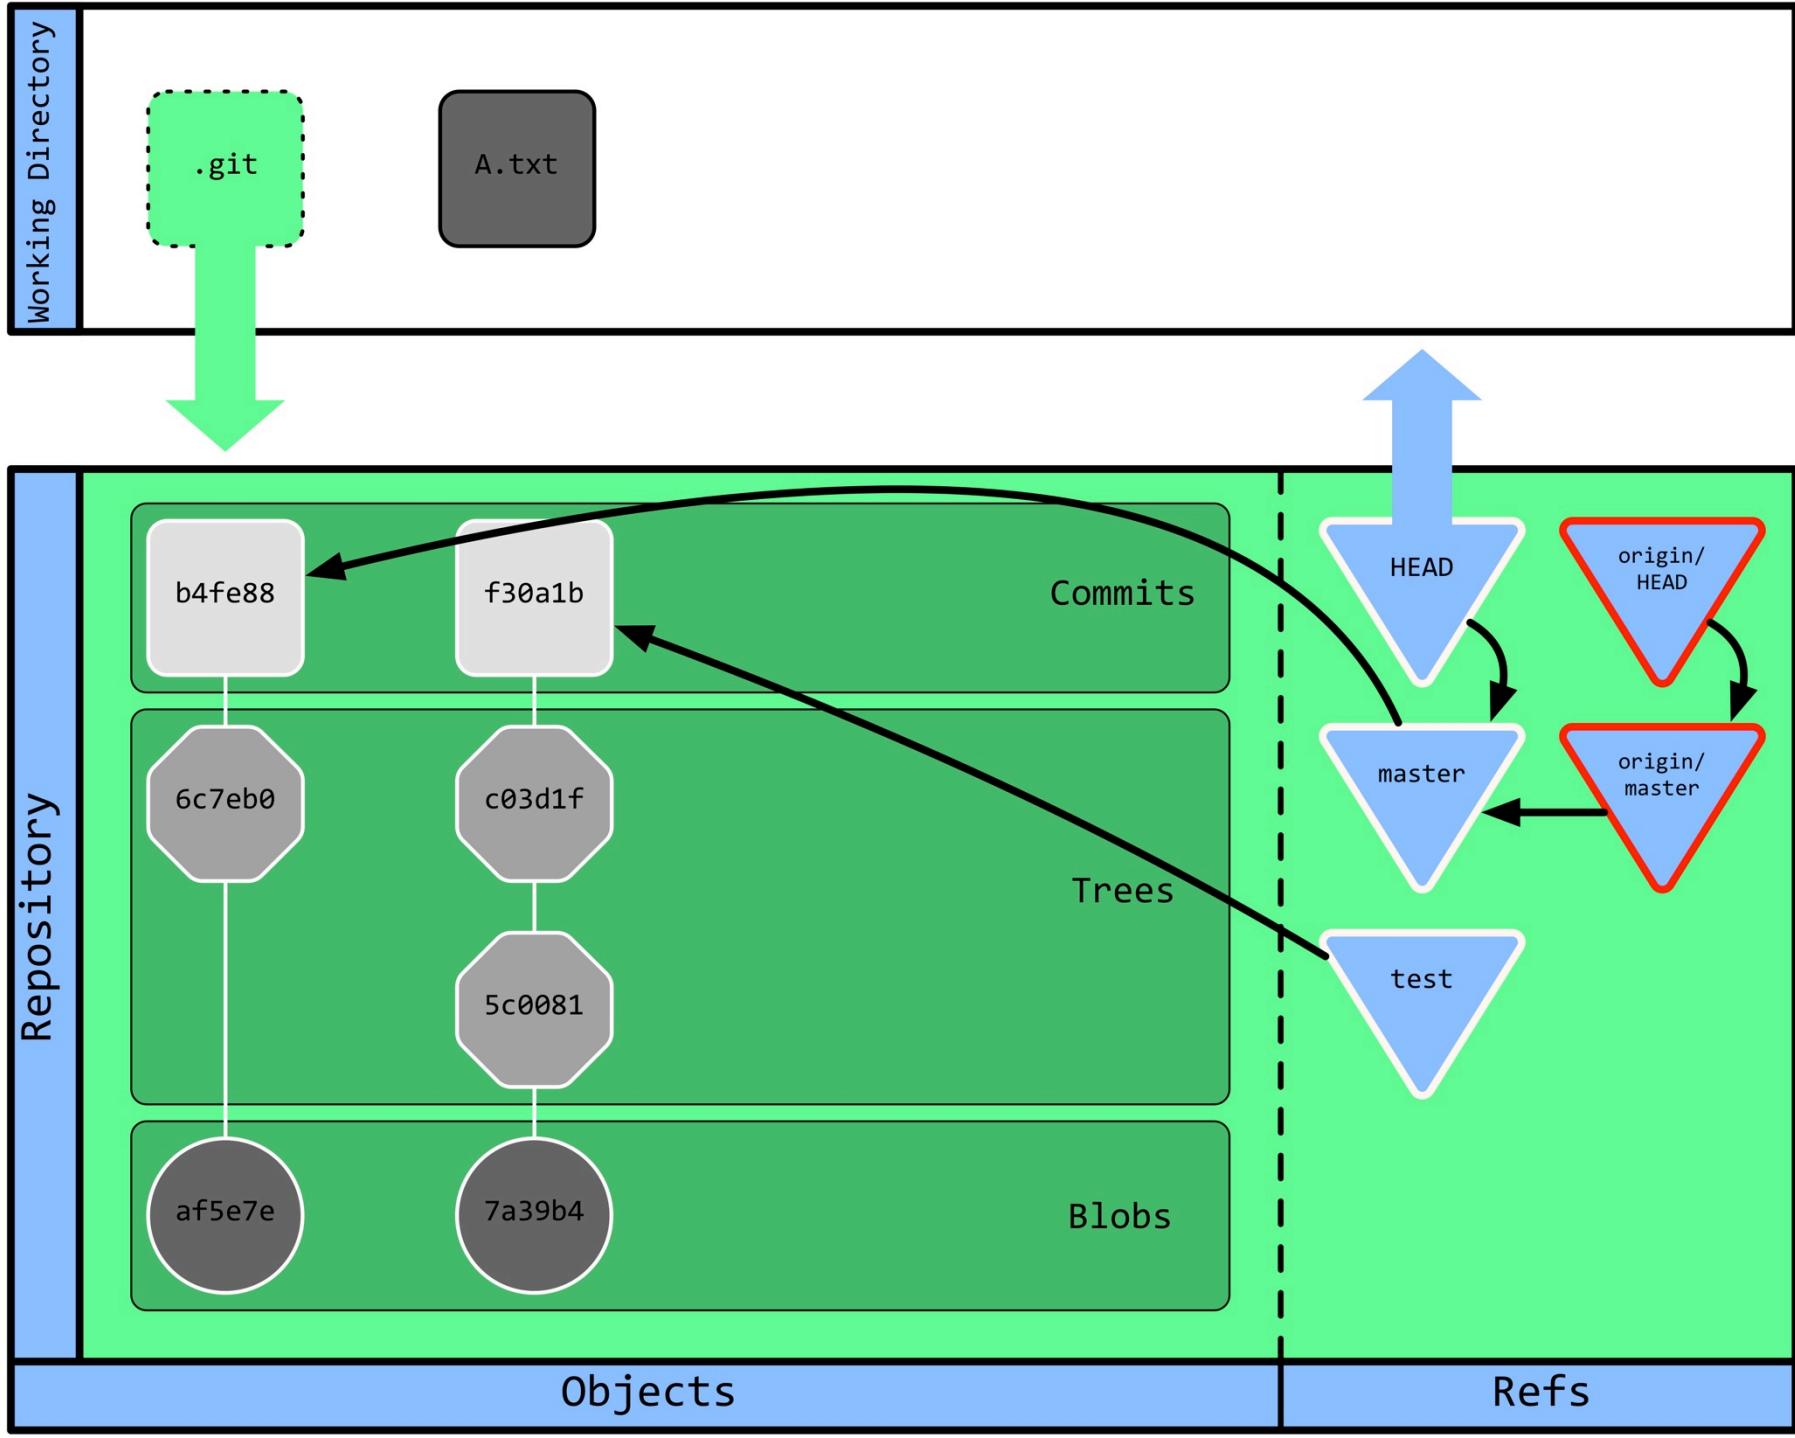
$ git cat-file -p 3e6d961
tree 43222cd2e96acac8c83f5162226721f3351c2208
parent 901afadf676ff8790af9b520708be50d0c10f645
author Joe Gibson <joseph.gibson@nasa.gov> 1469156527 -0400
committer Joe Gibson <joseph.gibson@nasa.gov> 1469156527 -0400

Add file B
```

Let's try merging test into master

git checkout master

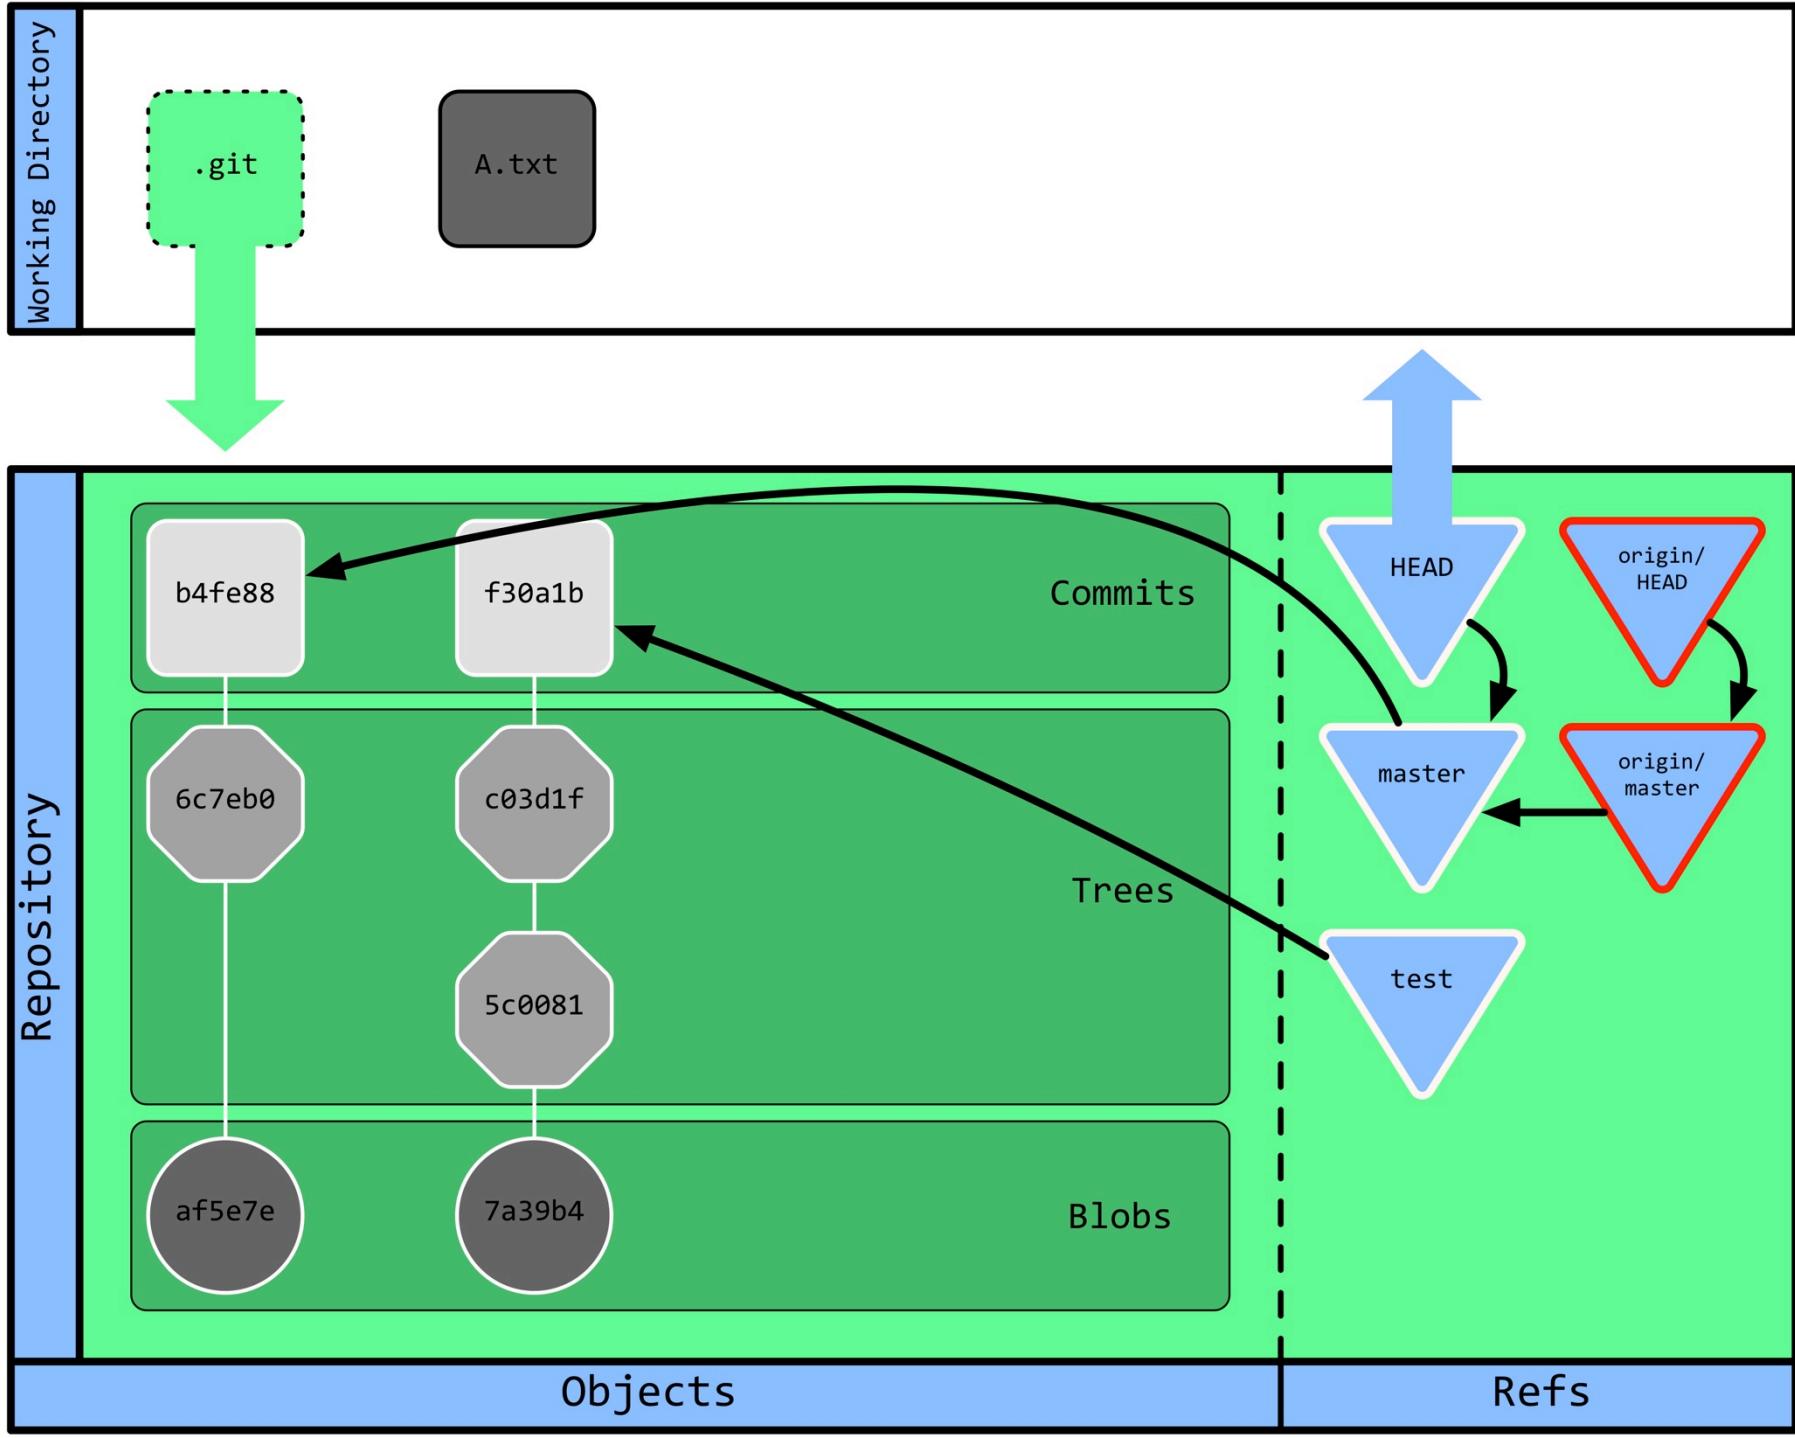


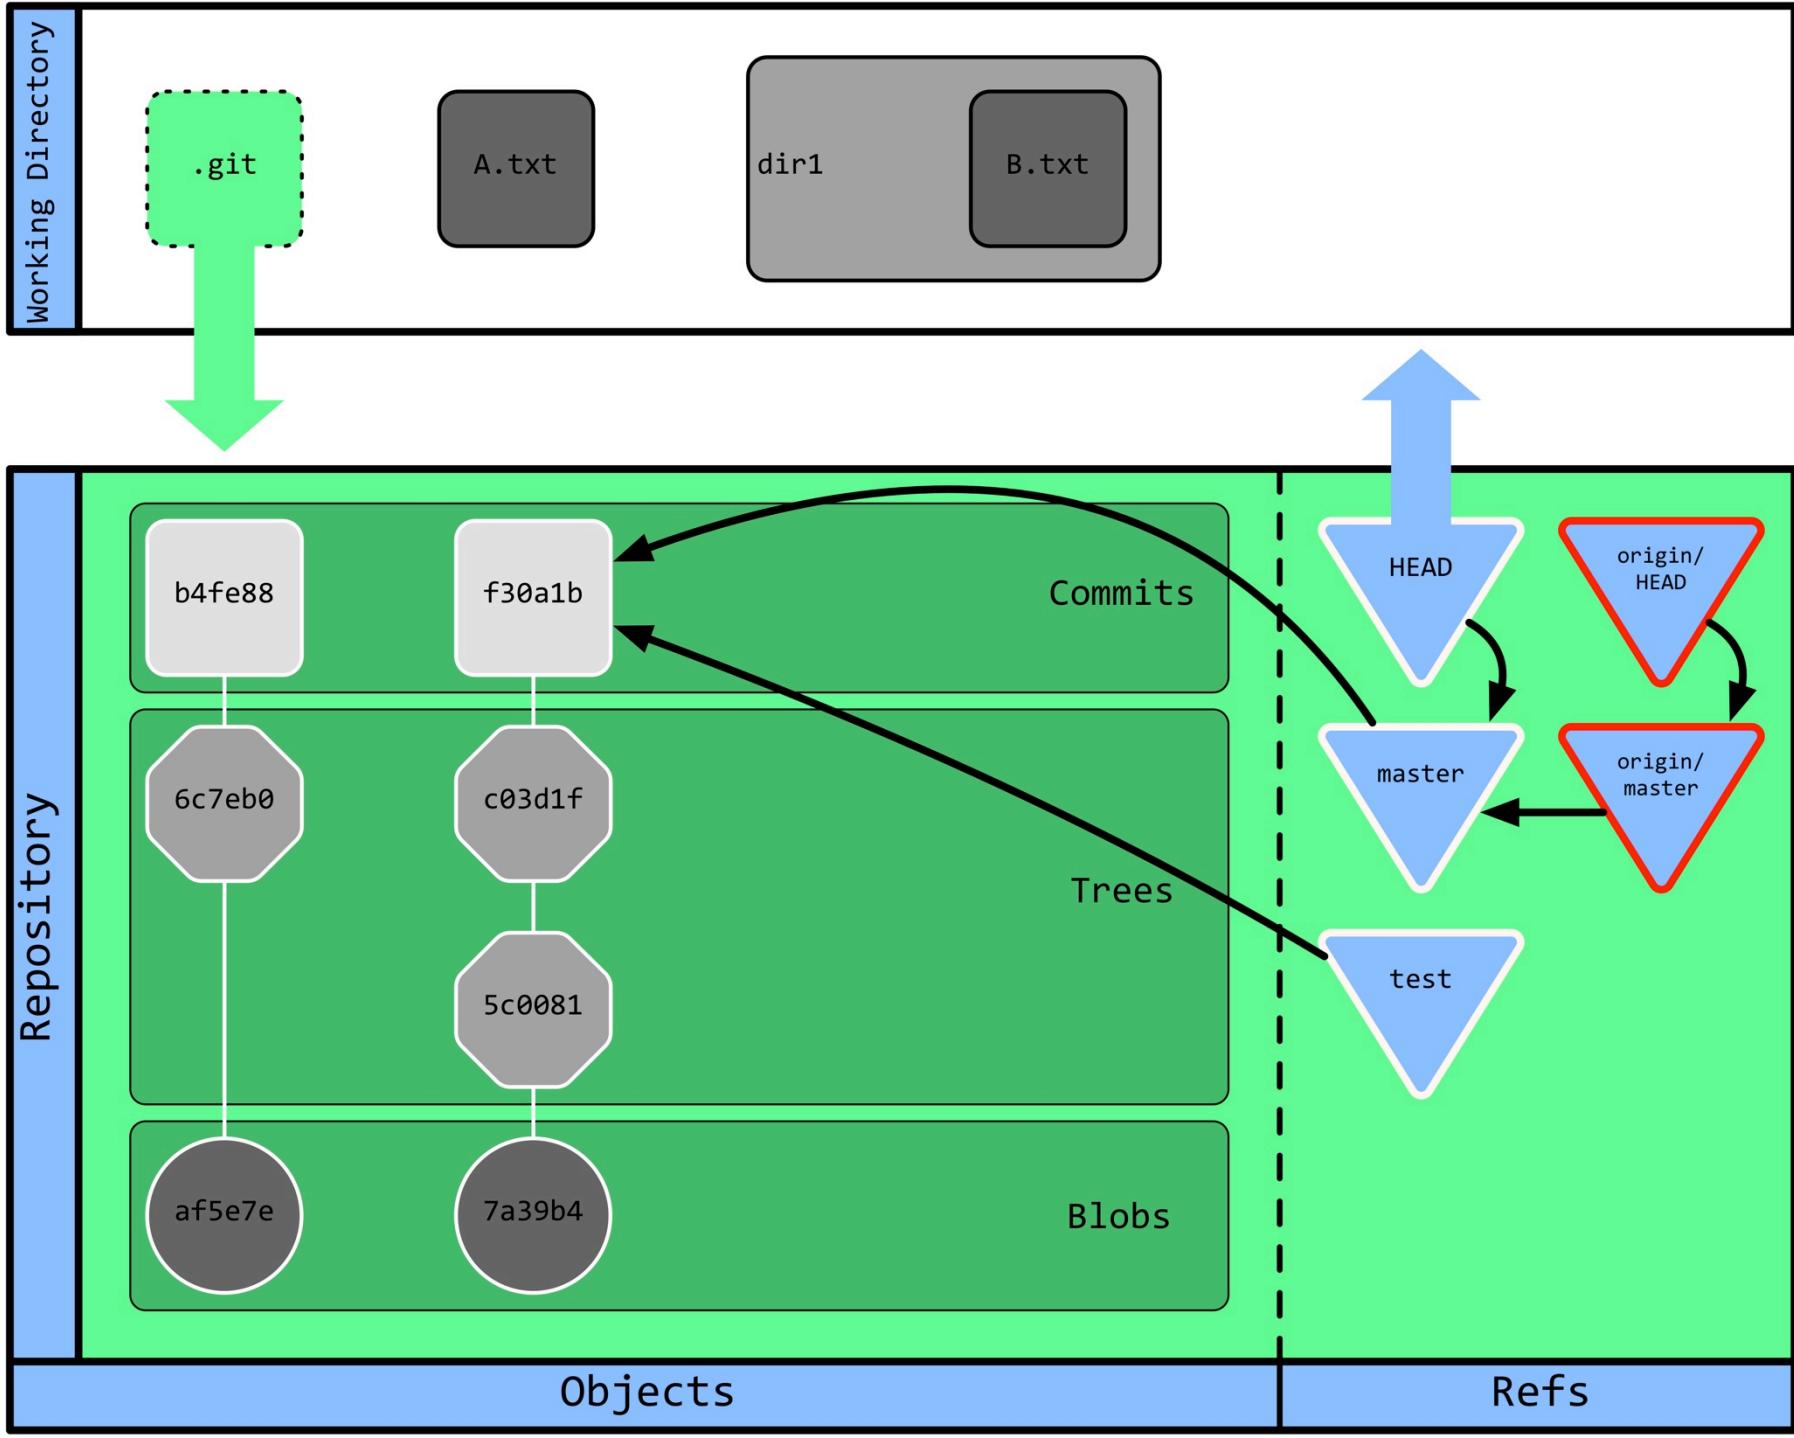


Notice how A.txt is the only file now...

Hint: HEAD points to master, and
whatever HEAD points to ends up in
your working directory.

git merge test





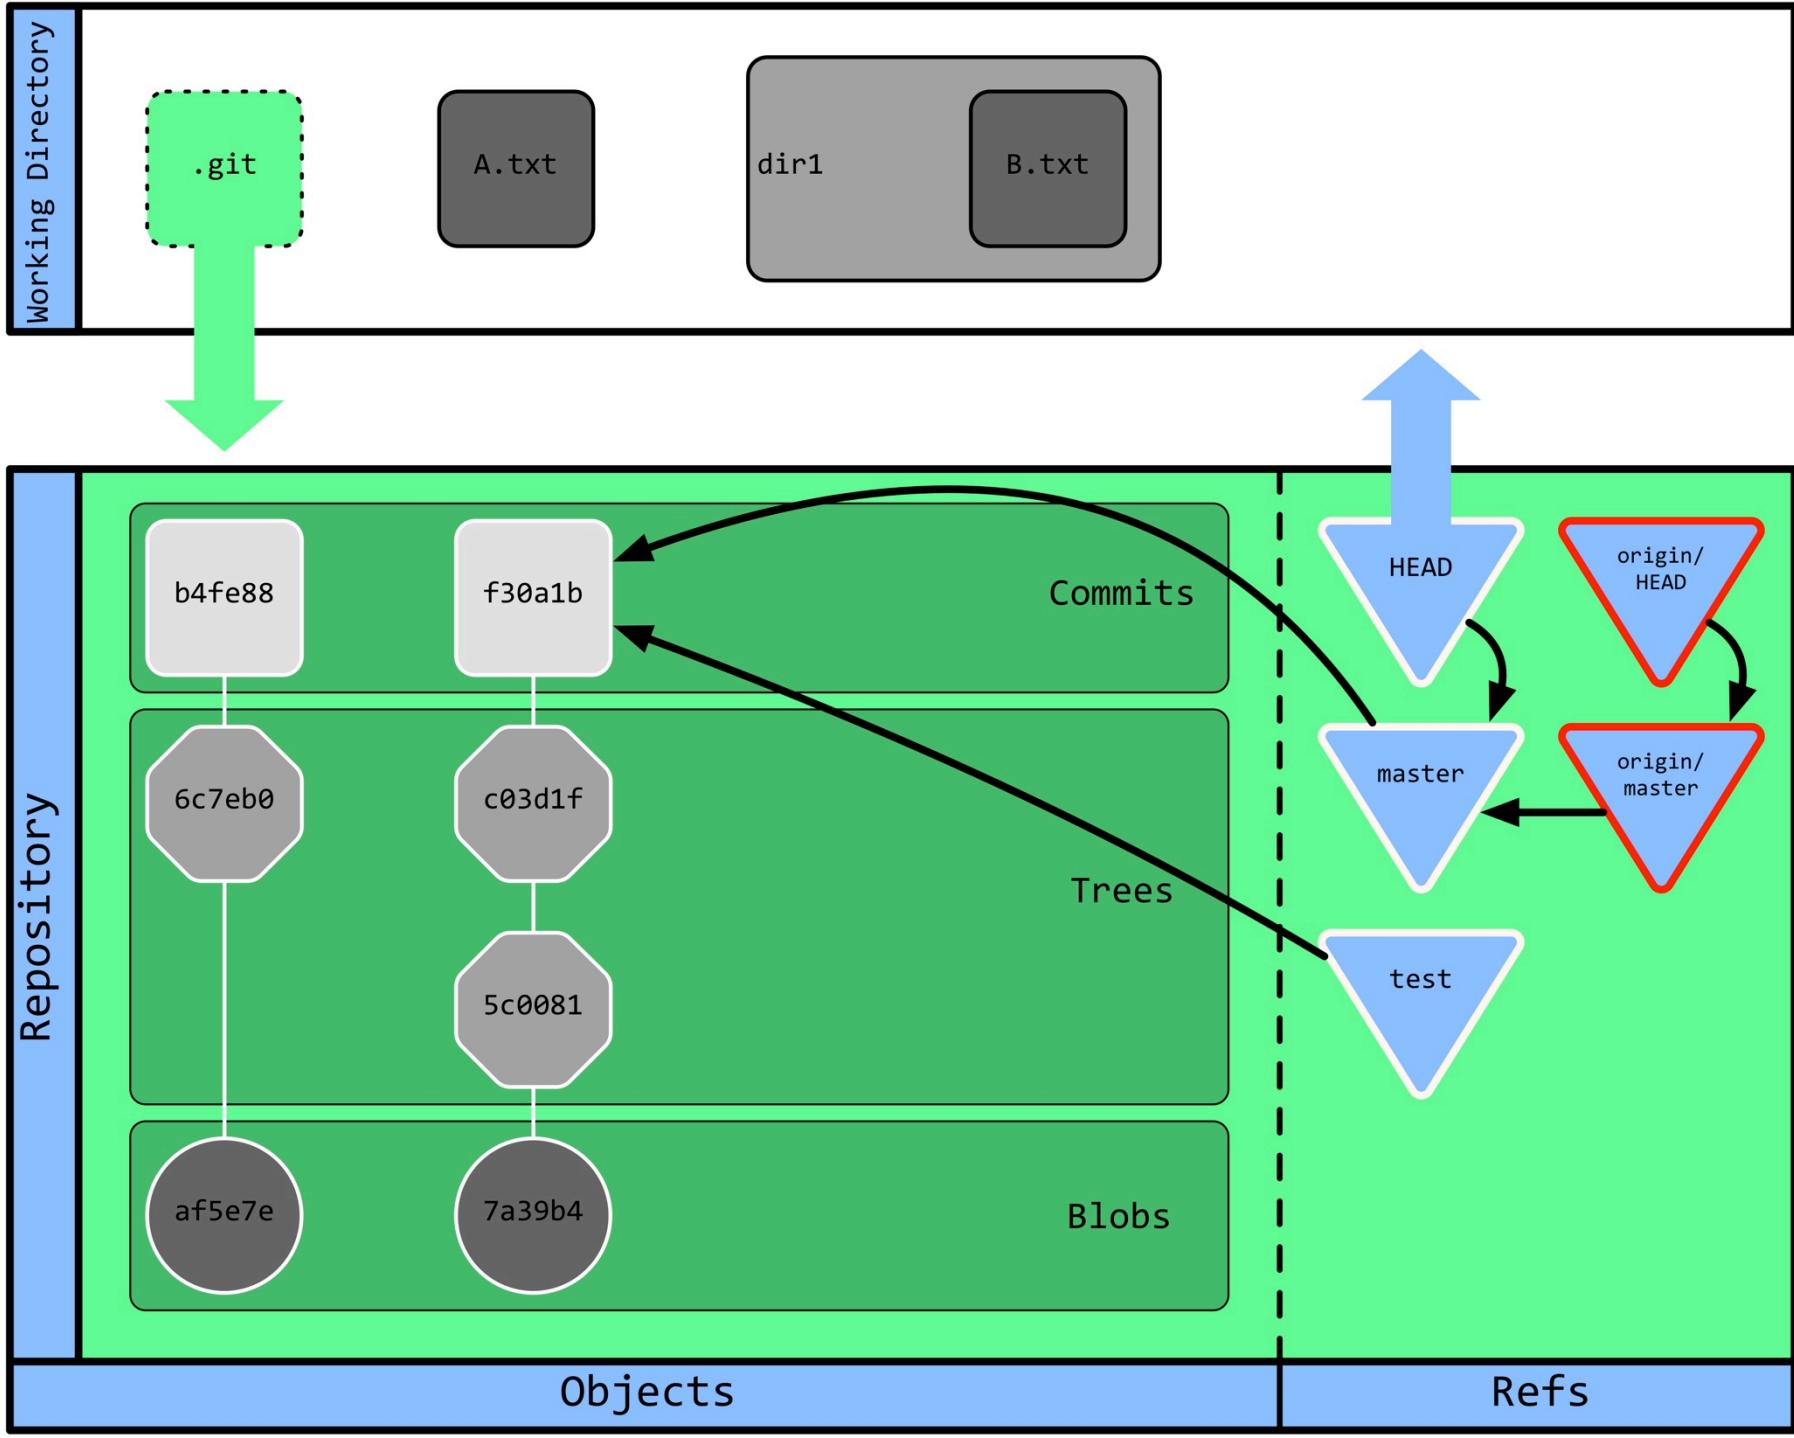
Easy, right?

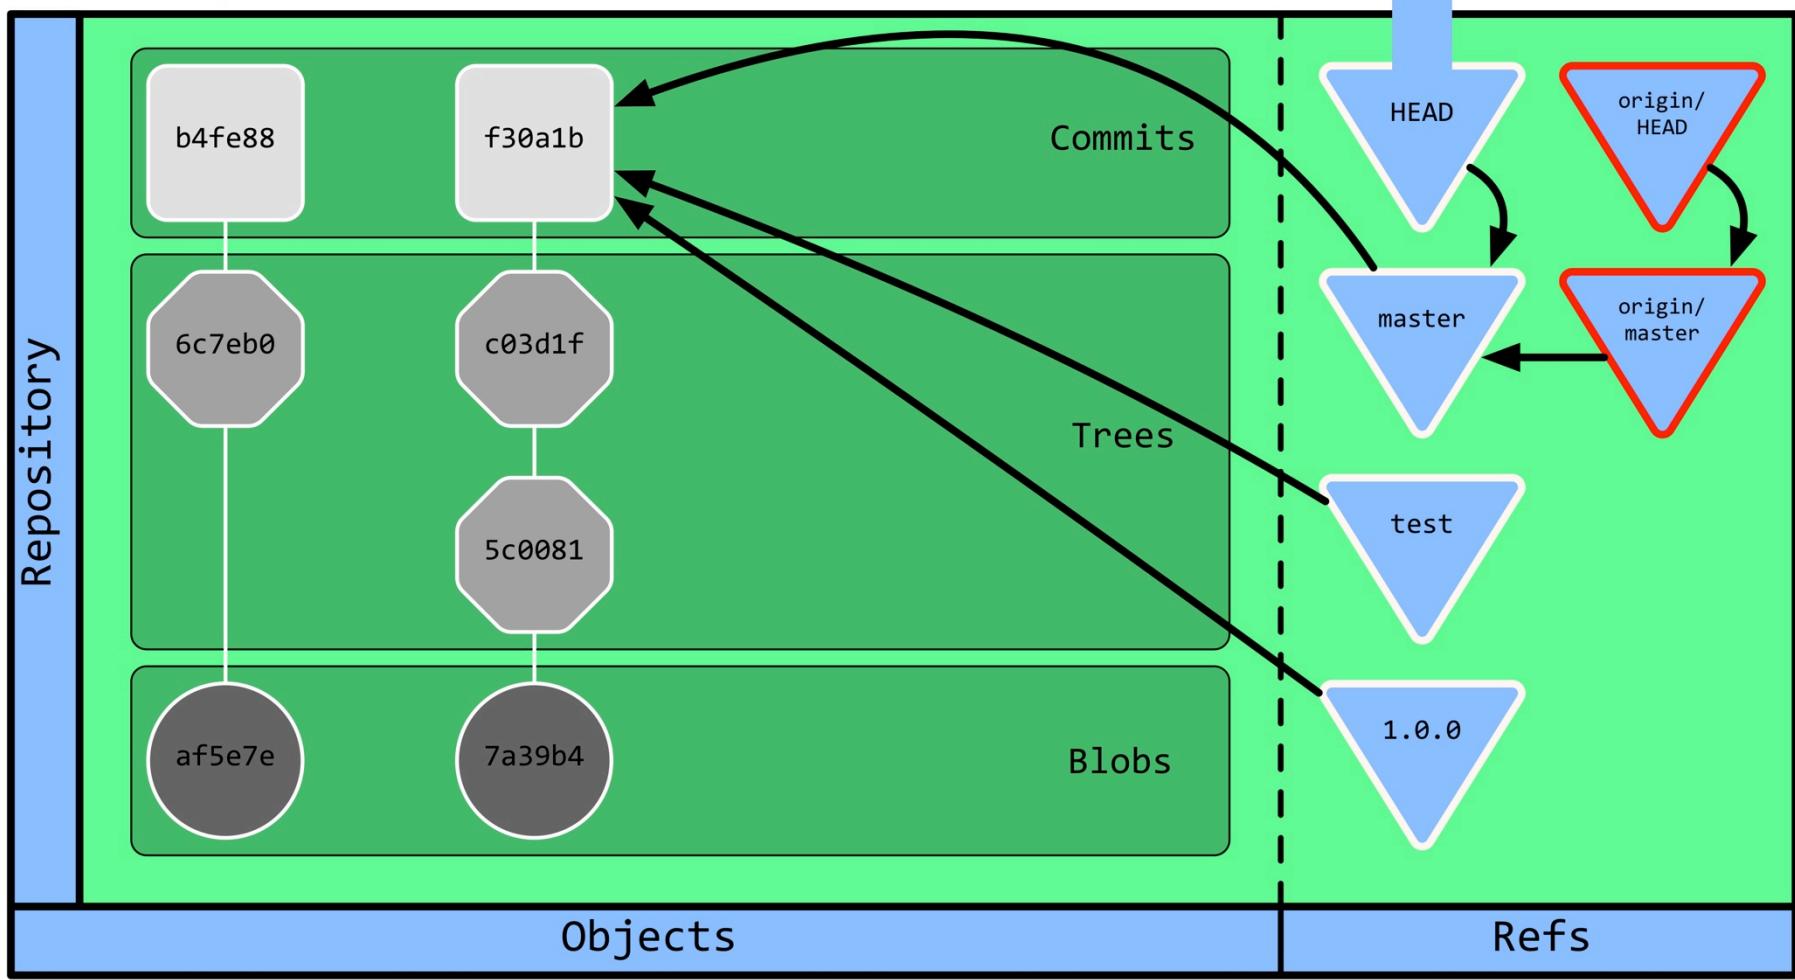
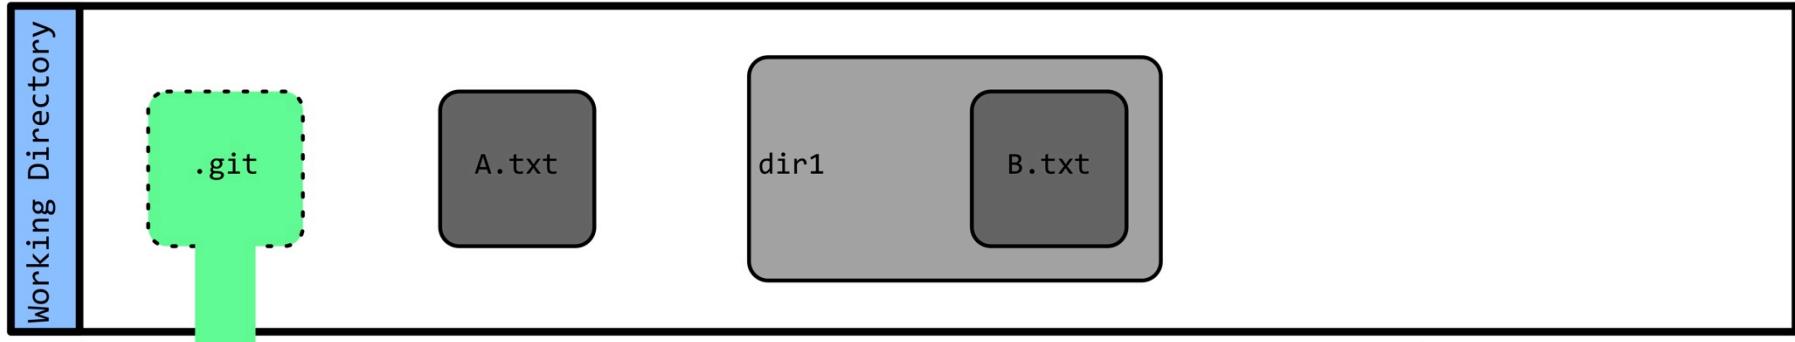
That was called a “fast-forward” merge.

Other merges can require an additional commit, but we’ll stick with the simple case.

We're happy with the state of the repo, so
let's tag it as 1.0.0

git tag 1.0.0





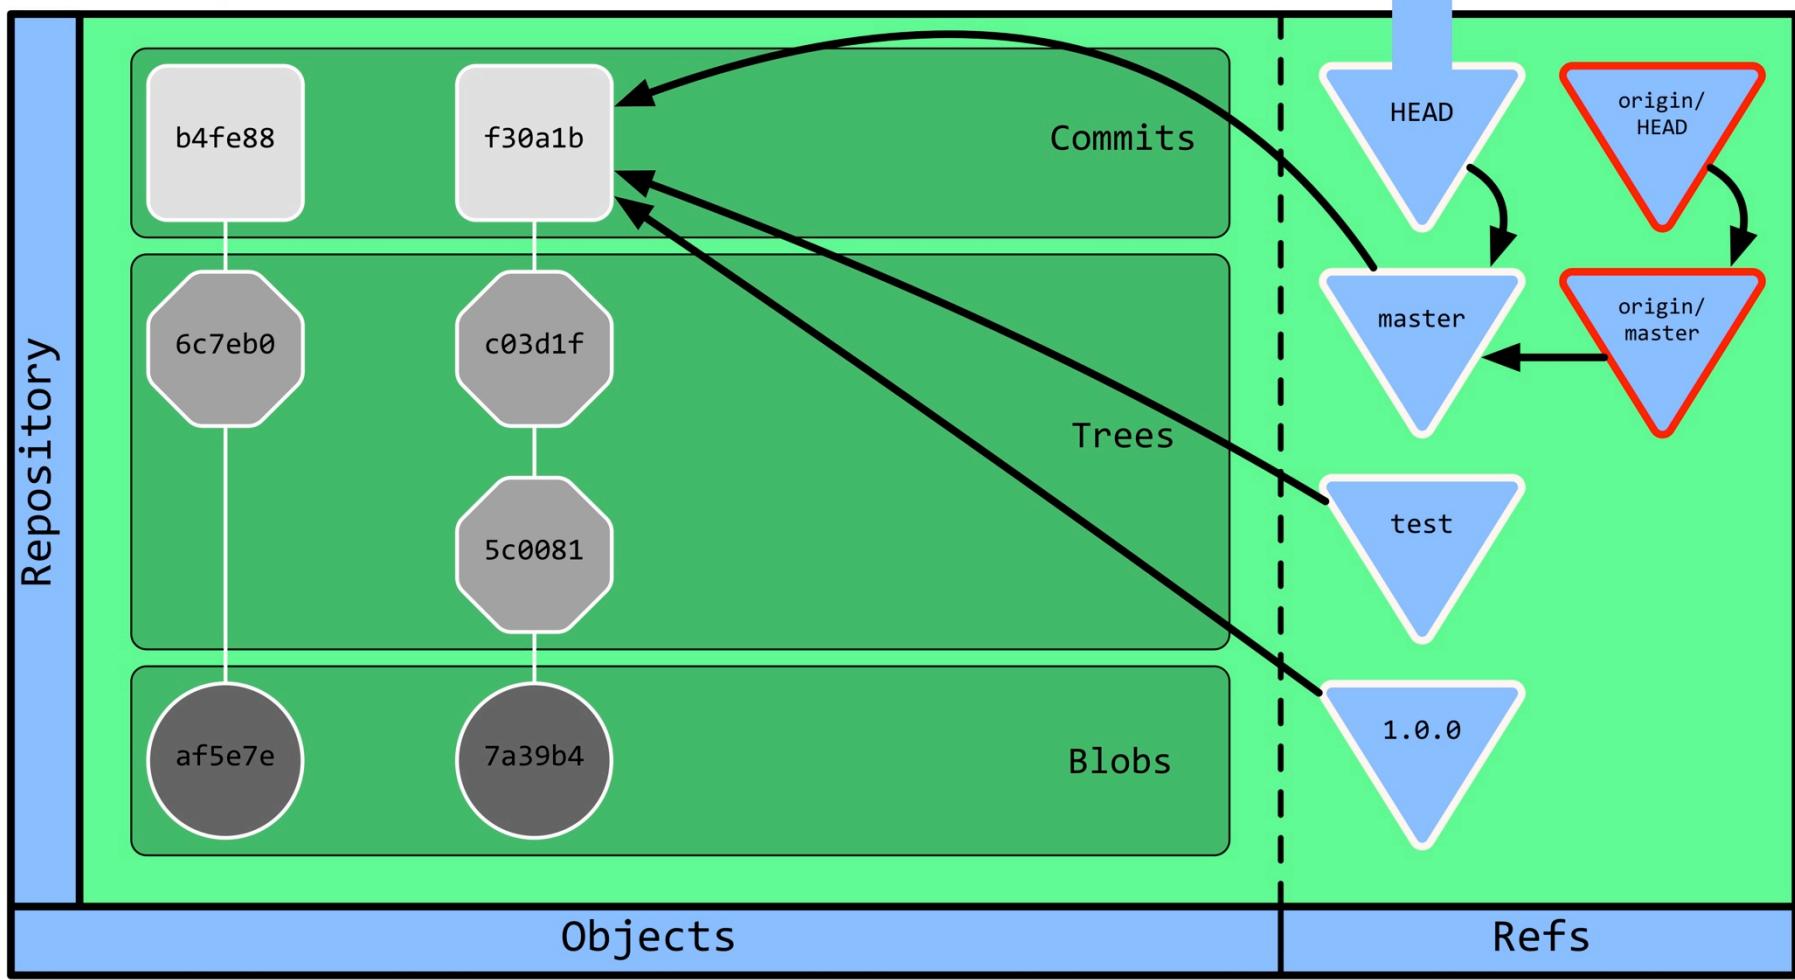
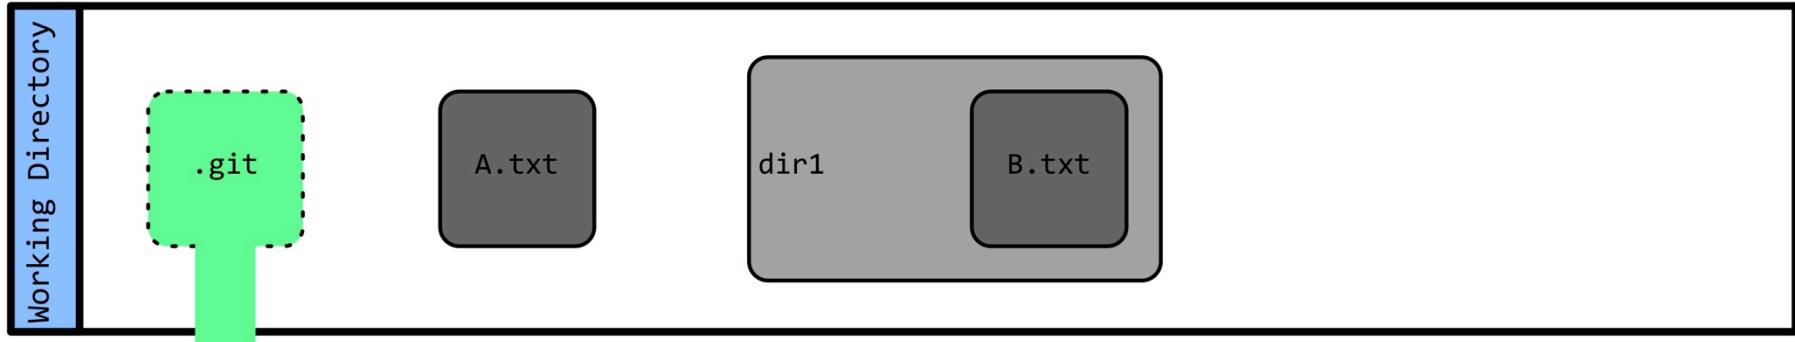
Yes, it's really that simple.

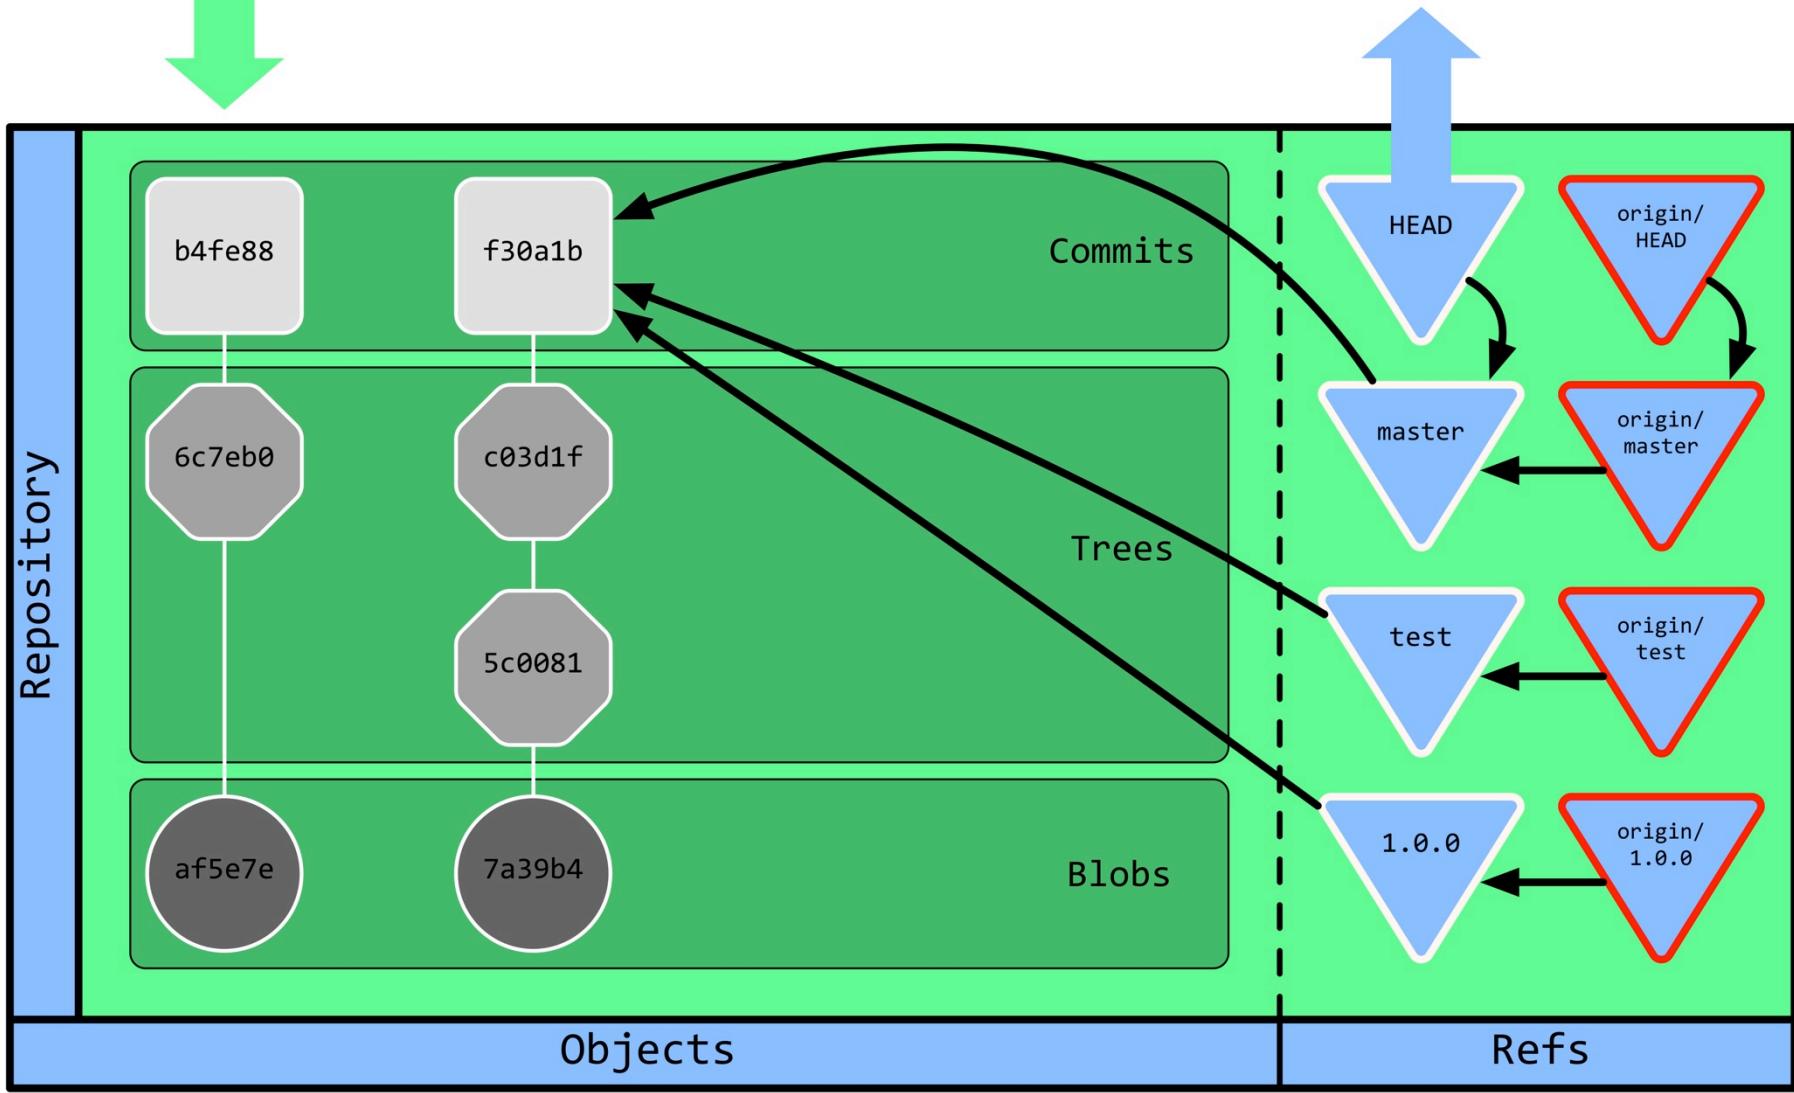
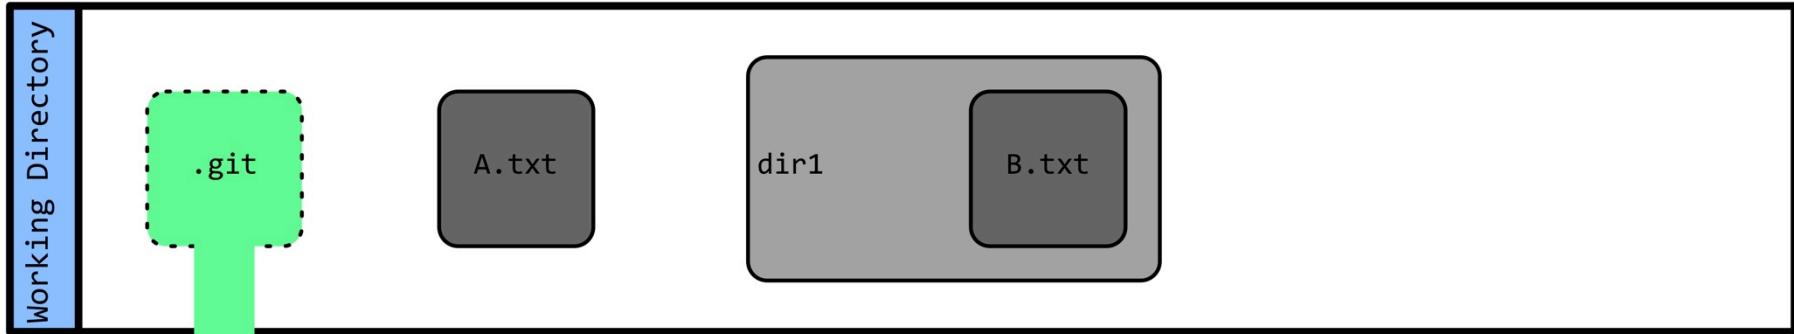
We'll be a team player and push our changes to the remote

```
git push origin master 1.0.0
```

Explicitly specify tag,
or use the **--tags** option

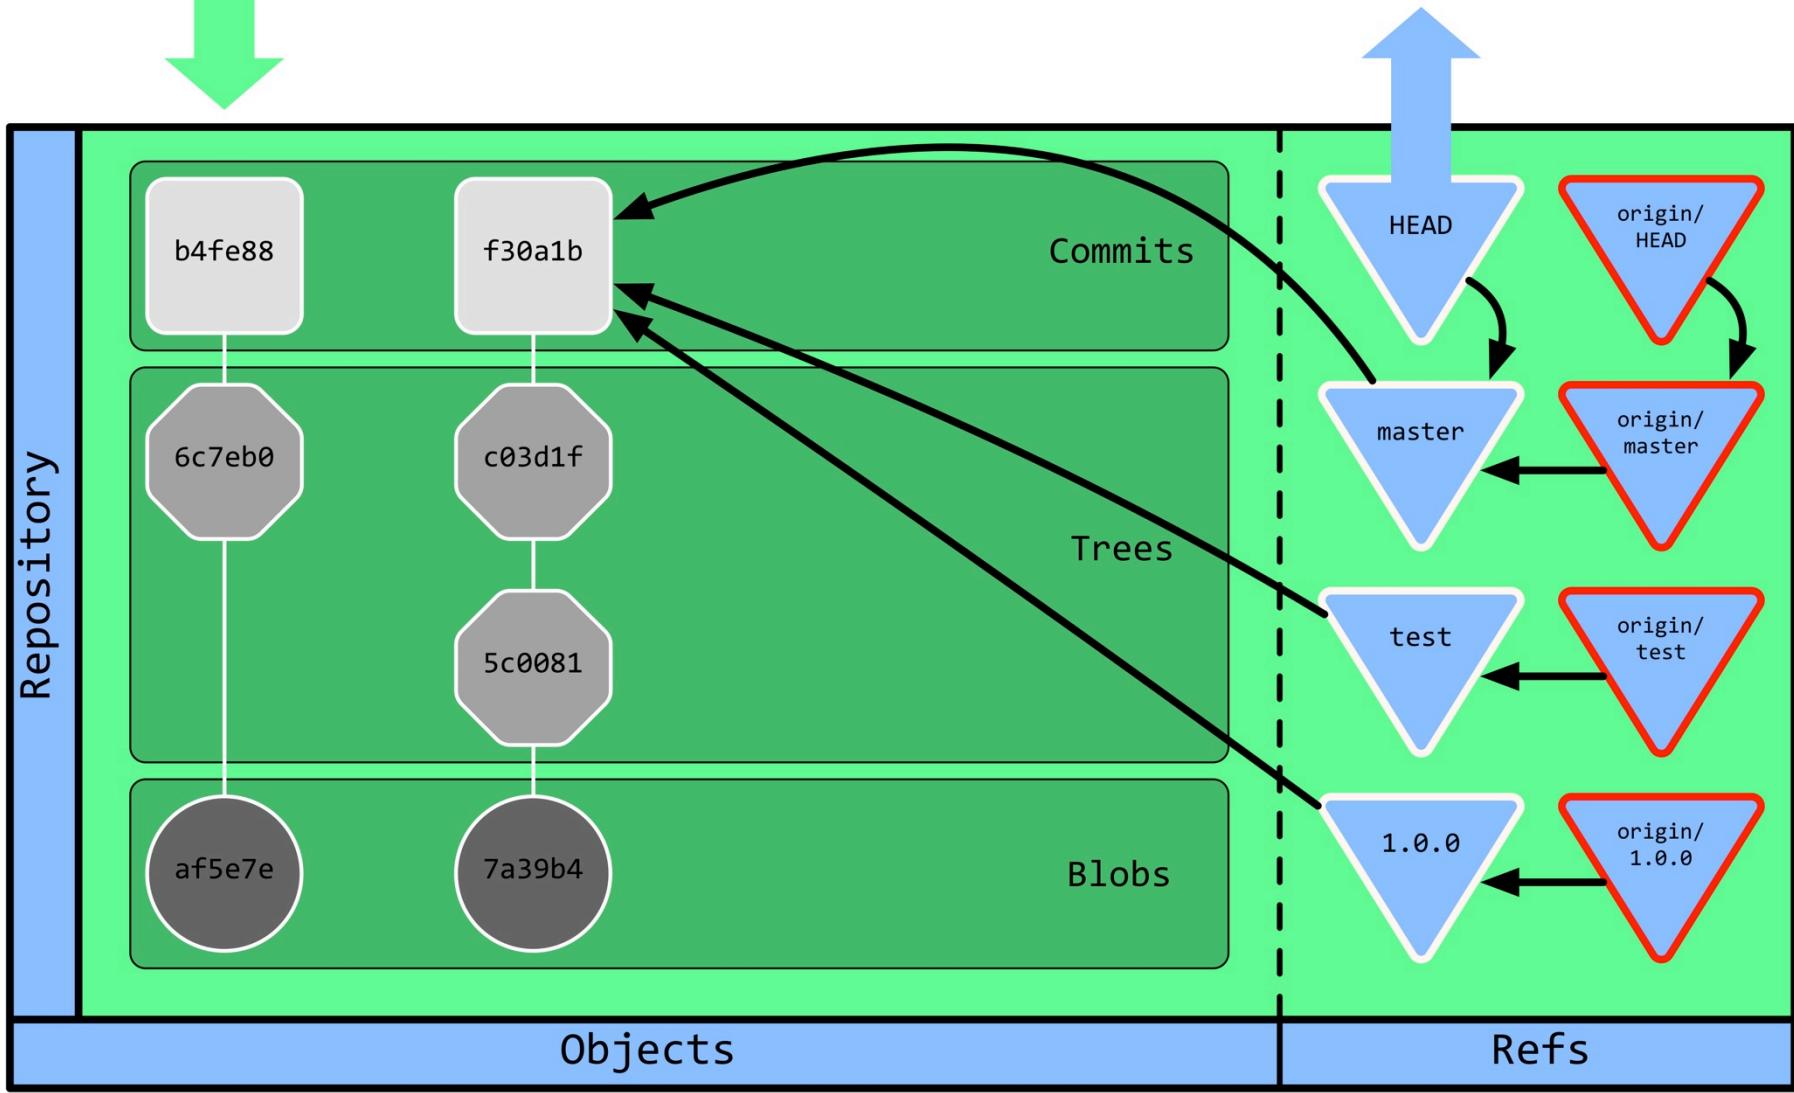
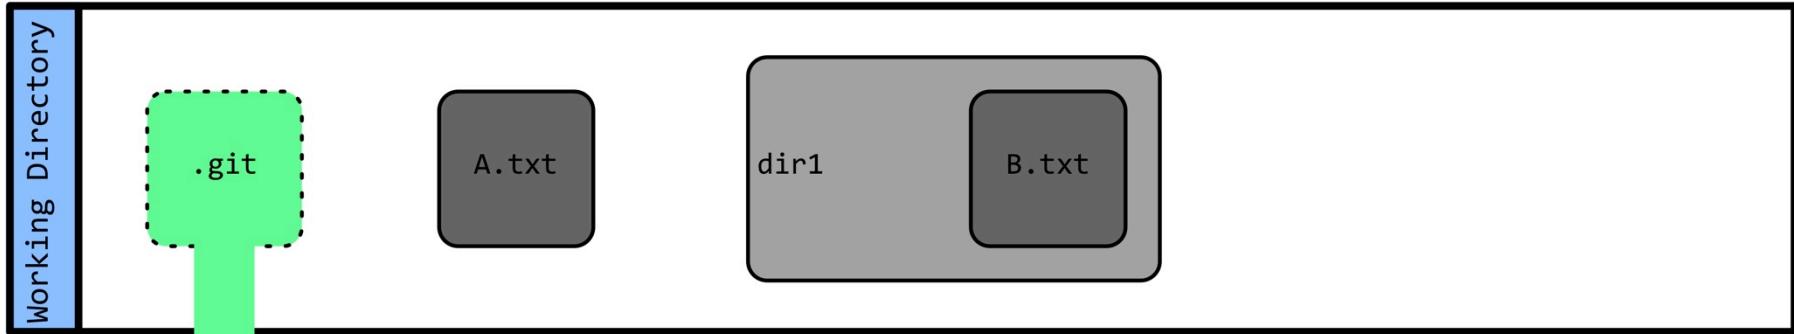


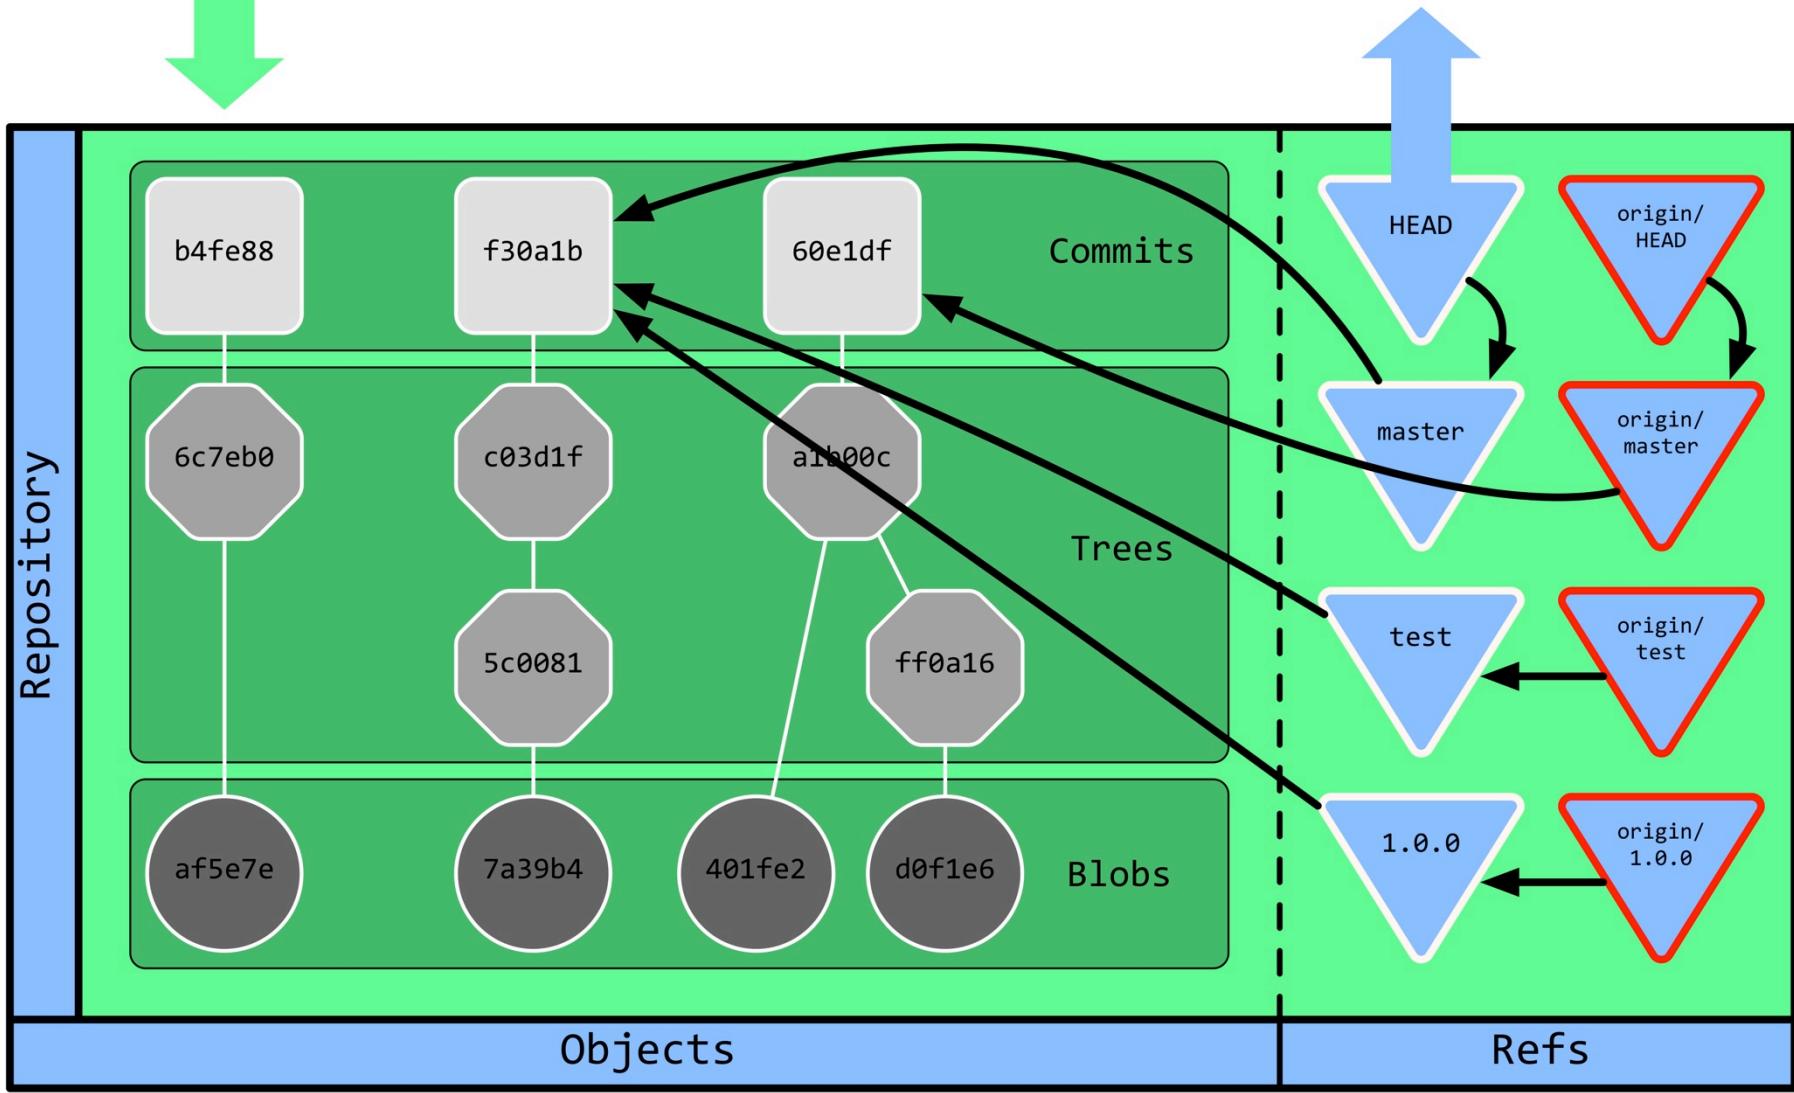
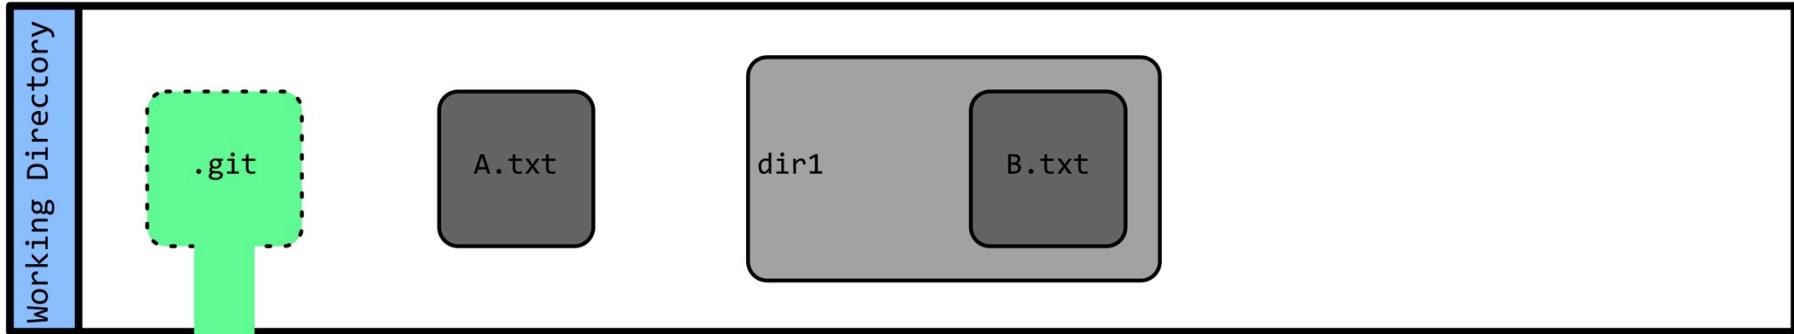




After some time, we can fetch from the remote to grab any changes others have made.

git fetch

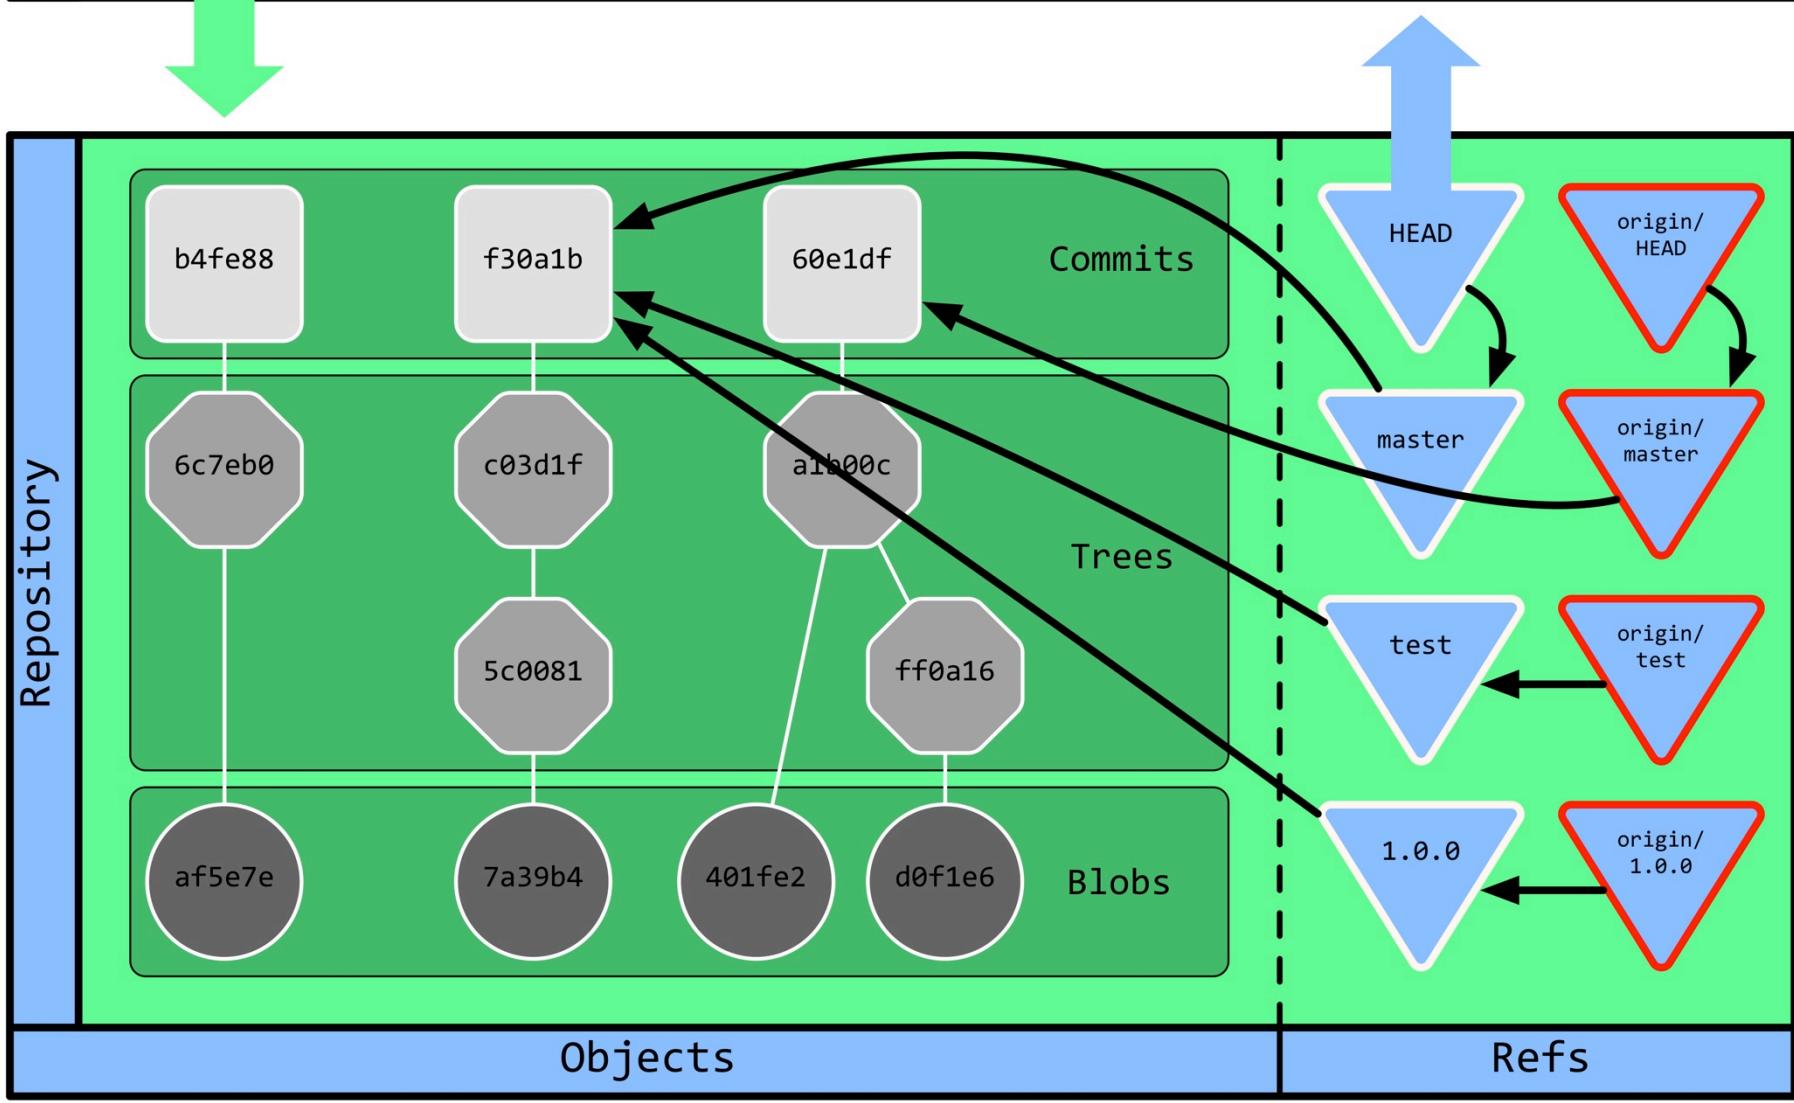
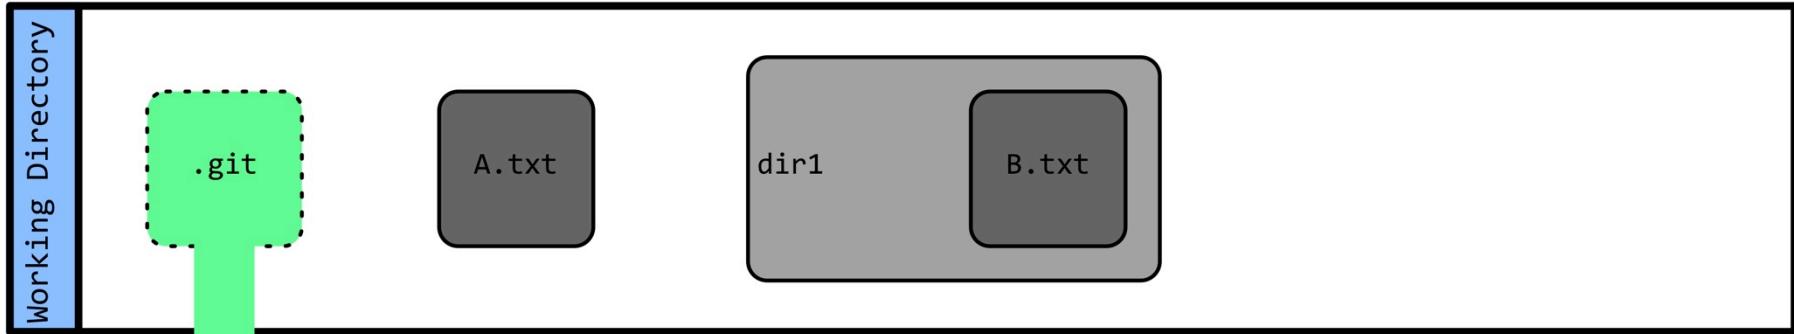


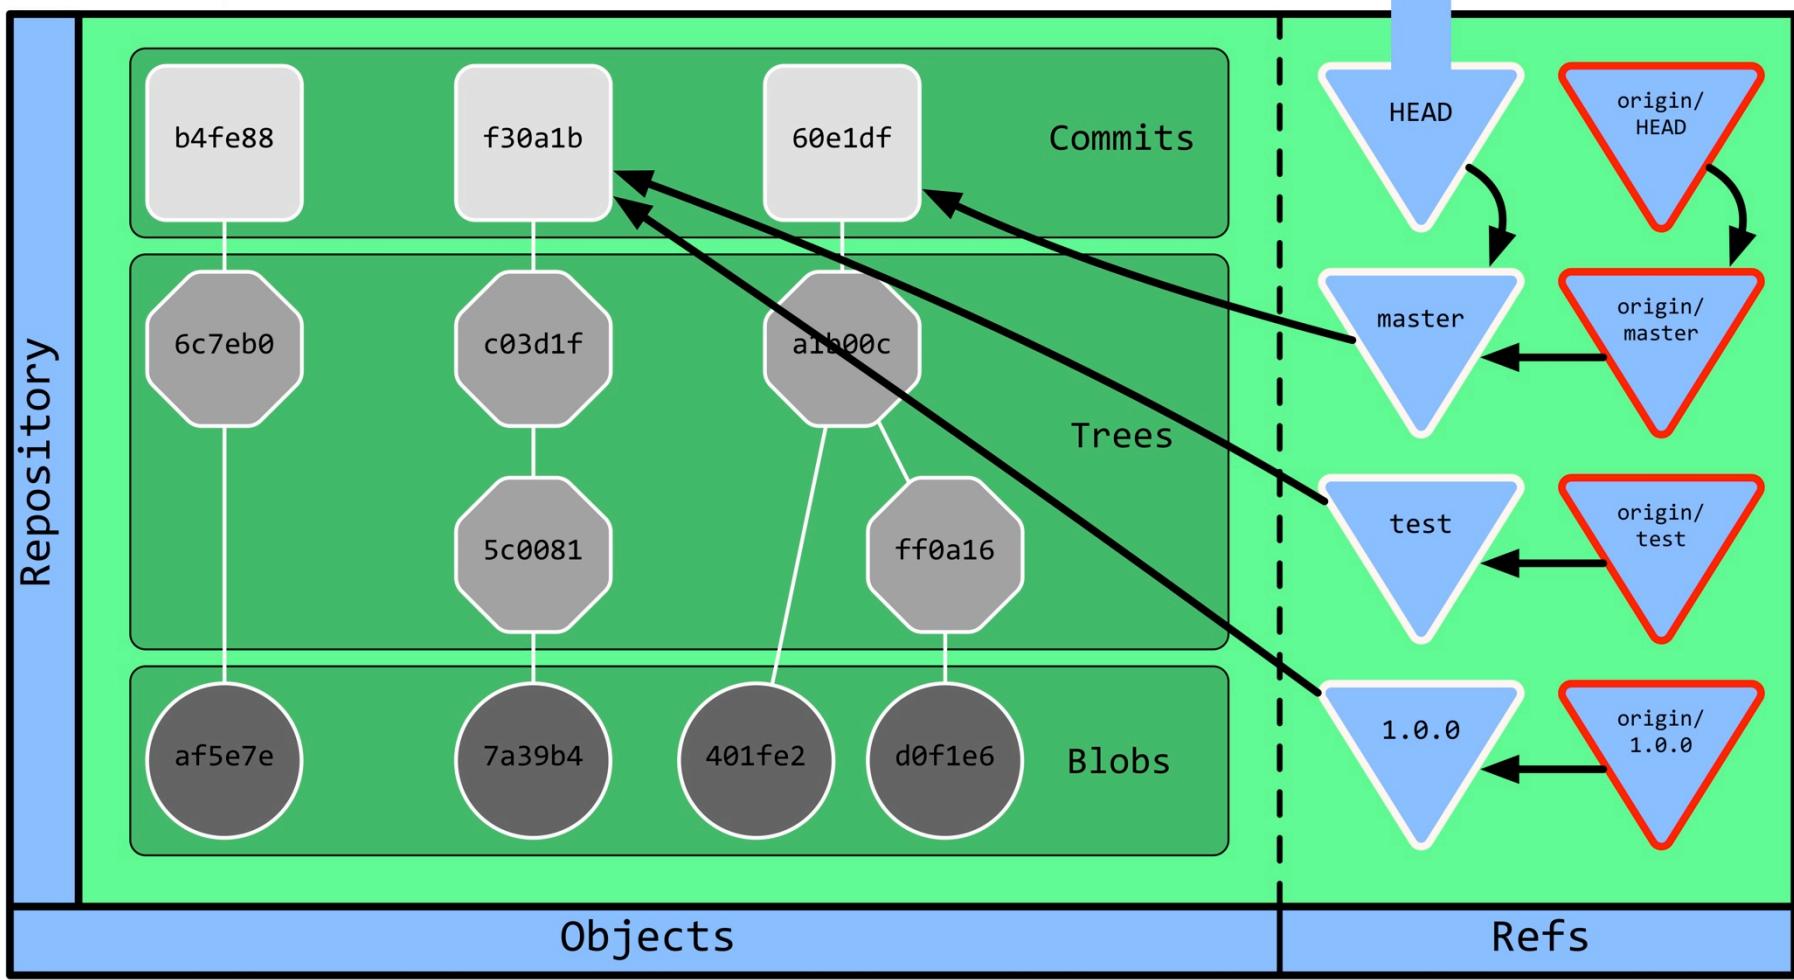
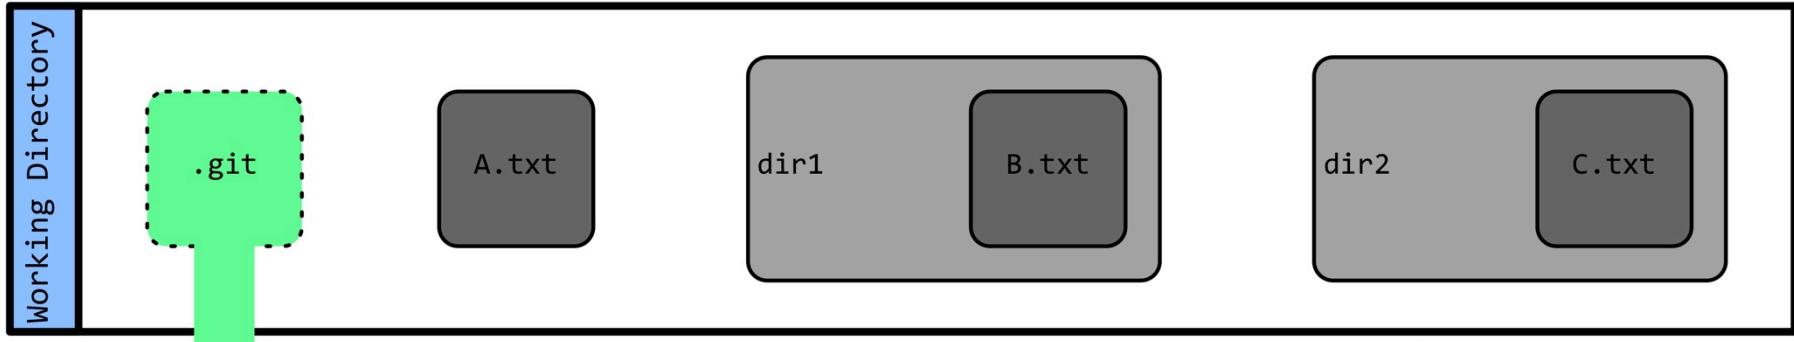


**It looks like someone made a commit,
adding dir2 and C.txt, and modified
the contents of A.txt**

To actually update our files, we need to merge origin/master into our local master

git merge origin/master



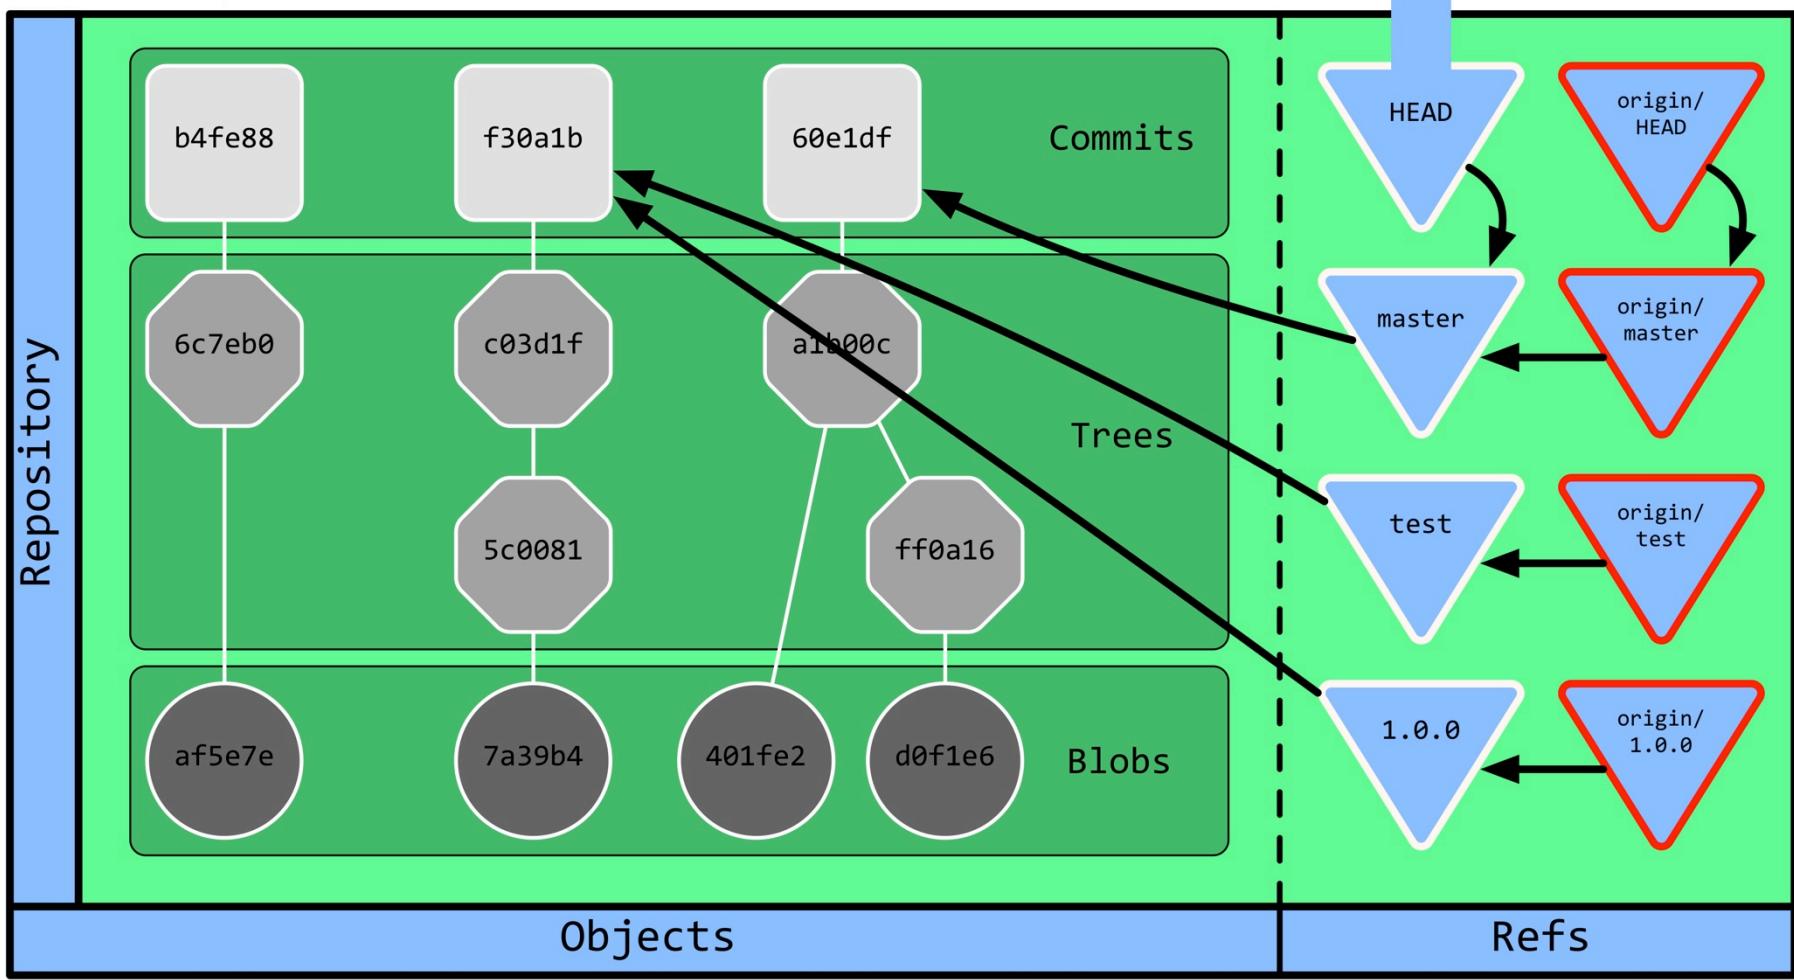
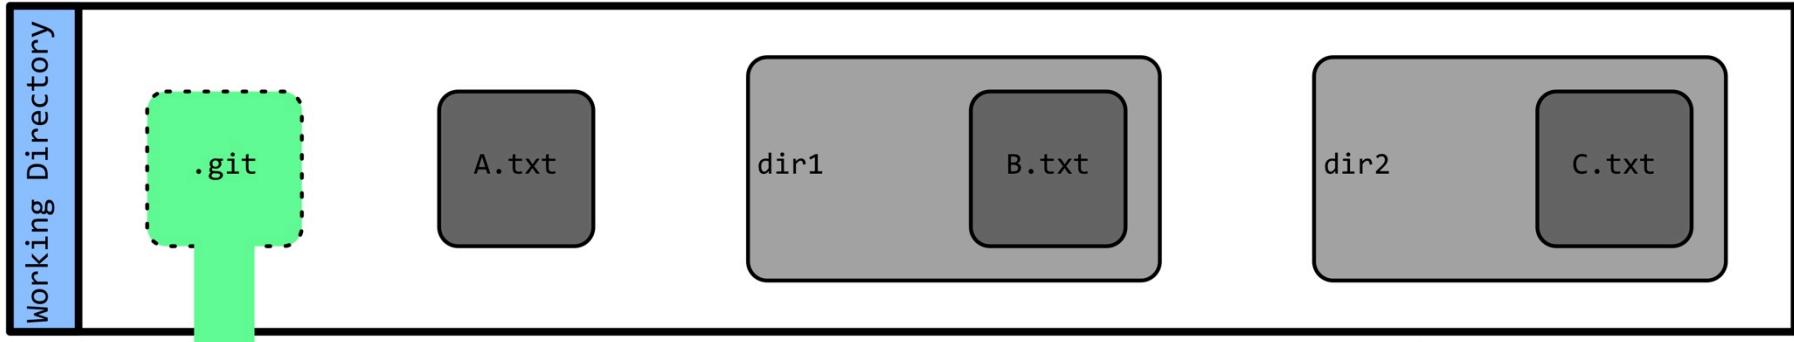


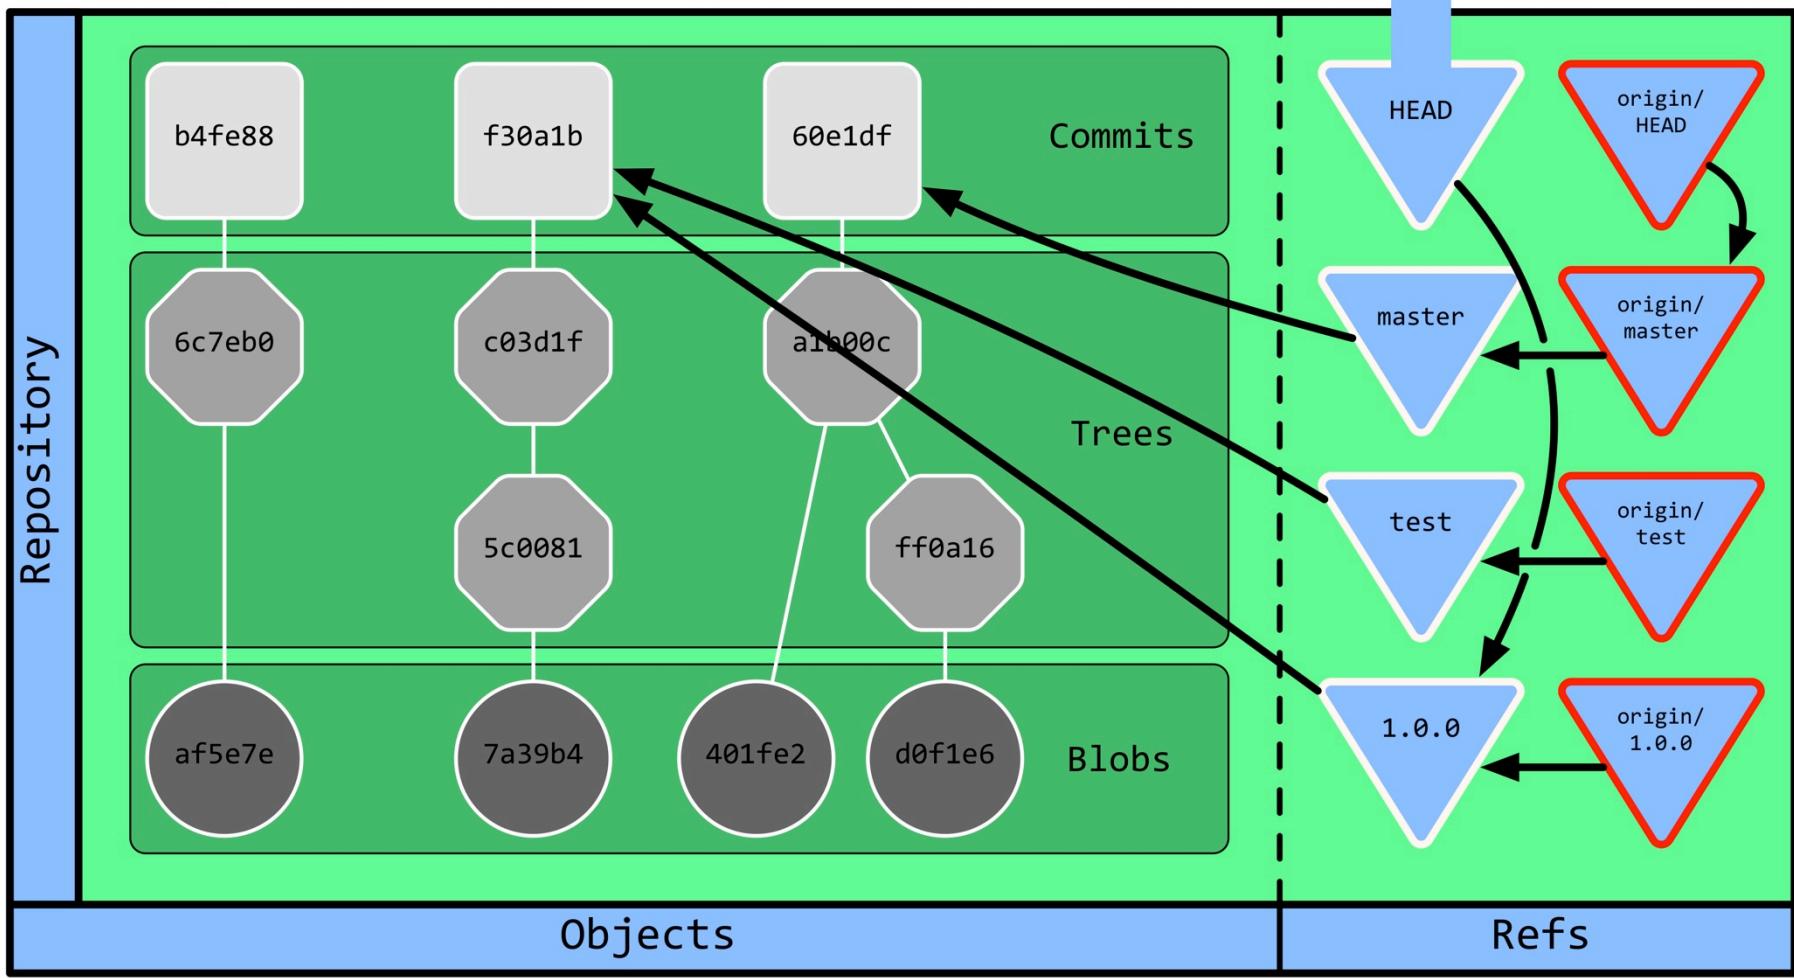
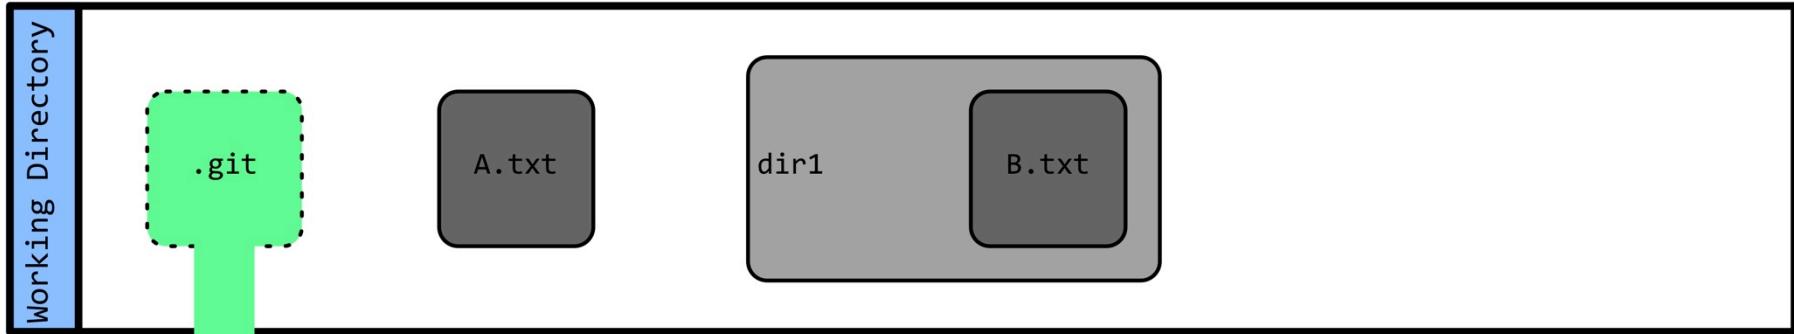
As a side note:

git pull is just a **fetch** followed by a
merge

Everything looks good, but let's see what happens when we want to switch to a previous commit, such as tag 1.0.0

git checkout 1.0.0





A Few Things to Note

- All the new refs, commits, trees, and blobs still exist, even though we “went back” to a previous commit
- **C.txt** and **dir2** **STILL EXIST!** They are blob **d0f1e6** and tree **ff0a16**, respectively
- To switch back to the latest commit, we can just type **git checkout master**

A Few Things to Note

- The `checkout` command did two things:
 - 1. `HEAD` pointed to the ref `1.0.0`
 - 2. The working directory (the files you see) were replaced with ones from the blobs referenced by the commit/trees that `1.0.0` points to
- Typing any of the following three commands would have done the same thing:
 - `git checkout test`
 - `git checkout f30a1b` # <-- detached HEAD!
 - `git checkout HEAD^` # <-- detached HEAD!
- That last one is shorthand for “`HEAD`’s parent”

Detached HEAD State

- At some point you may see this message:

```
$ git checkout HEAD^  
Note: checking out 'HEAD^'.
```

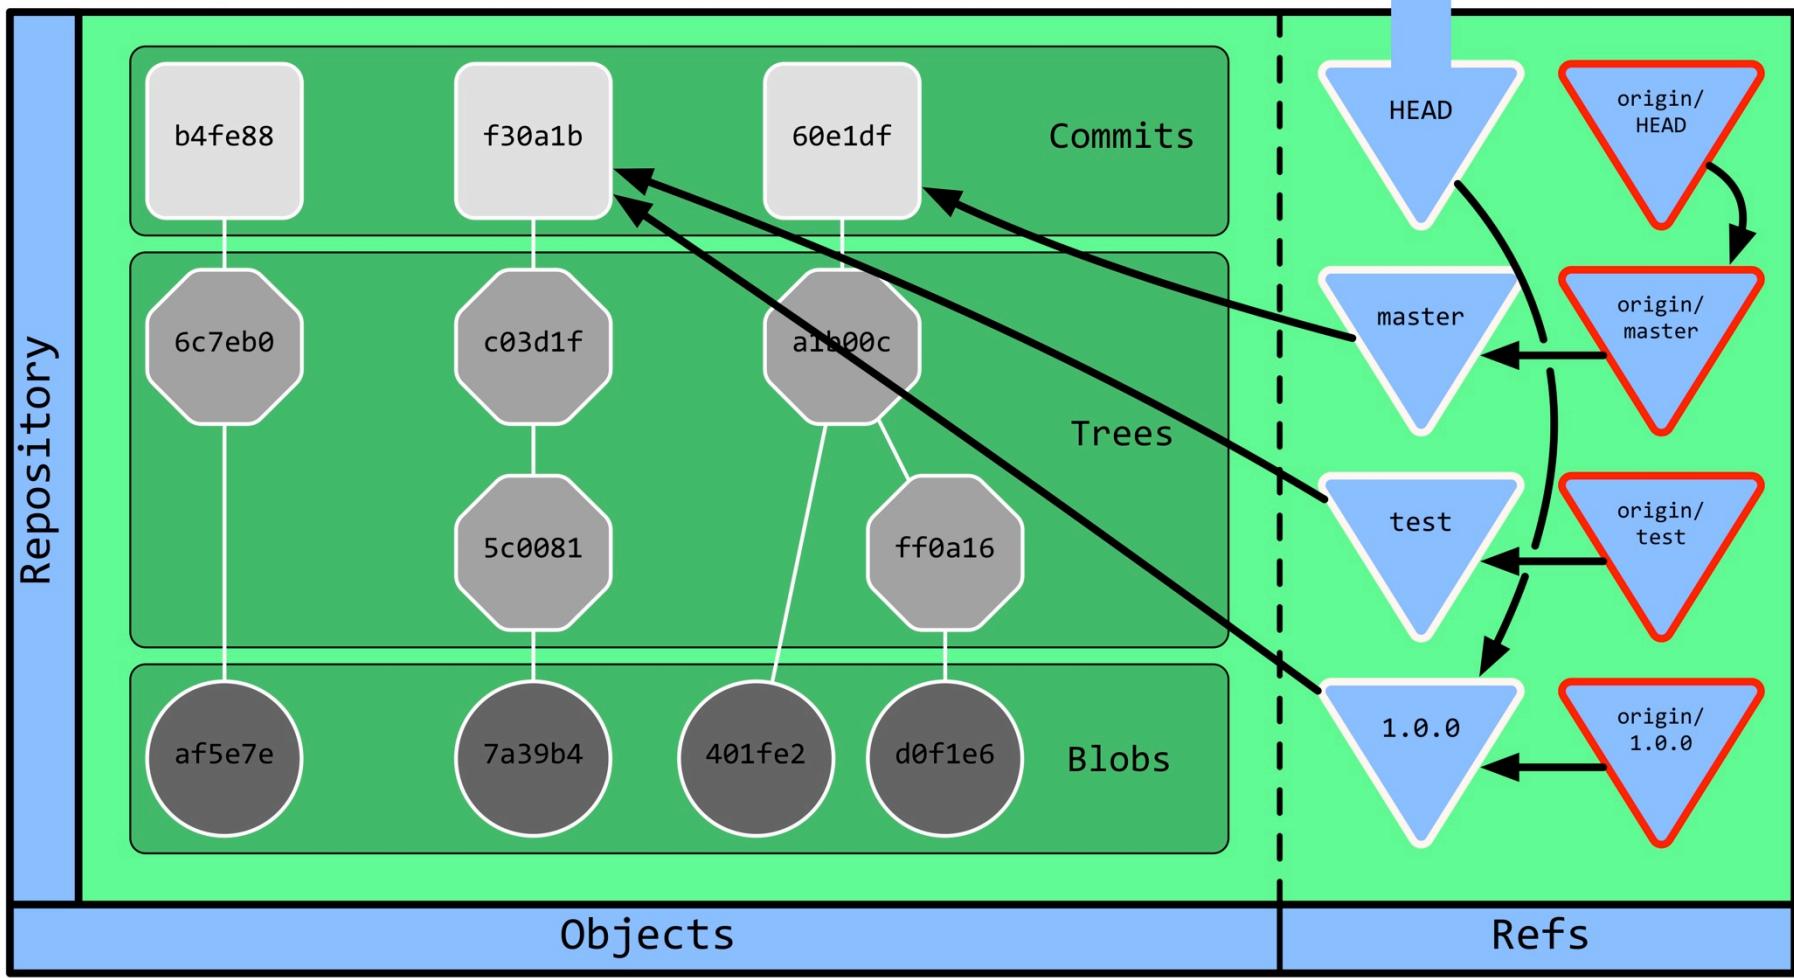
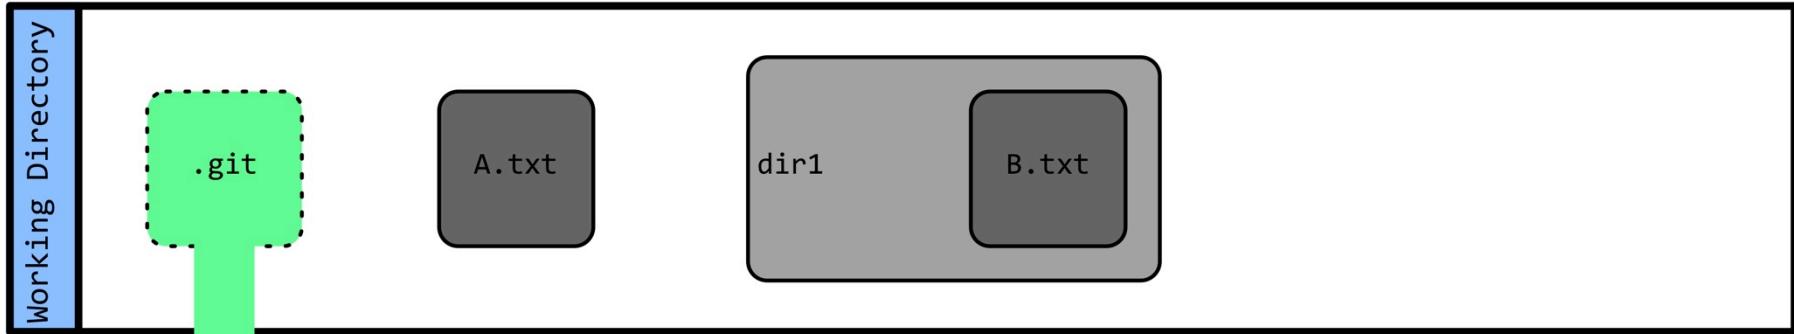
You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by performing another checkout.

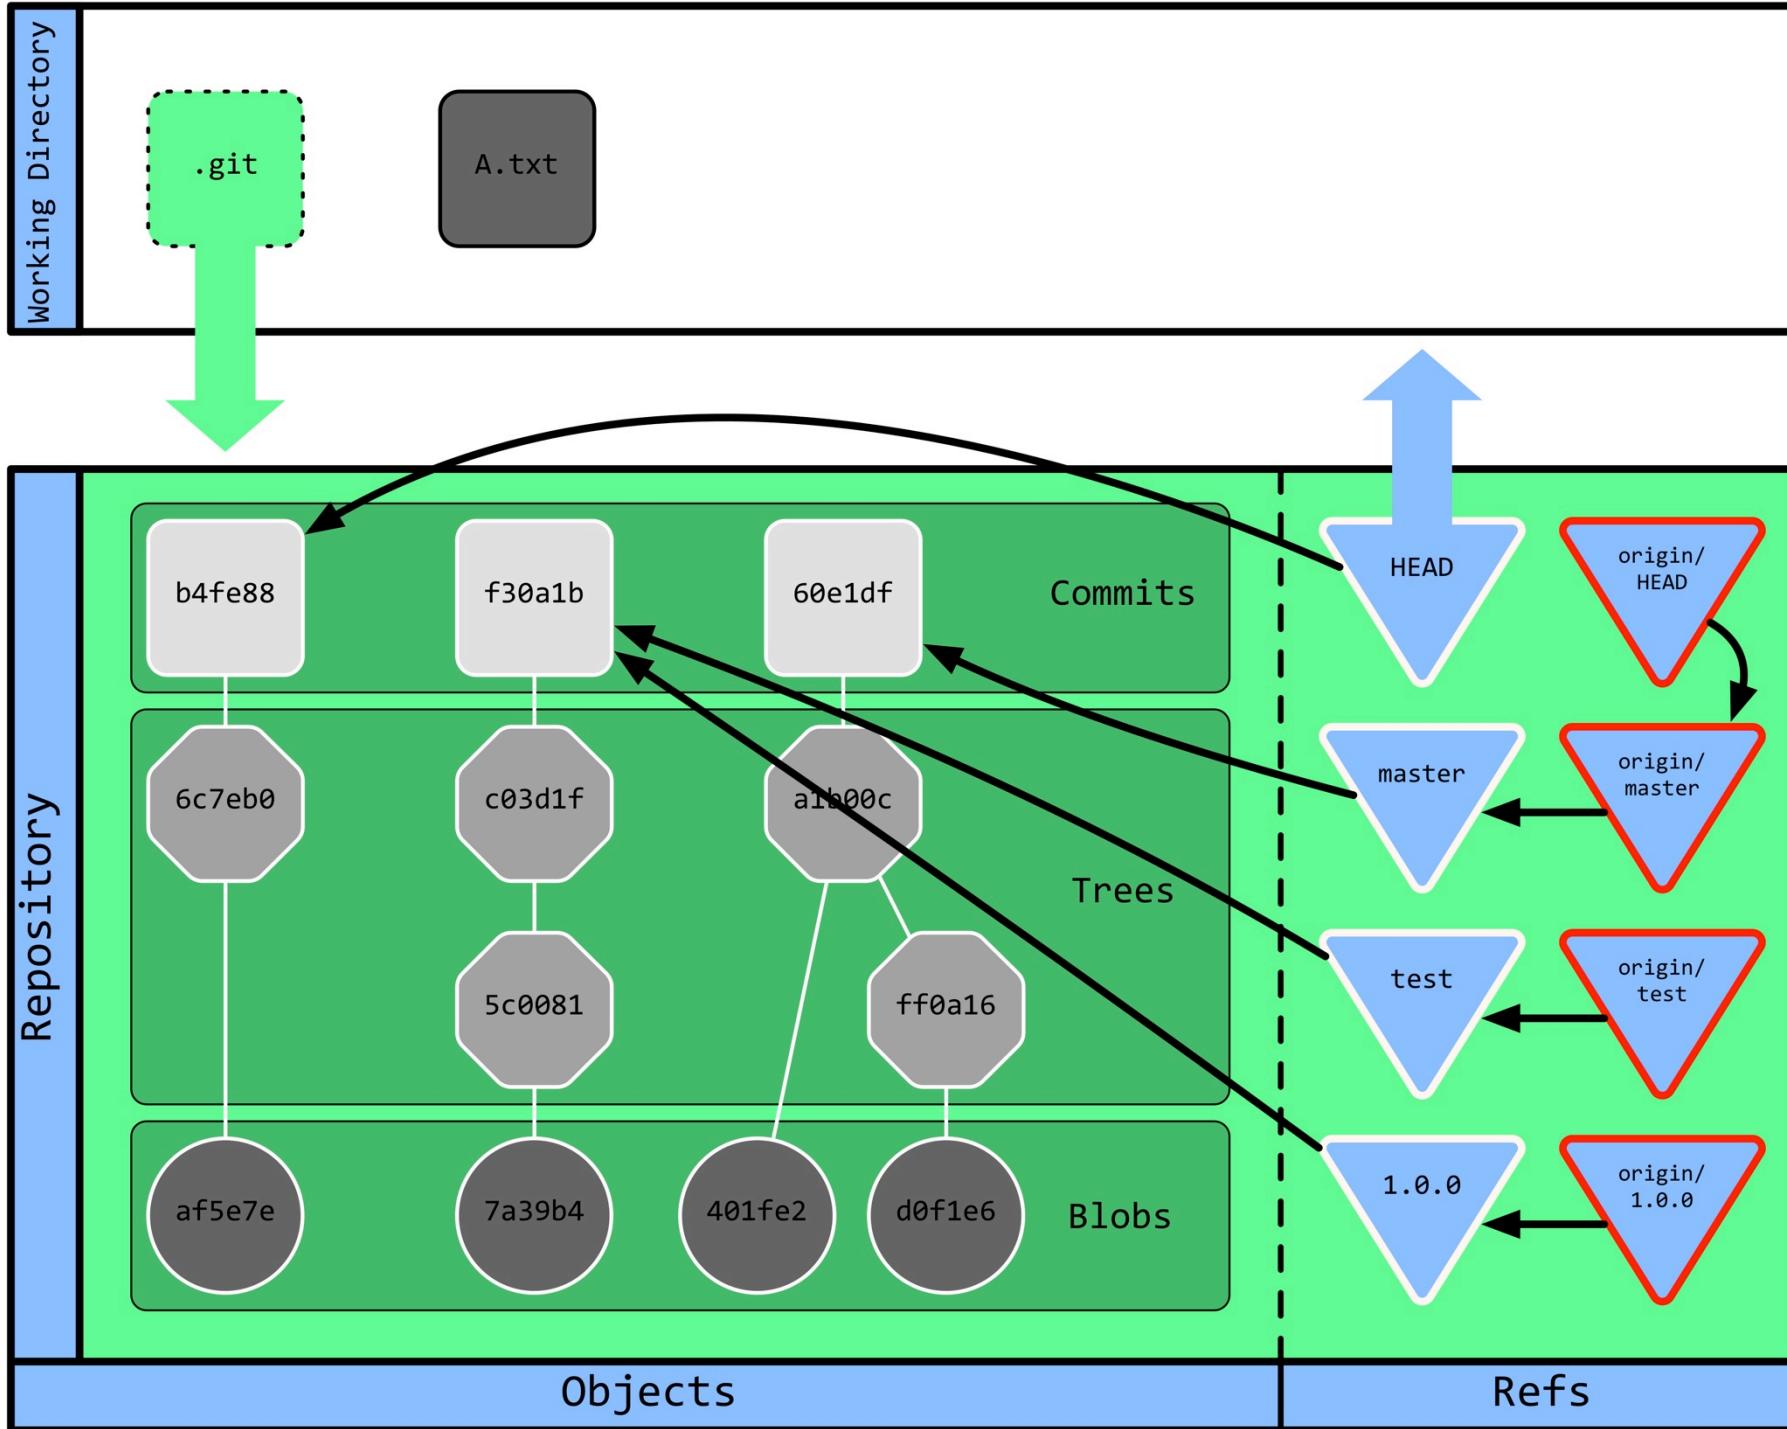
If you want to create a new branch to retain commits you create, you may do so (now or later) by using -b with the checkout command again. Example:

```
git checkout -b <new-branch-name>  
  
HEAD is now at 3e6d961... Add file B  
  
$ cat .git/HEAD  
3e6d9614ec3fc76e440aca6dafa8e6e0c14f34a
```

- A detached HEAD state is when HEAD points directly to a commit instead of a branch ref.

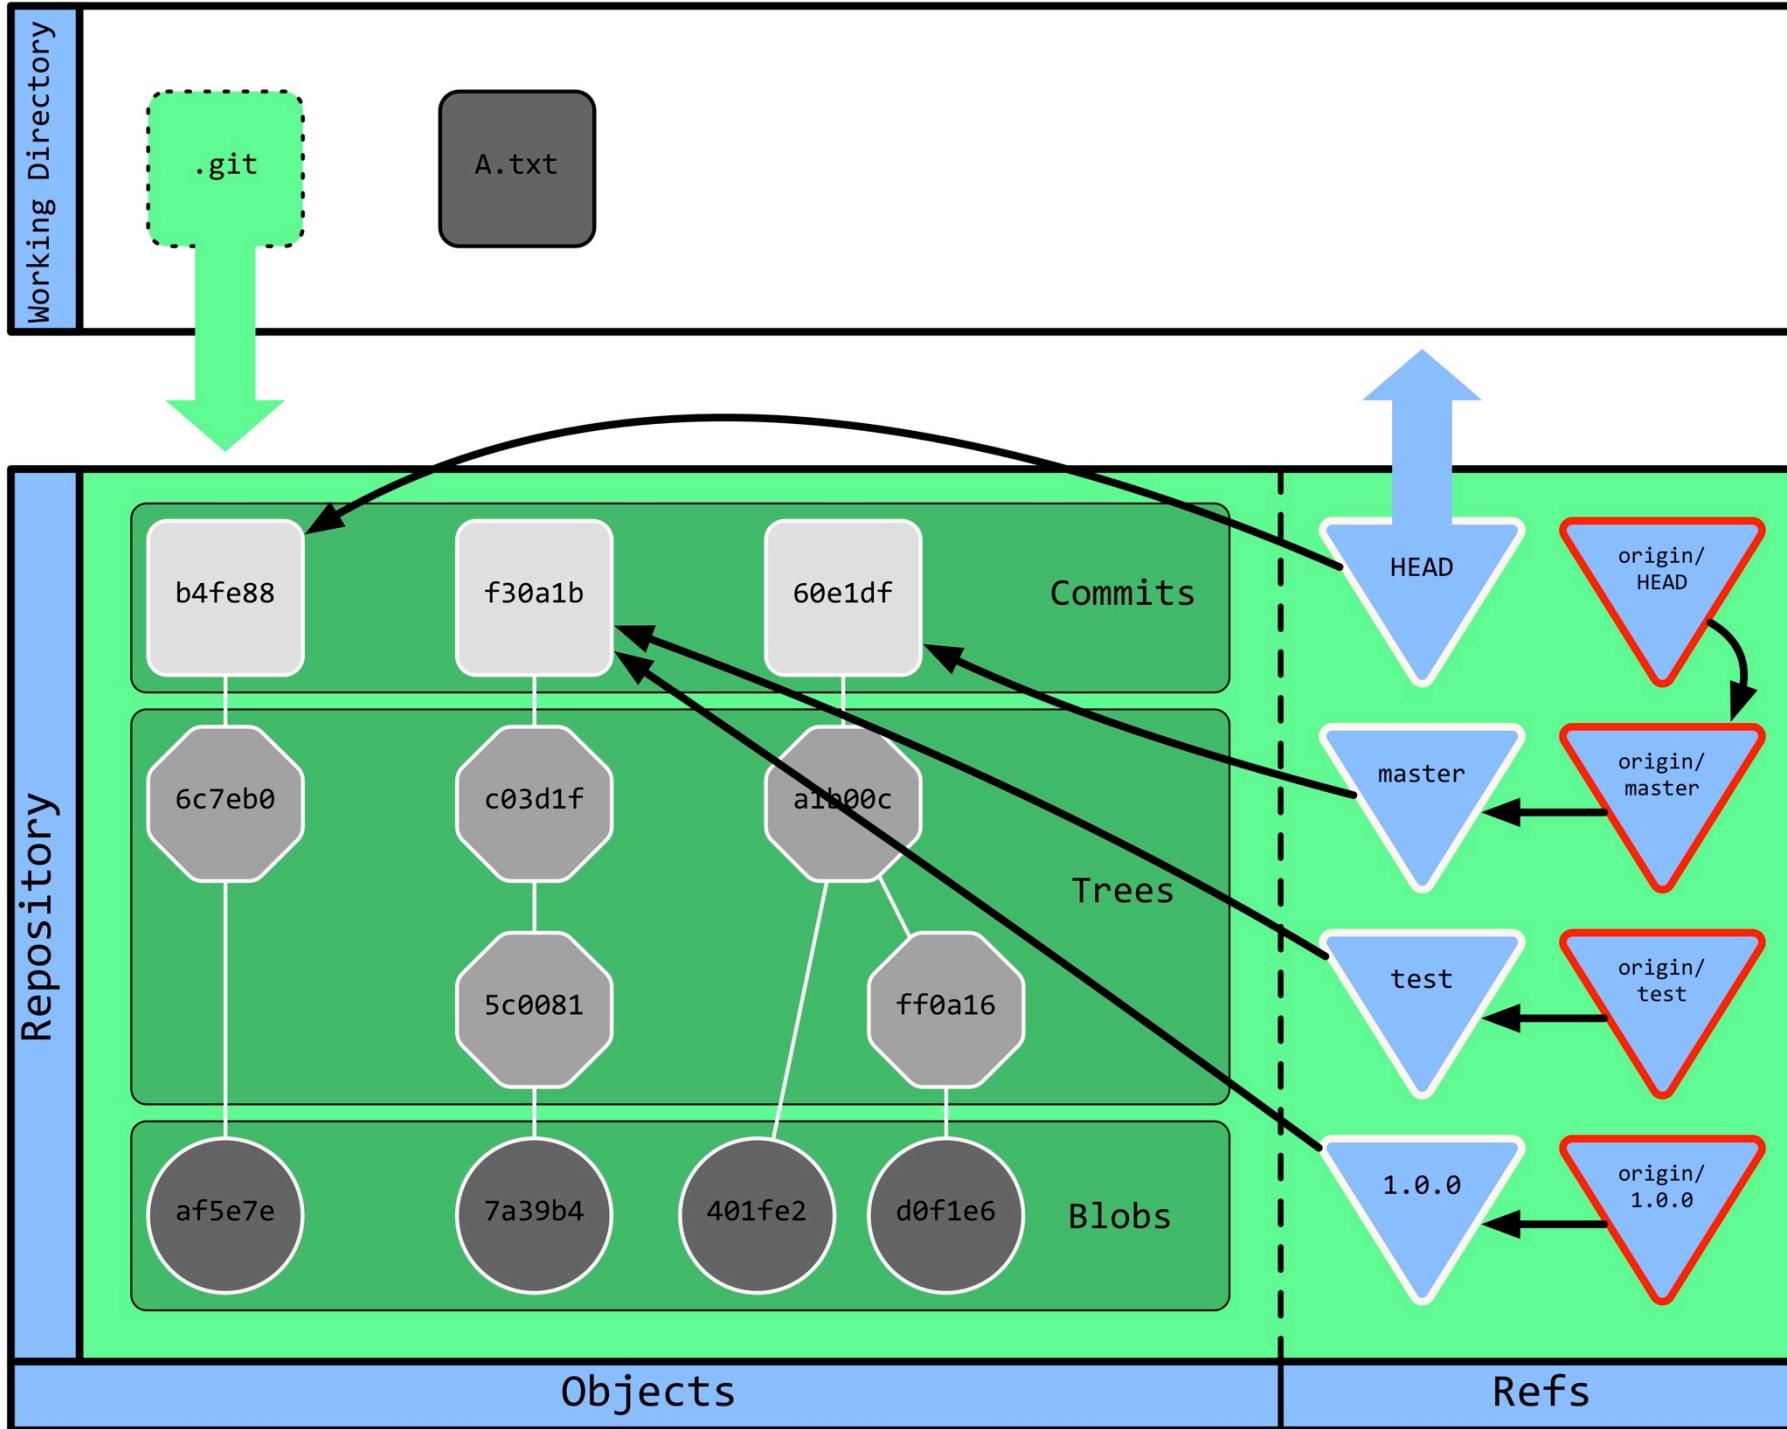
git checkout b4fe88

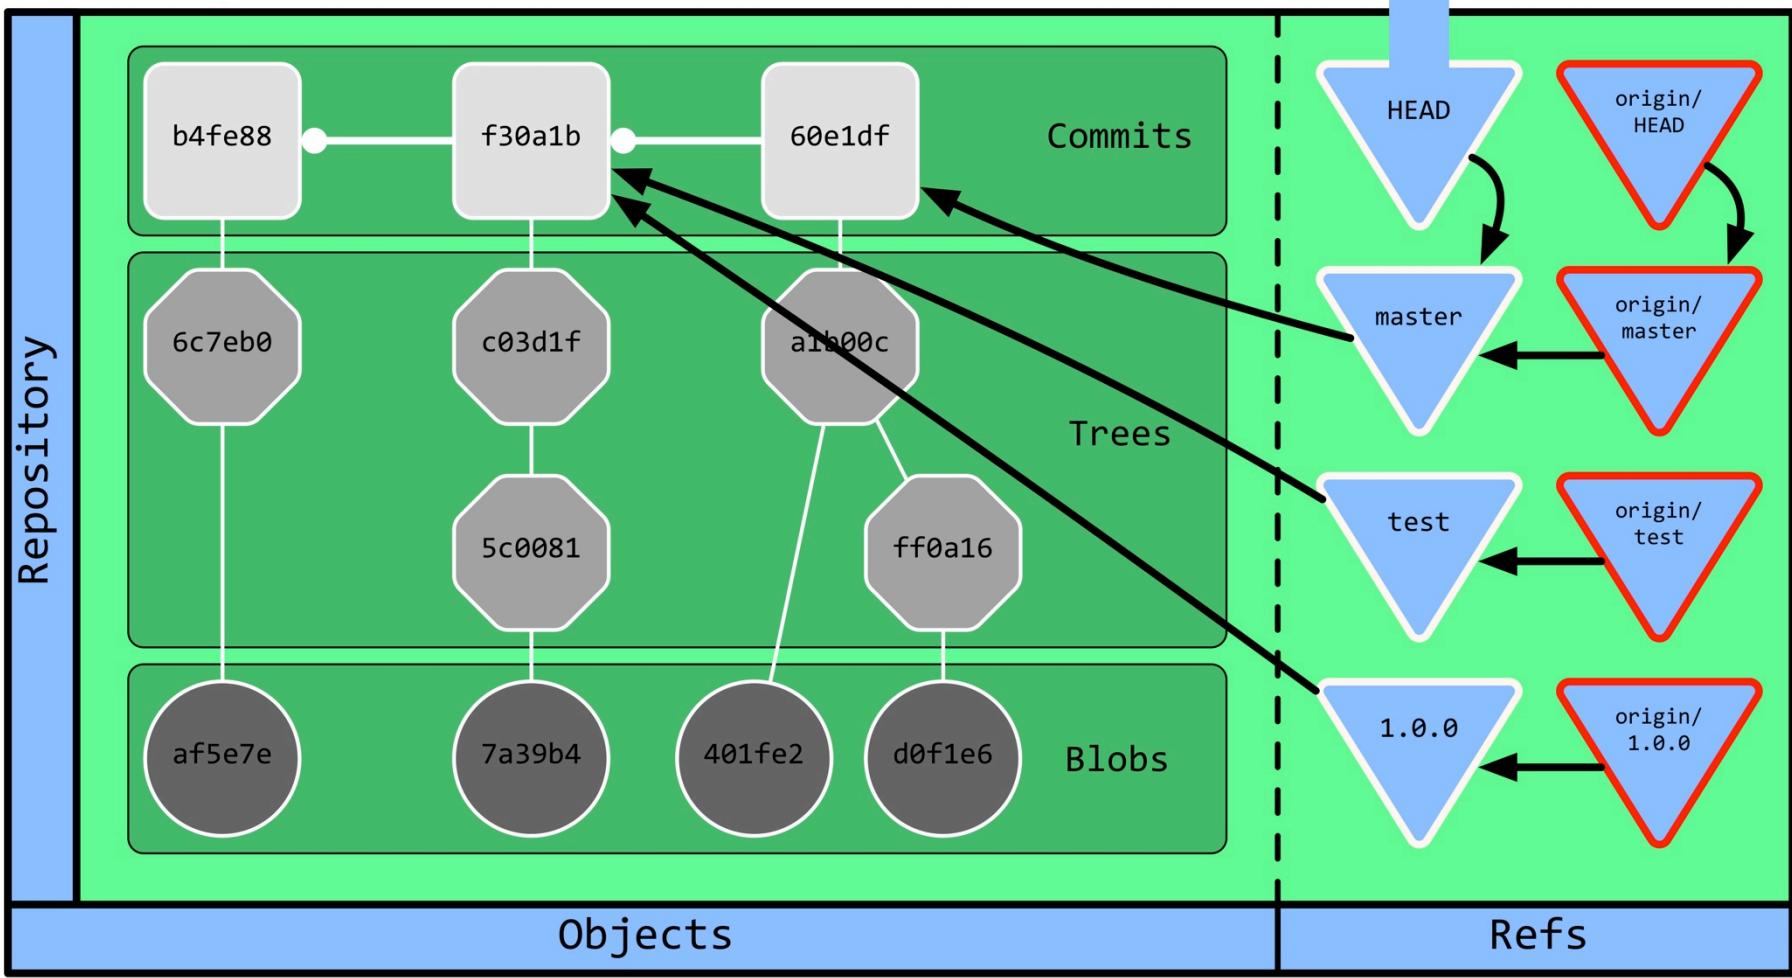
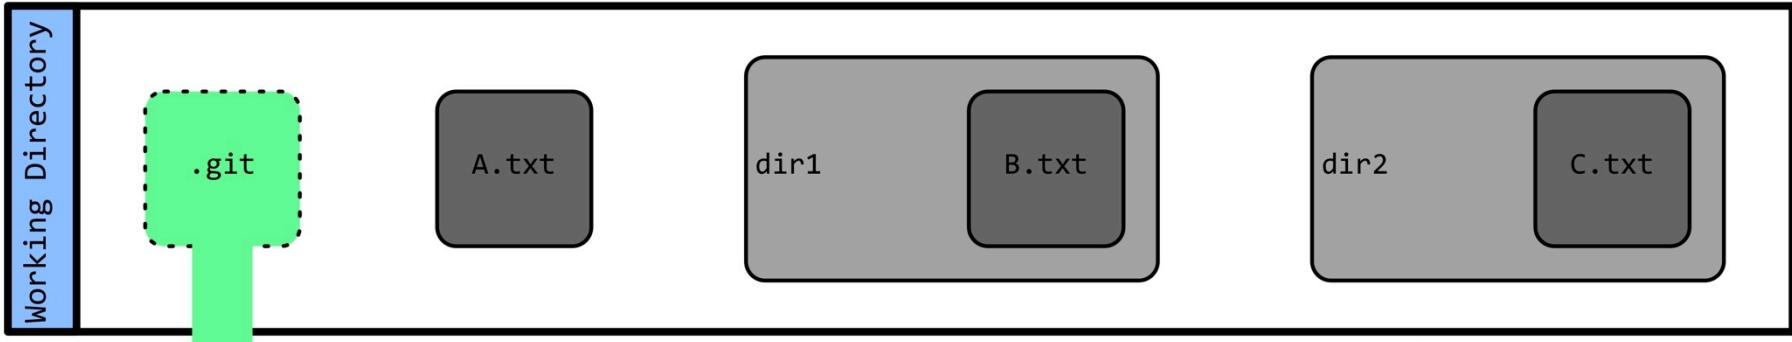




Finally, we'll “attach our HEAD” and go back to master

git checkout master





Questions?

References

- <https://git-scm.com/book/en/v2/Git-Internals-Git-Objects>
- <https://jwiegley.github.io/git-from-the-bottom-up/>